

**CTSpinoPelvic1K: spine, pelvis, ribs and femora in one coordinate frame, annotated
for lumbosacral transitional anatomy**

[Author names removed for double-anonymized review]

[Affiliations removed for double-anonymized review]

(Dated: 10 September 2026)

Purpose: A vertebra at the lumbosacral junction is named by counting. The gold standard counts caudally from C2 on whole-spine imaging¹; however, a lumbar surgical case is planned on lumbar-only imaging²⁻⁴ (T12 to S1) without C2. Abdominopelvic computed tomography (CT) holds that span plus the lowest ribs and pelvis, standing in for the preoperative view. Where a lumbosacral transitional vertebra (LSTV) alters the count, the local anatomy is ambiguous. Four rib-free vertebrae may be an L1 with a lumbar rib or an L5 assimilated to the sacrum, and six may be a sixth lumbar vertebra, a T12 with aplastic ribs, or a lumbarized S1⁵. CTSpinoPelvic1K asks whether local morphology resolves that ambiguity without the count. Two public label sets covered this imaging, CTSpine1K’s vertebrae and CTPelvic1K’s pelvis, on the same colonography patients, never joined; this release joins them on one series and annotates the bones neither had. It provides 802 CT records, per-level ribs and femora, lumbar levels anchored on the lowest rib-bearing vertebra and S1, and classes for L6, T13, a separate S1 and lumbar ribs, so anomalies are recorded as themselves.

Acquisition and Validation Methods: Records pair CTSpine1K⁶ and CTPelvic1K⁷ labels on each patient’s bone-richest acquisition under a VerSe-native scheme. Validation covered geometric invariants (802/802 pass), rib–vertebra incidence across 5,749 ribs (0.035% offset), and spinopelvic measures matching published values.

Data Format and Usage Notes: NIfTI image/label pairs with patient-grouped LSTV-stratified five-fold splits and a loader; archived at <https://doi.org/10.5281/zenodo.22139642>.

Potential Applications: Classifying a vertebra from local features; updating cadaveric morphometry; spinopelvic assessment; opportunistic screening; and, absent a public preoperative lumbar cohort, surgical planning research (377 records prone). *Limitations:* thoracic ground truth is field-of-view limited; postural angles are supine; no held-out test set; ribs are triaged-review pseudolabels; Castellvi grades are two-reader consensus on 33 records.

6 I. INTRODUCTION

7 The spine commonly carries seven cervical, twelve thoracic and five lumbar vertebrae, and
8 a vertebra is identified by where it falls in that sequence. Transitional variants sit at the
9 two ends of the lumbar column, the thoracolumbar transitional vertebra (TLTV) above and
10 the lumbosacral transitional vertebra (LSTV) below, and they disturb identification in two
11 ways at once. They change what the border vertebra and its ribs *look like*: a thoracolumbar
12 border vertebra may carry a full rib, a hypoplastic stump, or none at all, and a lumbosacral
13 one may be partly or completely assimilated to the sacrum (*sacralization*), with an anteriorly
14 wedged body, a reduced disc beneath it and hypoplastic facet joints, or conversely separated
15 from it (*lumbarization*), with a squared body, lumbar-type facets and a full-height disc⁵.
16 They also change *how many* vertebrae each region contains, eleven or thirteen thoracic and
17 four or six lumbar.

18 Either change alone would be tractable. Together they make it difficult to state defini-
19 tively what any vertebra at these borders *is*. Four rib-free vertebrae above the sacrum
20 may mean an L1 bearing a lumbar rib or an L5 assimilated to it; six may mean a true
21 sixth lumbar vertebra, a T12 whose ribs are aplastic, or an S1 lumbarized away from the
22 sacrum. Each set yields the same count (Fig. 2), and the morphology that would separate
23 the readings is not decisive to a human reader in every case. Whether it differs enough for
24 a learned classifier to separate them, or to assign each reading a probability, is the question
25 this release is built to support. Absent that, what remains is the count itself. The gold
26 standard for numeration in transitional anatomy is whole-spine imaging, counting caudally
27 from C2¹, which fixes how many vertebrae and how many rib pairs the column holds, and no
28 spine-limited acquisition can supply either count. LSTV is common, reported at 4–30% in
29 the general population⁵ and in 16.3% of a whole-spine CT series,⁸ and wrong-level surgery
30 is its most serious consequence.

31 Existing public collections cannot address this, and size is not the reason (Table I): spine
32 datasets omit the pelvis, pelvic datasets do not number the vertebrae, TotalSegmentator
33 has no class for a sixth lumbar vertebra or for a rib on a lumbar vertebra, and VerSe, the
34 one collection with an L6 class, stops at the sacrum and carries no ribs. No public collection
35 can therefore record a six-lumbar spine together with both of its borders, and a segmenter
36 trained on them must, faced with an enumeration anomaly, shift the whole column by a level

or absorb a vertebra into its neighbor. The size of that problem has now been measured on CT, where prior labeling methods assigned every vertebra correctly in 77% of subjects and an anomaly-aware extension of SPINEPS⁹ raised that to 99%¹⁰. Its weights label L6 and T13 on CT, but its 1,536-scan CT cohort is in-house and unreleased, so the capability exists as weights, not as data.

The intraoperative side has prior art of its own. LevelCheck registers the intraoperative radiograph to the preoperative CT and projects the CT’s vertebral labels onto it^{11–13}, but those labels are placed by hand and verified by the surgeon¹¹, so the method assumes a correctly labeled CT and does not supply one. Automating that step where the anatomy is anomalous is what a dataset carrying the anomaly classes is for.

CTSpinoPelvic1K places the spine, pelvis, ribs and femora in one coordinate frame and gives a class to each structure an enumeration anomaly produces, a sixth lumbar vertebra and a rib borne by a lumbar vertebra. A scheme without those classes lacks a class vocabulary expressive enough to describe anatomic anomalies.

The project began with a surgeon’s question, whether a vertebra can be named from the anatomy in front of the reader, on the images a lumbar surgical case is actually planned on. By guideline those are lumbar-only imaging studies^{2,3,14}, T12 to S1, the span the ACR spine parameters state for the lumbar spine⁴, and whole-spine coverage is reserved for trauma with an identified injury^{15,16} and for deformity radiographs. An abdominopelvic CT runs from the diaphragm to the pelvic floor, so it holds that span together with the lowest ribs and the whole pelvis, and it lacks C2 exactly as the planning study does. Every record comes from one prospective trial protocol, ACRIN 6664^{17,18}, which scanned 802 patients aged 50 and over, supine and prone, on multidetector CT at 15 centers, on five scanner vendors and nine models, with manufacturer, model and reconstruction kernel recorded per record. To our knowledge it is the largest spine-and-pelvis annotated CT cohort acquired under a single protocol, where the comparators pool collections, are multi-site by design, or are routine clinical scans under many protocols (Table I). With the protocol held fixed, scanner effects on a model become measurable rather than confounded. The price is an abdominal field of view, in which only the lowest thoracic levels are present (Sec. V).

The gap was in plain sight. CTPelvic1K⁷ and CTSpine1K⁶, released months apart by an overlapping author group, each annotated the COLONOG collection under radiologist supervision, one the pelvis and the other the vertebrae, on the same patients, and no one

69 joined them, although the join needs no new imaging and no new radiologist. It is not a
70 file merge, because each annotation had to be traced back to the series it was drawn on
71 (Sec. II.A). Done once, that crosswalk yields a spine-and-pelvis frame at 802 patients, and
72 the bones neither set had, ribs, femora, S1 and hardware, could then be annotated against
73 it rather than from scratch.

74 I.A. Vertebra labeling when the scan cannot settle the count

75 *The limitation, stated plainly.* No scan in this corpus contains C2, so the conventional
76 count cannot be performed, a property of abdominal imaging rather than of the annotation.
77 A thirteenth thoracic vertebra and a lumbar rib are different phenotypes, distinguished as
78 separate subtypes in cadaveric classifications of the thoracolumbar junction^{19,20}: the first is
79 an additional rib-bearing segment, usually above five lumbar vertebrae, so the column holds
80 25 presacral vertebrae; the second is a rudimentary rib on a lumbar-type L1 in a column
81 of the usual 24. Between the two borders used here, the last rib-bearing vertebra and the
82 sacrum, none of them is separable by counting: four rib-free vertebrae fit a lumbar rib, a
83 cranially shifted T13 or an L5 assimilated to the sacrum, and six fit an aplastic twelfth
84 rib or a sixth lumbar vertebra. A *stump rib*, hypoplastic and a cardinal indicator of a
85 thoracolumbar transitional vertebra, is still a rib, so only its length records the anomaly.
86 Rib anomalies travel with the lumbosacral border: on whole-spine CT hypoplastic twelfth
87 ribs occur with sacralization and lumbar ribs with lumbarization.⁸

88 *What the release does in addition.* Every vertebra carries the identifier its radiologist-
89 sourced annotation assigned, corrected where the source was wrong; those labels are the
90 ground truth. Every released measurement is also expressed against the two landmarks
91 present in essentially every abdominal scan, the sacrum below and the lowest rib-bearing
92 vertebra above, so that the measurement is positional rather than nominal: a vertebra by
93 how many rib-free vertebrae separate it from the sacrum, a lowest rib by its length as a
94 fraction of the rib above it, a transverse process by its span and its distance from the ala.
95 None of these changes with the name a reader gives the junction, so cases that disagree
96 about the name remain comparable.

97 *But the vertebrae themselves are known to differ.* Level-specific morphometry is long
98 established: pedicle width and height change systematically from the thoracic to the lumbar

TABLE I Public CT collections at the lumbosacral junction. Black, class present; outlined, graded only to exclude it²⁹.

Collection	Scans	Count	L6	Sacrum	S1	Pelvis	Ribs	L. rib	Femora	Grade
CTSpine1K ⁶	1,005									
CTPelvic1K ⁷	1,184									
VerSe ³⁰	374									
RibSeg v2 ³¹	660									
TotalSegmentator ³²	1,204									
CTSpinoPelvic1K	802									

spine,²¹ vertebral body, endplate and canal dimensions differ by level,²⁴⁻²⁶ and the transverse process and the iliolumbar ligament mark the last lumbar vertebra independently of any count.^{5,27} If those differences hold at the boundary itself the phenotype is decidable locally, and on this corpus it does. Separating T12 from L1 on shape alone gives an area under the curve of 0.990 and 97.3% accuracy over 1485 vertebrae: a logistic regression on five released per-level measures and six ratios between them, each measure divided by the same patient's median across their levels, so that size, which separates thoracic from lumbar trivially and says nothing at the junction, is removed; folds are grouped by case and the curve is taken on pooled out-of-fold scores (`test_morphometric_separability.py` in the released code). Schinz et al. reach the same conclusion from the other side: on 1,242 whole-thoracolumbar CTs a shape-based labeling of the junction matched nerve morphology in every case, against 92.6–97.2% for counting- and rib-based rules.²⁸

II. ACQUISITION AND VALIDATION METHODS

II.A. Sources, and why a crosswalk was necessary

Both sources draw part of their imaging from The Cancer Imaging Archive (TCIA) CT colonography collection (COLONOG)^{17,18}. CTSpine1K's 1,005 volumes come from four collections and CTPelvic1K's 1,184 from seven; COLONOG is the only collection common to both and the only one acquired under a single protocol, meaning one scan prescription and patient preparation for every case, so this release is its COLONOG subset, 782 patients




FIG. 1 How the dataset was built. n counts scans, or corrections, at each step.

with a CTSpine1K vertebral annotation and 20 with a CTPelvic1K pelvic annotation only, 802 in all. Neither source released the mapping from an annotation to the specific CT series it was drawn on.

That mapping matters because every patient in CT COLONOGRAPHY was scanned *twice*, prone and supine, and an annotation carries a patient identifier and nothing finer, so neither collection recorded which acquisition it had labeled. Turning a patient from supine to prone alters lumbar lordosis and the alignment of each segment against the next, so outlines drawn on one series are the wrong *shape* for the other and no rigid realignment recovers them. The crosswalk (Fig. 1) therefore resolves each annotation to a patient and then, within that patient's volumes, to the series whose bone *agrees* with it, where candidates are thresholded at bone attenuation and scored by overlap with the mask. It mattered for

129 351 patients, nearly half the cohort, whose two collections had labeled different acquisitions
130 of the same person; those records are released but held out of any analysis that assumes a
131 single posture, two labeled postures of one patient being a resource rather than a defect.

132 II.B. Densification, and why it is asymmetric

133 A corpus in which half the records lack a pelvis is not a spinopelvic dataset, so the missing
134 structures were pseudolabeled, asymmetrically (Fig. 1). Pelvis onto spine-only records is the
135 safe direction, since the sacrum and innominate bones carry no enumeration to get wrong,
136 and quality is reported as held-out performance on the *pelvic_only* records, withheld from
137 training for that purpose. The reverse direction carries the whole problem this dataset
138 addresses, because a pseudolabeled spine must commit to a count, and on a transitional
139 vertebra it commits to the commoner reading; those 20 records were corrected by hand.

140 II.C. Ribs

141 No public rib annotation covers this cohort, and the two available automatic sources^{32,33}
142 fail in complementary ways.

143 Möller’s binary rib network³³ reports high accuracy, and its authors show that stump
144 ribs can be classified from a partial field of view; that result concerns the morphological
145 features computed on a rib once segmented. On this corpus the segmentation itself failed
146 on ribs that were *partially outside* the field of view. A rib crossing the boundary of an
147 abdominal acquisition was liable to be missed altogether rather than segmented to the edge.
148 In a colonography cohort the lower ribs are routinely clipped, and those are the ribs the
149 count depends on.

150 TotalSegmentator³² does segment clipped ribs and supplies the per-rib numbering that a
151 binary network cannot. Its characteristic failure is at the other end of the bone, where the
152 most medial portion of the rib (roughly the last tenth to sixth approaching the costovertebral
153 joint) is frequently absent. A rib truncated at its medial end never reaches the vertebra it
154 articulates with, so the rib–vertebra incidence check (the quality-control gate on the rib class
155 labels, which matches each rib to the vertebra its number implies) has nothing to test.

156 The two were therefore unioned (Fig. 1), TotalSegmentator supplying the numbering and

the medial-clipped cases, Möller the detection where its output is more complete.

The result is a pseudolabel whose review was triaged rather than exhaustive. A label-based rule flagged every record in which one rib bone carried more than one identifier, or in which a rib failed to reach the vertebra it articulates with, since a single bone carrying two numbers is a numbering error whatever the anatomy. The rule selected 152 records; 148 were reviewed and corrected, and four were not and are identified in the release. Reviewers were medical students working through a student research consortium³⁴, trained against a written labeling protocol, with every correction attributed to its annotator. Records passing the rule were not individually inspected; the residual rib-vertebra offsets reported in Sec. II.F are the independent check on what the rule missed.

II.D. Label scheme and counting anchors

The scheme is VerSe-native. Vertebrae keep their VerSe identifiers (C1–C7 = 1–7, T1–T12 = 8–19, L1–L6 = 20–25, sacrum 26, coccyx 27, T13 = 28), and every non-VerSe structure takes a fixed identifier above that range, S1 carved from the sacrum (29), hips (30–31), femora (32–33), thirteen ribs per side (34–46 left, 47–59 right), lumbar ribs (60–61) and surgical hardware (62–68). The space is contiguous. Rib 13 is the true rib of a T13 and is empty in this release, so a thirteenth thoracic vertebra and a lumbar rib are recorded as the distinct phenotypes they are.

Two anchors give every record the same positional frame, and Fig. 2 shows both, the lowest rib-bearing vertebra above, **S1** carved from the sacrum below. S1 supplies the sacral endplate landmark that sacral slope and pelvic incidence are measured from, and with both brackets present L5 and L6 are distinguishable where a spine-limited view alone would not be.

How S1 is cut. The sacrum’s outer boundary is the source annotation and is never altered; S1 is the cranial part of that same bone, separated by a plane. The plane’s normal is the sacrum’s own cranio-caudal axis, made orthogonal to a left-right axis measured directly rather than assumed, so as to remove the patient’s rotation within the scanner, which is not negligible: a median of 4.5°, a maximum of 16.1°, and 508 of 801 records above 3°. The cut preserves the sub-division volume of the preceding release, changing the orientation of the S1/S2 boundary and not how much sacrum is called S1; per-record geometry ships with the




FIG. 2 The two anchors, from the released labels. Red, lowest rib-bearing vertebra and its rib; blue, S1; yellow, the rib-free vertebrae between. (b) and (c) both count four. In (b) a long rib sits on L1; in (c) the twelfth ribs are stumps and the segment below the fourth lumbar body is fused to the sacrum (Castellvi IIIb) and carries the S1 identifier.

archive.

VerSe³⁰ carries L6, whose identifier this release keeps, so its L6 *is* VerSe's.

A rib on a lumbar body receives its own class, and this is the one class here with no published counterpart. A lumbar rib and a stump rib are the same object under two counts, so a scheme that numbers every rib 1–12 leaves the annotator either calling it rib 12, asserting the vertebra beneath it is thoracic, the very question at issue, or discarding it. TotalSegmentator carries ribs 1–12 per side and no such class, and VerSe's T13 is a vertebra rather than a rib. This scheme records both: VerSe's T13 (28) with its rib pair (46, 59) for a thoracic-type thirteenth vertebra, and 60/61 for a lumbar-type body bearing a rudimentary rib. No record here carries a T13; the sixteen lumbar-rib cases were read as lumbar-type bodies, and where a border vertebra could not be settled by its readers the label stands as read.

Figure 2 shows why the interval is the primitive and the label is derived. Among the records with four rib-free vertebrae, some reach that count from above, through a rib on L1, and some from below, through a lumbosacral segment fused to the sacrum, in Fig. 2(c) together with stump twelfth ribs; the count alone cannot tell them apart, and the release can because ribs, rib lengths and the S1 boundary are all labeled.

Sixteen records carry a lumbar rib, thirteen bilateral and three unilateral, all on the right, too few to speak to side preference.

Eighteen records carry an L6. Seventeen of them carry a consensus Castellvi grade; the eighteenth was not graded. The source transitional label (pelvic where CTPelvic1K flagged

one, otherwise the lumbar count) reads lumbarization in fourteen, sacralization in two and semi-sacralization in one, and is unremarkable in the ungraded record. A spine can carry six lumbar-type vertebrae and still be read as a sacralization, depending on where the reader started the count.

Both counts, and the manifest fields that report them (`has_l6`, `has_lumbar_rib`, `n_lumbar_labels`, `lumbar_rib_side`), are computed from the released label volumes by counting identifiers 25 and 60–61.

II.E. Computational tools

Access. Every step is performed by open-source code archived with the dataset, together with the reference loader, the label scheme and the quality-control scripts that generated the tables in this manuscript.

Versions and parameters. Segmentation of femora, the sacral sub-division and the per-level rib numbering used TotalSegmentator³² (`total` task, release 2.x) on a single graphics processing unit. The binary rib network is Möller’s³³ released weights, applied through the nnU-Net v2 framework³⁵. The pelvic completion for records whose pelvis was absent used a five-fold nnU-Net v2 ensemble applied out-of-fold, so that no record is completed by a model that has seen it. The five fold checkpoints, with their nnU-Net plans and dataset descriptor, are released alongside the data (the repository URL identifies the authors; it is on the title page and will be restored here for the camera-ready version). Image handling throughout used NiBabel and SciPy; no image is resampled or reoriented when a label is written, so every label shares its image’s grid and affine exactly.

Measurement. Dimensions and spinopelvic parameters come from code released with the dataset. A sacral plate whose fitted normal lies more than 60° off the cranial axis is rejected and the case reported missing: such a fit has found the near-vertical anterior face of the promontory, not the endplate. Gating on the geometry rather than on the value leaves unusual patients in. an open-source measurement toolkit is released separately (the URL identifies the authors and is on the title page) as a reusable implementation measuring in each vertebra’s own frame; no open-source implementation computing these parameters from CT segmentations was found.

Generative artificial intelligence. A large language model (Claude, Anthropic) assisted

with drafting and editing the manuscript text and with writing the analysis and figure code,
all under author review. It was not used to generate, annotate or interpret imaging data;
every number reported here is computed by the released code from the released volumes.

II.F. Validation

Validation has three layers. Geometry is checked first, because a measurement taken on
a corpus with a geometric fault still returns a plausible-looking number.

Geometric invariants. Every case is checked for label geometry agreeing with its CT in
shape, affine and voxel spacing; for identifiers within the published scheme; for a non-empty
label; and for sided structures falling on the correct sides, with the left–right axis read from
the affine rather than assumed.

The check compares *each sided pair separately*. Testing the ribs alone passed all 802
records, four of which carried a `left_hip` label on the patient’s right side; pooling every
sided structure into one test still missed one of those four records, because a large correctly-
sided structure masks a smaller transposed one. A gate that passes everything is not evidence
of a clean corpus until it has been shown capable of failing.

Rib–vertebra incidence. Each rib is matched to its nearest vertebra and compared against
the vertebra its number implies. Of 11,548 ribs, 5,788 are not evaluable because the expected
vertebra lies outside the field of view and 11 have no vertebra within the anchor distance;
of the 5,749 evaluable, 5,747 match, two (0.035%) are offset by one level, and no case is
misnumbered. The denominator is stated because quoting the two offsets against all 11,548
ribs would halve the rate by counting ribs the check never examined. Both offsets are
field-of-view truncations, one at +1 level and one at −1.

The same check carries an invariant that reads as a null result. Of the 5,749 evaluable
ribs, *zero* are assigned to a lumbar vertebra, not because no lumbar ribs exist, since 16
records carry one, but because such a rib takes class 60 or 61 and never enters the numbered
set the test examines. A numbered rib landing on a lumbar vertebra would be a counting
error of exactly the kind this dataset exists to expose.

Agreement with published values. Pelvic incidence is the strongest of these checks because
it is a morphological property rather than a posture; it agrees to within half a degree with
the automated supine CT series of Veilleux *et al.*²³ and sits below the standing figure, the

TABLE II Spinopelvic measures, median over the 575 of 802 records passing the geometric gates, against automated supine CT.²³ Standing values are 55.0, 41.0, 13.0 and 60.0°.³⁶

Measure	This dataset (supine, $n = 575$)	Veilleux (supine CT)
Pelvic incidence	52.6°	52.1
Sacral slope	33.1°	36.5
Pelvic tilt	19.6°	15.6
Lumbar lordosis	49.3°	—
PI – LL mismatch	4.1°	—

direction reported for subjects imaged both ways.²² Sacral slope and lumbar lordosis also sit below their standing references, as recumbency predicts, and pelvic tilt above.

II.G. Castellvi typing

Every record whose vertebral count departs from five carries a Castellvi grade³⁷ in the released manifest, 33 cases typed I–IV with the *a/b* unilateral–bilateral qualifier. The grade and the count are *different axes*. Grade IIIb occurs here at rib-free counts of four, five and six alike, and seven of the 33 graded cases carry a normal count of five.

Two radiology resident physicians graded every case independently and resolved their six disagreements by consensus; both reads and the consensus are in the manifest. The distinction the readers disagreed on, type II (articulation) against type III (bony fusion)⁵, is the one that matters clinically, and four of the six disagreements were III against II.

II.H. Surgical hardware, and why a threshold is not enough

Metal is trivial to detect and hard to interpret, and the difference matters here because **an iatrogenic fusion is indistinguishable from a congenital one to a distance measurement**. A cage-bridged interspace reads as “no gap” exactly as a fused transitional vertebra does.

II.H.0.a. Detection over-calls by a factor of five. A 2500 HU threshold inside a shell around the labeled skeleton flags 52 of the 802 records. All 52 were reviewed manually, confirming 11 and rejecting 41 (79%) as contrast, calcification and reconstruction artefact.




FIG. 3 Instrumentation in the release, one exemplar each, drawn through bone. (a) Hip arthroplasty, 66, $n = 8$. (b) Femoral-neck screws, 68, $n = 1$. (c) Sacroiliac screws, 67, $n = 1$. (d) Interbody cages, 63, $n = 1$. Bar, 5 cm.

Attenuation does not separate the two groups (both saturate); site and volume do, every confirmed implant at 2,586 mm³ or above and every rejection at 1,768 mm³ or below.

II.H.0.b. The subtype block had to be extended. Identifiers 62–65 name spinal instrumentation (generic, cage, screw-and-rod, plate) and cannot describe what this cohort holds, so **66 arthroplasty**, **67 sacroiliac screw** and **68 osteosynthesis** were added. Osteosynthesis holds parts of one bone together where arthroplasty replaces a joint, and a shape rule calls a femoral stem a rod. The confirmed cases are eight arthroplasties, one osteosynthesis, one sacroiliac fixation and one pair of interbody cages (Fig. 3).

II.H.0.c. Metal outranks bone, and that is a subtraction. In every confirmed case the implant was already inside a bone label, so naming it reclaims 1,538,852 voxels rather than adding them. In eight of these cases the femoral head that pelvic incidence and tilt are measured from is an implant.

III. DATA FORMAT AND USAGE NOTES

Each record is a gzipped NIfTI image/label pair sharing an identical 4×4 affine, canonicalized to PIR, requiring no resampling. Masks are NIfTI rather than DICOM, which is what volumetric segmentation tooling expects.

Splits are patient-grouped and LSTV-stratified five-fold, frozen and shipped with the data as `splits_5fold.json`. The five validation folds partition all 802 records, so the release provides cross-validation and *not* a held-out test set; users reporting a single held-out number should carve one out and state which records it holds.

Stratification is on the transitional subtype rather than a binary LSTV flag. Across all

TABLE III Data types and metadata, and the records populating each, from the released manifest.

Field	Records Available	
Identifiers and paths, configuration, match type, image/label alignment check	802	100.0%
Patient position, tube potential, section thickness, kernel	802	100.0%
Transitional label and class, annotation origin, instrumentation flags	802	100.0%
Spine series identifier; bone-fraction check	782	97.5%
Scanner manufacturer and model	768	95.8%
Sex	749	93.4%
Age	709	88.4%
Pelvic series identifier	713	88.9%
Castellvi grade (consensus of two readers)	33	4.1%

802 records the source transitional label reads lumbarization in 14, sacralization in 17 and semi-sacralization in two; separately, 18 records carry a sixth lumbar-type vertebra, nine carry only four lumbar labels, and 33 carry a consensus Castellvi grade. A record called sacralization by its four lumbar labels enters the same stratum as one called sacralization by the source label, so the two rare strata hold 15 records each and every validation fold carries three of each. That is the most stratification can guarantee, and it is why a fold-level metric on the rare classes carries an error bar far wider than its own decimal places.

The archive of record is Zenodo (<https://doi.org/10.5281/zenodo.22139642>). Code and documentation are on GitHub at the tagged release commit; any model-hub copy is a convenience mirror, not the archive.

IV. DESCRIPTIVE ANALYSIS

The cohort is a colorectal cancer screening population of 802 records: 393 recorded female and 345 male, 11 carrying DICOM’s *other* value and 53 no sex field, with age present for 709 (median 59, range 50–89). All 802 are counted in totals; the eleven outside the female/male dichotomy are excluded from sex-stratified measures.

Count-free measures (Fig. 4). Five rib-free vertebrae above the sacrum is typical; four and six are where transitional anatomy sits, and which is present does not by itself determine what to call it. The lowest-rib length ratio is bimodal, and its lower mode, below 0.33 of




FIG. 4 Count-free measures. (a) Vertebrae between the lowest rib and the sacrum. (b) Lowest rib length over the rib above. (c) Lowest lumbar transverse span against its gap to the ala; transitional labels marked.




FIG. 5 Spinopelvic measures against standing reference bands (mean $\pm$ SD)³⁶.

the rib above, marks a stump twelfth rib in 98 of the 788 records with a measurable pair (12.4%, against 12.6% on whole-spine CT⁸); 16 records carry a lumbar rib (2.0%, against 1.5%). The co-occurrence reported there is testable here only where the junction was read; stump ribs accompany four of the 14 source-labeled sacralizations and 94 of the other 758 records (odds ratio 2.8, $p = 0.09$, Fisher's exact test), and no lumbarization carries a lumbar rib.

Spinopelvic measures (Fig. 5). Pelvic incidence sits on the standing reference, as a morphological property rather than a posture should. Sacral slope falls below it and pelvic tilt above, the direction recumbency predicts.

Opportunistic measures. Every scan carries a vertebral bone density measurement at no additional dose.³⁸ L1 trabecular attenuation is reported at the published standard site.

V. POTENTIAL APPLICATIONS AND LIMITATIONS

The principal use is the *spinopelvic interface*. Spine and pelvis carry radiologist-sourced annotation on one coordinate frame in 802 records, making sacral morphology measurable against the lumbar column rather than in isolation, among them the sacral endplate and its slope, the alar corridors constraining S2-alar-iliac and iliac screw trajectories, and the L5–S1 geometry a lumbosacral fusion must cross. Pelvic incidence is derived from the segmented sacrum and femoral heads, so it is reproducible from the labels themselves.

The transitional layer serves that use rather than competing with it, since no trajectory can be planned across a junction whose segments cannot be named.

A proxy for the preoperative view. No public preoperative lumbar CT cohort exists that we know of, so this release is the closest available stand-in for the view a lumbar case is planned on: it holds the T12-to-S1 span with the lowest ribs and the pelvis, 377 of its 802 records were acquired prone, the position of posterior lumbar surgery, and it carries classes for the anomalies that make level identification fail. It is a screening population at colonography dose, not a surgical series, so it supports method development for level identification and instrumentation geometry rather than validation of a planning decision.

V.A. Three uses the released measurements support

Naming a level from its own geometry. Wrong-level spine surgery runs at roughly one in 3,110 spinal procedures, and among the predominant contributing factors is the transitional anatomy this cohort was assembled around.^{39,40} Per-level body height, canal width, endplate width, transverse-process span and wedge ratio are released for T11–L5 across all 802 records, and because the labels were assigned against the twelfth rib rather than by enumeration, a model trained on them is not learning the convention that fails.

Reference morphometry with its spread. The values a surgeon consults for endplate and canal dimensions and pedicle width come from cadaveric series of ten to thirty specimens^{24–26} redrawn in the textbooks as one curve per level.⁴¹ A standard error describes the mean, not the spread of individuals, so neither tells the reader whether the patient in front of them is unusual. The same measures are given here from 802 records with the distribution attached (Fig. 6), reproducing the classical pedicle widening into the lower lumbar spine and adding




FIG. 6 Morphometry by level, T11–L5, as median, interquartile range and 5th–95th percentile, with n at right. (a) Endplate width. (b) Canal, width and depth. (c) Transverse pedicle width, the mean of the two sides. Black lines are Panjabi’s cadaveric means^{24,25}; the endplate series does not reach T11–T12.




FIG. 7 Records carrying each vertebral level, cranial to caudal. Counts per identifier in Table S1.

the width of each distribution.

Patient-specific models. Finite-element studies increasingly apply patient-specific modeling to alignment planning and implant design.⁴² Such a model needs spine, pelvis and femoral heads in one frame, which is this release’s construction; in 351 patients two postures give a within-patient change in alignment against which a predicted postural response can be checked.

Limitations: coverage and cohort. Thoracic ground truth is field-of-view limited (Fig. 7) and does not reach T1. Postural angles are supine; pelvic incidence is not postural, but it is modality-sensitive, reading below radiography for the same subjects.²² The cohort is one acquisition protocol of a colorectal screening population aged 50 and over, so its distributions

376 should not be read as a surgical one.

377 *Label strength varies by structure, and the release does not average over that.* Vertebral
378 labels derive from radiologist-supervised source annotations, corrected where those sources
379 were wrong; pelvic labels on records that lacked one are pseudolabeled. The rib layer
380 is a pseudolabel whose review was triaged, so its guarantee is the absence of the failure
381 modes that rule detects. Transitional labels come from two independent sources and are not
382 uniformly adjudicated, which is why the measures reported here are count-free; the Castellvi
383 grades are a consensus of two radiology resident physicians. S1 is carved by an automatic
384 estimate of the S1–S2 boundary; in 100 records that estimate is implausible, making S1 more
385 than half the sacrum’s height or under fifteen percent of it, and those records are flagged
386 in the quality-control table for exclusion from any measurement that depends on the sacral
387 endplate.

388 *Absent structures, and no classifier at the lower border.* Sacrum, hips and femora are
389 present in all 802 records; one record carries no S1 and one no T12, each outside the field
390 of view. Nine records carry no L5 identifier because their source annotation counted four
391 lumbar vertebrae, all nine are Castellvi IIIb, and the fused transitional segment carries the
392 S1 identifier as the caudal anchor; the manifest names them. No shape-based classifier is
393 attempted at the lumbosacral junction, where 15 records in each rare stratum leave three
394 per fold, too few to estimate from; the 0.990 area under the curve reported above is the
395 thoracolumbar boundary, not this one.

396 VI. DISCUSSION

397 *Comparison with related material.* VerSe^{29,30} comes closest in intent but has no sacral
398 mask for the structure a Castellvi grade describes; this release segments it, as L6 or as a
399 carved S1. VerSe carries neither the pelvis and femora a spinopelvic measurement needs nor
400 prone acquisitions, so it could not have been the starting point here; what this release adds
401 over TotalSegmentator³² is the classes an anomaly needs, not more anatomy per scan.

402 *Future directions.* Applying the released weights to VerSe would add the S1 anchor and
403 the anomaly classes to a bone-kernel collection. A Castellvi read of the whole cohort would
404 test the co-occurrence of rib and lumbosacral anomalies at full power and supply the labels
405 a lumbosacral classifier needs.

406 **VII. CONCLUSION**

407 CTSpinoPelvic1K places spine, pelvis, per-level ribs and femora on one coordinate frame
408 in 802 records, and gives the anatomy that makes lumbar numbering ambiguous a class of
409 its own. Because both counting anchors are explicit, a level can be identified from a field of
410 view that cannot support the conventional count downward from C2.

411 **DATA AVAILABILITY**

412 The release described here is v10, permanently archived at [https://doi.org/10.5281/](https://doi.org/10.5281/zenodo.22642578)
413 [zenodo.22642578](https://doi.org/10.5281/zenodo.22642578); the concept DOI <https://doi.org/10.5281/zenodo.22139642> resolves
414 to the current version. The archive carries the loader, the label scheme, the quality-control
415 tables and the analysis code under an open-source licence. The repository URL identifies
416 the authors; it is on the title page and will be restored here for the camera-ready version.

417 **ACKNOWLEDGEMENTS**

418 Acknowledgements are given on the title page, per the journal’s double-anonymized re-
419 view policy. This work received no funding.

420 **CONFLICT OF INTEREST**

421 The authors have no relevant conflicts of interest to disclose.

422 **ETHICS**

423 This work uses publicly available, de-identified imaging and annotations derived from it.
424 The authors’ institutional review board determined that it is not human participant research
425 under 45 CFR 46 and requires no oversight.

REFERENCES

- ¹J. Lian, N. Levine, and W. Cho, A review of lumbosacral transitional vertebrae and associated vertebral numeration, *Eur. Spine J.* **27**, 995–1004 (2018), doi:10.1007/s00586-018-5554-8.
- ²T. A. Hutchins et al., ACR Appropriateness Criteria Low Back Pain: 2021 Update, *J. Am. Coll. Radiol.* **18**, S361–S379 (2021), doi:10.1016/j.jacr.2021.08.002.
- ³American College of Radiology, ACR–ASNR–ASSR–SPR practice parameter for the performance of computed tomography (CT) of the spine, revised 2022 (Resolution 23), American College of Radiology, Reston, VA. <https://www.asnr.org/wp-content/uploads/2019/06/CT-Spine-1.pdf>, 2022.
- ⁴American College of Radiology, ACR–ASNR–SABI–SSR practice parameter for the performance of magnetic resonance imaging (MRI) of the adult spine, revised 2023 (Resolution 6), American College of Radiology, Reston, VA. <https://www.asnr.org/wp-content/uploads/2019/06/MR-Adult-Spine.pdf>, 2023.
- ⁵G. Konin and D. Walz, Lumbosacral transitional vertebrae: classification, imaging findings, and clinical relevance, *AJNR Am. J. Neuroradiol.* **31**, 1778–1786 (2010), doi:10.3174/ajnr.A2036.
- ⁶Y. Deng et al., CTSpine1K: a large-scale dataset for spinal vertebrae segmentation in computed tomography, *Mach. Learn. Biomed. Imaging* **3**, 824–832 (2025), doi:10.59275/j.melba.2025-gf84.
- ⁷P. Liu et al., Deep learning to segment pelvic bones: large-scale CT datasets and baseline models, *Int. J. Comput. Assist. Radiol. Surg.* **16**, 749–756 (2021), doi:10.1007/s11548-021-02363-8.
- ⁸S. Nagata et al., Lumbosacral transitional vertebra on whole-spine CT: prevalence and association with rib abnormalities, *Skeletal Radiol.* **54**, 2169–2177 (2025), doi:10.1007/s00256-025-04952-z.
- ⁹H. Möller et al., SPINEPS—automatic whole spine segmentation of T2-weighted MR images using a two-phase approach to multi-class semantic and instance segmentation, *Eur. Radiol.* **35**, 1178–1189 (2025), doi:10.1007/s00330-024-11155-y.
- ¹⁰H. Möller et al., VERIDAH: Solving Enumeration Anomaly Aware Vertebra Labeling across Imaging Sequences, *arXiv*, 2026, arXiv:2601.14066.

doi:10.48550/ARXIV.2601.14066.

¹¹Y. Otake et al., Automatic localization of vertebral levels in x-ray fluoroscopy using 3D-2D registration: a tool to reduce wrong-site surgery, *Phys. Med. Biol.* **57**, 5485–5508 (2012), doi:10.1088/0031-9155/57/17/5485.

¹²S.-f. L. Lo et al., Automatic Localization of Target Vertebrae in Spine Surgery: Clinical Evaluation of the LevelCheck Registration Algorithm, *Spine (Phila. Pa. 1976)* **40**, E476–E483 (2015), doi:10.1097/BRS.0000000000000814.

¹³T. De Silva et al., Utility of the LevelCheck Algorithm for Decision Support in Vertebral Localization, *Spine (Phila. Pa. 1976)* **41**, E1249–E1256 (2016), doi:10.1097/BRS.0000000000001589.

¹⁴D. S. Kreiner et al., An evidence-based clinical guideline for the diagnosis and treatment of degenerative lumbar spinal stenosis (update), *Spine J.* **13**, 734–743 (2013), doi:10.1016/j.spinee.2012.11.059.

¹⁵S. Sixta et al., Screening for thoracolumbar spinal injuries in blunt trauma: an Eastern Association for the Surgery of Trauma practice management guideline, *J. Trauma Acute Care Surg.* **73**, S326–S332 (2012), doi:10.1097/TA.0b013e31827559b8.

¹⁶American College of Surgeons Committee on Trauma, ACS TQIP best practices guidelines in imaging, American College of Surgeons, Chicago. https://www.facs.org/media/oxdjw5zj/imaging_guidelines.pdf, 2018.

¹⁷K. Smith et al., Data from CT COLONOGRAPHY (ACRIN 6664), The Cancer Imaging Archive, version 2 [dataset], 2015, doi:10.7937/K9/TCIA.2015.NWTESAY1.

¹⁸K. Clark et al., The Cancer Imaging Archive (TCIA): maintaining and operating a public information repository, *J. Digit. Imaging* **26**, 1045–1057 (2013), doi:10.1007/s10278-013-9622-7.

¹⁹A. M. Du Plessis, L. M. Greyling, and B. J. Page, Differentiation and classification of thoracolumbar transitional vertebrae, *J. Anat.* **232**, 850–856 (2018), doi:10.1111/joa.12781.

²⁰A. M. Poolman, Q. Wessels, A. V. Schoor, and N. Keough, Thoracolumbar transitional vertebrae: Quantitative differentiation and associated numeric variation in the vertebral column using skeletal remains, *J. Anat.* **243**, 311–318 (2023), doi:10.1111/joa.13865.

²¹M. R. Zindrick et al., Analysis of the morphometric characteristics of the thoracic and lumbar pedicles, *Spine (Phila. Pa. 1976)* **12**, 160–166 (1987), doi:10.1097/00007632-198703000-00012.

²²C. M. Lee and R. W. Liu, Comparison of pelvic incidence measurement using lateral
x-ray, standard CT versus CT with 3D reconstruction, *Eur. Spine J.* **31**, 241–247 (2022),
doi:10.1007/s00586-021-07024-7.

²³N. J. Veilleux, N. V. Kalore, J. A. Vossen, and J. S. Wayne, Automatic characterization of
pelvic and sacral measures from 200 subjects, *J. Bone Joint Surg. Am.* **102**, e130 (2020),
doi:10.2106/JBJS.20.00343.

²⁴M. M. Panjabi et al., Thoracic human vertebrae: quantitative three-dimensional anatomy,
Spine (Phila. Pa. 1976) **16**, 888–901 (1991), doi:10.1097/00007632-199108000-00006.

²⁵M. M. Panjabi et al., Human lumbar vertebrae: quantitative three-dimensional anatomy,
Spine (Phila. Pa. 1976) **17**, 299–306 (1992), doi:10.1097/00007632-199203000-00010.

²⁶J. L. Berry, J. M. Moran, W. S. Berg, and A. D. Steffee, A morphometric study of human
lumbar and selected thoracic vertebrae, *Spine (Phila. Pa. 1976)* **12**, 362–367 (1987),
doi:10.1097/00007632-198705000-00010.

²⁷R. J. Hughes and A. Saifuddin, Numbering of lumbosacral transitional vertebrae on
MRI: role of the iliolumbar ligaments, *AJR Am. J. Roentgenol.* **187**, W59–W65 (2006),
doi:10.2214/AJR.05.0415.

²⁸D. Schinz et al., A vertebral shape-based classification for labeling the thoracolumbar
junction on CT imaging, *Eur. Spine J.* **35**, 2195–2203 (2026), doi:10.1007/s00586-025-
09364-0.

²⁹H. Liebl et al., A computed tomography vertebral segmentation dataset with anatomical
variations and multi-vendor scanner data, *Sci. Data* **8**, 284 (2021), doi:10.1038/s41597-
021-01060-0.

³⁰A. Sekuboyina et al., VerSe: a vertebrae labelling and segmentation bench-
mark for multi-detector CT images, *Med. Image Anal.* **73**, 102166 (2021),
doi:10.1016/j.media.2021.102166.

³¹L. Jin et al., RibSeg v2: a large-scale benchmark for rib labeling and
anatomical centerline extraction, *IEEE Trans. Med. Imaging* **43**, 570–581 (2024),
doi:10.1109/TMI.2023.3313627.

³²J. Wasserthal et al., TotalSegmentator: robust segmentation of 104 anatomic structures
in CT images, *Radiol. Artif. Intell.* **5**, e230024 (2023), doi:10.1148/ryai.230024.

³³H. Möller et al., Automated thoracolumbar stump rib detection and analysis in a large
CT cohort, *arXiv*, 2025, arXiv:2505.05004. doi:10.48550/arXiv.2505.05004.

³⁴[Reference removed for double-anonymized review: a published description of the authors' student research program.]

³⁵F. Isensee, P. F. Jaeger, S. A. A. Kohl, J. Petersen, and K. H. Maier-Hein, nnU-Net: a self-configuring method for deep learning-based biomedical image segmentation, *Nat. Methods* **18**, 203–211 (2021), doi:10.1038/s41592-020-01008-z.

³⁶R. Vialle et al., Radiographic analysis of the sagittal alignment and balance of the spine in asymptomatic subjects, *J. Bone Joint Surg. Am.* **87**, 260–267 (2005), doi:10.2106/JBJS.D.02043.

³⁷A. E. Castellvi, L. A. Goldstein, and D. P. Chan, Lumbosacral transitional vertebrae and their relationship with lumbar extradural defects, *Spine (Phila. Pa. 1976)* **9**, 493–495 (1984), doi:10.1097/00007632-198407000-00014.

³⁸P. J. Pickhardt et al., Opportunistic screening for osteoporosis using abdominal computed tomography scans obtained for other indications, *Ann. Intern. Med.* **158**, 588–595 (2013), doi:10.7326/0003-4819-158-8-201304160-00003.

³⁹M. G. Mody, A. Nourbakhsh, D. L. Stahl, M. Gibbs, M. Alfawareh, and K. J. Garges, The Prevalence of Wrong Level Surgery Among Spine Surgeons, *Spine (Phila. Pa. 1976)* **33**, 194–198 (2008), doi:10.1097/BRS.0b013e31816043d1.

⁴⁰N. Epstein, A perspective on wrong level, wrong side, and wrong site spine surgery, *Surg. Neurol. Int.* **12**, 286 (2021), doi:10.25259/SNI.402_2021.

⁴¹E. C. Benzell, *Biomechanics of Spine Stabilization*, Thieme, New York, 3rd edition, 2015, ISBN 978-1-60406-924-2.

⁴²E. Beaulieu et al., A review of finite element analysis in spine surgery decision-making, *J. Clin. Med.* **15**, 2584 (2026), doi:10.3390/jcm15072584.

FIGURE CAPTIONS

Figure 1. How the dataset was built. n counts scans, or corrections, at each step.

Figure 2. The two anchors, from the released labels. Red, lowest rib-bearing vertebra and its rib; blue, S1; yellow, the rib-free vertebrae between. (b) and (c) both count four. In (b) a long rib sits on L1; in (c) the twelfth ribs are stumps and the segment below the fourth lumbar body is fused to the sacrum (Castellvi IIIb) and carries the S1 identifier.

Figure 3. Instrumentation in the release, one exemplar each, drawn through bone. (a)

551 Hip arthroplasty, 66, $n = 8$. (b) Femoral-neck screws, 68, $n = 1$. (c) Sacroiliac screws, 67,
552 $n = 1$. (d) Interbody cages, 63, $n = 1$. Bar, 5 cm.

553 **Figure 4.** Count-free measures. (a) Vertebrae between the lowest rib and the sacrum. (b)
554 Lowest rib length over the rib above. (c) Lowest lumbar transverse span against its gap to
555 the ala; transitional labels marked.

556 **Figure 5.** Spinopelvic measures against standing reference bands (mean $\pm$ SD)³⁶.

557 **Figure 6.** Morphometry by level, T11–L5, as median, interquartile range and 5th–95th
558 percentile, with n at right. (a) Endplate width. (b) Canal, width and depth. (c) Transverse
559 pedicle width, the mean of the two sides. Black lines are Panjabi’s cadaveric means^{24,25}; the
560 endplate series does not reach T11–T12.

561 **Figure 7.** Records carrying each vertebral level, cranial to caudal. Counts per identifier in
562 Table S1.