

CTSpinoPelvic1K: spine, pelvis, ribs and femora in one coordinate frame, annotated for lumbosacral transitional anatomy

Article type: Dataset and Software Article

Authors

1. Gregory Schwing, MD, PhD ¹
2. Ashley Schehr, BS ²
3. Annika Tekumulla, BS ²
4. Margret Khoushi, BS ²
5. Ryan Christian, BS ²
6. Dane Hubers, BS ²
7. Faris Mahjoub, BS ²
8. Hassan Saad, BS ²
9. Mia Sooch, BA ²
10. Sathya Siddapureddy, BS ²
11. Michael McLellan, BS ²
12. Jerick Kim, BS ²
13. Miraziz Ismoilov, MD ³
14. Nizar Alnabahneh, MD ³

¹Department of Surgery, Detroit Medical Center and Wayne State University, Detroit, Michigan, United States of America

²School of Medicine, Wayne State University, Detroit, Michigan, United States of America

³Department of Radiology, Detroit Medical Center and Wayne State University, Detroit, Michigan, United States of America

Corresponding author

Gregory Schwing

Department of Surgery, Detroit Medical Center and Wayne State University

Detroit, Michigan, United States of America

Email: [corresponding author email]

Author contribution statement

G.S. designed the dataset and the label scheme, implemented the processing pipeline and the quality-control procedures, and drafted the manuscript. A.S., A.T., M.K., R.C., D.H., F.M., H.S., M.S., S.S., M.M. and J.K. performed annotation correction and quality-control review of the released segmentations under supervision. M.I. and N.A. performed the radiological reads — the Castellvi grading and the adjudication of candidate instrumentation. [BEFORE SUBMISSION: confirm each contribution individually with the author concerned and expand to CRediT roles.] All authors read and approved the final version of the manuscript.

Conflict of interest statement

The authors have no relevant conflicts of interest to disclose.

Acknowledgements

The imaging analyzed here was made available by The Cancer Imaging Archive and by the investigators who published the source annotations. [Funding sources to be listed.]

Data and code availability (unblinded)

Dataset, archived before submission: <https://doi.org/10.5281/zenodo.22139642> (concept identifier, resolves to the current version); the version described in the manuscript is v7, <https://doi.org/10.5281/zenodo.22242745>.

Code: <https://github.com/Gregory-Schwing-MD-PhD/CTSpinoPelvic1K>, tagged at the release commit.

Mirror: <https://huggingface.co/datasets/OpenSpineConsortium/CTSpinoPelvic1K> — a convenience mirror, not the archive of record.

Ethics

This work uses publicly available, de-identified imaging and annotations derived from it. The Wayne State University Institutional Review Board issued a Notice of IRB Administrative Determination of Non-Human Participant Research on 2 September 2026 (WSU IRB HPR Number 2026-269, “Open Spine Consortium: Secondary Analysis of Publicly Available, De-Identified Imaging Datasets and the Derived Annotations, Software, and Models Produced From Them”), finding that the project does not constitute human participant research under the Common Rule at 45 CFR 46 and that IRB review and oversight are not required.

**CTSpinoPelvic1K: spine, pelvis, ribs and femora in one coordinate frame, annotated
for lumbosacral transitional anatomy**

[Author names removed for double-anonymized review]

[Affiliations removed for double-anonymized review]

(Dated: 3 September 2026)

Purpose: Lumbar levels cannot be numbered reliably when a lumbosacral transitional vertebra (LSTV) is present. Levels are conventionally established by counting downward from C2, the reference standard for numeration¹. An abdominal field of view lacks that landmark, leaving a specific ambiguity: four rib-free lumbar vertebrae may mean an L1 bearing a lumbar rib or an L5 assimilated to the sacrum, and six may mean a true sixth lumbar vertebra or a T12 whose ribs are aplastic². CT-SpinoPelvic1K is built to test whether local morphology can resolve them without counting from the C2 anchor. It provides 802 CT records anchoring each lumbar level on the lowest rib-bearing vertebra above and the first sacral segment below, with radiologist-sourced spine *and* pelvic annotation on one frame, per-level ribs and femora. Classes for a sixth lumbar vertebra, a thirteenth thoracic vertebra, a separate first sacral segment, and lumbar ribs make the set expressive enough for anatomical anomalies.

Acquisition and Validation Methods: Records pair TCIA CT COLONOGRAPHY imaging with CTSpine1K³ vertebral and CTPelvic1K⁴ pelvic labels on the acquisition with highest bone coverage, under a VerSe-native scheme. Validation: geometric invariants (802/802 pass); rib-vertebra incidence across 5,749 evaluable ribs (2 offsets, 0.035%); and agreement of spinopelvic measures with published values.

Data Format and Usage Notes: Gzipped NIfTI image/label pairs sharing one affine, canonicalized to PIR, with frozen patient-grouped LSTV-stratified five-fold splits and a reference loader; archived at 10.5281/zenodo.22139642.

Potential Applications: Testing whether a network can classify a vertebra from local features alone; updating textbook morphometry drawn from small cadaveric series; spinopelvic assessment at the lumbosacral interface; variant studies; opportunistic screening. *Limitations:* thoracic ground truth is field-of-view limited; postural angles supine; no held-out test set; the rib layer is a triaged-review pseudolabel; Castellvi grades assigned by two radiology resident physicians.

6 I. INTRODUCTION

7 The spine commonly carries seven cervical, twelve thoracic and five lumbar vertebrae, and
8 a vertebra is identified by where it falls in that sequence. Transitional variants sit at the two
9 ends of the lumbar column — the thoracolumbar transitional vertebra (TLTV) above, the
10 lumbosacral transitional vertebra (LSTV) below — and they disturb identification in two
11 ways at once. They change what the border vertebra and its ribs *look like*: a thoracolumbar
12 border vertebra may carry a full rib, a hypoplastic stump, or none at all, and a lumbosacral
13 one may be partly or completely assimilated to the sacrum, with an anteriorly wedged body,
14 a reduced disc beneath it and hypoplastic facet joints, or conversely separated from it with
15 a squared body, lumbar-type facets and a full-height disc². They also change *how many*
16 vertebrae each region contains — eleven or thirteen thoracic, four or six lumbar.

17 Either change alone would be tractable. Together they make it difficult to state defini-
18 tively what any vertebra at these borders *is*. Four rib-free vertebrae above the sacrum may
19 mean an L1 bearing a lumbar rib or an L5 assimilated to it; six may mean a true sixth
20 lumbar vertebra or a T12 whose ribs are aplastic. Each pair yields the same count, and the
21 morphology that would separate the two readings is not decisive to a human reader in every
22 case. Whether it nonetheless differs enough for a learned classifier to separate them, or at
23 least to assign each reading a probability, is the question this release is built to support.
24 Absent that, what remains is the sequence — conventionally established by counting down-
25 ward from C2, the reference standard for numeration¹, and unavailable in any spine-limited
26 acquisition. LSTV is common, reported between 4% and 30% depending on definition, and
27 wrong-level surgery is its most serious consequence.

28 Existing public collections cannot address this, and size is not the reason (Fig. 1). Spine
29 datasets omit the pelvis; pelvic datasets do not number the vertebrae; and TotalSegmentator,
30 the whole-body scheme, has no class for a sixth lumbar vertebra or for a rib on a lumbar
31 vertebra. A six-lumbar spine therefore cannot be recorded in any of them, however good
32 the segmentation. VerSe^{5,6} makes the point most clearly (Fig. 1b): it deliberately enriched
33 for anatomical variants and graded them by Castellvi, yet left the graded vertebrae fused to
34 the sacrum unsegmented.

35 CTSpinoPelvic1K places the spine, pelvis, ribs and femora in one coordinate frame and
36 gives a class to each structure an enumeration anomaly produces: a sixth lumbar vertebra,

37 and a rib borne by a lumbar vertebra. A scheme without those classes must record such
38 anatomy as something it is not.

39 *Where the material is.* Everything described here is v7 of a single archived dataset,
40 deposited at <https://doi.org/10.5281/zenodo.22242745>. Citations should use <https://doi.org/10.5281/zenodo.22139642>, the permanent identifier for the dataset across all
41 of its versions, which resolves to whichever version is current; the version identifier above
42 pins the exact release measured in this manuscript. Documentation, the label scheme, the
43 reference loader and the per-record quality-control tables are distributed with the archive
44 and are described in Sec. III.

46 I.A. Vertebra labeling when the scan cannot settle the count

47 *The limitation, stated plainly.* No scan in this corpus contains C2 or the whole cervical
48 spine, so the conventional count cannot be performed. That is a property of abdominal
49 imaging rather than of the annotation, and it means a variant at the thoracolumbar junction
50 cannot be *resolved by counting* here at all: a thirteenth thoracic vertebra, a lumbar rib, and
51 a *stump rib* — a hypoplastic twelfth rib, one of the cardinal indicators of a thoracolumbar
52 transitional vertebra — produce overlapping appearances.

53 *What the release does instead.* No vertebral label is asserted where it cannot be supported.
54 Every quantity is expressed against the two landmarks present in essentially every abdominal
55 scan — the sacrum below and the lowest rib-bearing vertebra above — so the description is
56 positional rather than nominal: a vertebra by how many rib-free vertebrae separate it from
57 the sacrum, a lowest rib by its length as a fraction of the rib above it, a transverse process
58 by its span and its distance from the ala. None of these changes according to whether a
59 reader would have said L5 or S1, so cases that disagree about the name remain comparable.
60 Where the junction is in view the count is determined and the conventional name follows,
61 and the two descriptions agree. The positional description is, by construction, blind to the
62 ambiguity: an L1 bearing a lumbar rib and an L5 assimilated to the sacrum both leave
63 four rib-free vertebrae above the sacrum, and a true sixth lumbar vertebra and a T12 with
64 aplastic ribs both leave six. The count alone cannot separate the readings.

65 *But the vertebrae themselves are known to differ.* Counting is not the only route to an
66 identity. Level-specific morphometry is long established: pedicle width and height change




FIG. 1 Public CT collections at the lumbosacral junction. (a) Size, stated first because it is the claim this dataset does *not* make: three of the five comparators carry more scans, and RibSeg v2¹² labels more ribs. (b) What each published scheme can *represent*, a property of its class list rather than of its segmentation quality. VerSe is graded-but-excluding: it applies the Castellvi classification and uses it to decide what to omit, leaving vertebrae partially fused to the sacrum unsegmented⁶ — 25 of the 33 graded records here.

systematically from the thoracic to the lumbar spine,⁷ vertebral body, endplate and canal dimensions differ by level,^{8–10} and the transverse process and the iliolumbar ligament mark the last lumbar vertebra independently of any count.^{2,11} If those differences hold at the boundary itself, the phenotype is decidable from the vertebra in front of the reader without any global count. On this corpus it does hold: separating T12 from L1 on shape alone, with each patient’s own size divided out and cross-validation grouped by case, gives an area under the curve of 0.990 and 97.3% accuracy over 1485 vertebrae. Size is deliberately removed, because a classifier given raw millimeters separates thoracic from lumbar immediately and has learned nothing that helps at the junction, where the two are nearly the same size.

That is reported as a property of the released measurements rather than as a method: the information needed to settle these variants is present locally.

II. ACQUISITION AND VALIDATION METHODS

II.A. Sources, and why a crosswalk was necessary

Both CTSpine1K³ and CTPelvic1K⁴ draw part of their imaging from the TCIA CT colonography collection^{13,14}, and the vertebral and pelvic annotations in both were produced under board-certified radiologist supervision. Neither released the mapping from an annotation to the specific CT series it was drawn on.

That mapping is not a bookkeeping detail. Every patient in CT COLONOGRAPHY was scanned *twice*, prone and supine, often with more than one kernel, so one patient identifier resolves to several volumes differing in posture, position, field of view and noise — and an annotation carries a patient identifier and nothing finer.

Placing it on the wrong volume of the same patient is not merely a displacement. The spine is mildly deformable, and turning a patient from supine to prone alters lumbar lordosis and the alignment of each segment against the next, so a vertebra moves by a different amount, and in a different direction, from the one above it. Outlines drawn on one series are the wrong *shape* for the other, and no rigid realignment recovers them. In isolation such a mask looks like competent work; it reveals itself only when overlaid on the image.

The crosswalk (Fig. S1) therefore resolves each annotation first to a patient and then, within that patient’s volumes only, to the series whose bone *agrees* with it: candidates are thresholded at bone attenuation and scored by overlap with the annotation mask. Bone rather than image similarity, because bone is what the annotation describes — two acquisitions of one abdomen resemble each other everywhere except in where the skeleton sits.

II.B. Fused and split records

The two annotation sets were placed independently, so they do not always land on the same series. Of 802 records, 342 carry spine and pelvic annotations on the same volume (*fused*), 351 are *separate*, 89 are *spine-only* with no pelvic annotation for that patient, and 20 are *pelvic-only*. The split is the crosswalk’s main finding rather than bookkeeping: in those separate records — nearly half of all patients — the two public collections annotated *different volumes of the same person*, and neither recorded which. A pelvic annotation placed on the other acquisition of a prone/supine pair cannot be combined with the spine

107 annotation without resampling across a change of posture. Those records are retained rather
108 than discarded, and held out of the postural analyses, because two acquisitions of one patient
109 differing in spinal alignment are a resource for a mobility study rather than a defect.

110 II.C. Densification, and why it is asymmetric

111 A corpus in which half the records lack a pelvis is not usable as a spinopelvic dataset,
112 so the missing structures were pseudolabeled — asymmetrically (Fig. S1), because the two
113 directions are not equivalent.

114 Pelvis onto spine-only records is the safe direction. The sacrum and innominate bones
115 are single objects with no enumeration: no count to get wrong, no naming convention to
116 choose, no analog of a transitional vertebra. Quality is reported as held-out performance on
117 the *pelvic_only* records, withheld from training for exactly this purpose.

118 The reverse direction carries the whole problem this dataset exists to address, since a
119 pseudolabeled spine must commit to a vertebral count and on a transitional vertebra will
120 commit to the commoner reading. Those 20 records were therefore corrected by hand rather
121 than densified automatically — tractable at twenty cases, which is the practical reason the
122 asymmetry runs the way it does.

123 II.D. Ribs

124 No public rib annotation covers this cohort, and the two available automatic sources fail
125 in complementary ways.

126 Möller’s binary rib network¹⁵ reports high accuracy, and its authors show that stump ribs
127 can be classified from a partial field of view; that result concerns the morphological features
128 computed on a rib once segmented. On this corpus the segmentation itself failed on ribs
129 that were *partially outside* the field of view: a rib crossing the boundary of an abdominal
130 acquisition was liable to be missed altogether rather than segmented to the edge. This is
131 not a limitation its authors report — their evaluation excluded subjects in which the last
132 rib was not visible, and their failure analysis attributes errors to costal-process fusion — so
133 it is recorded here as an observation from this cohort. Since this is a colonography cohort,
134 the lower ribs are routinely clipped, and those are precisely the ribs the count depends on.

135 TotalSegmentator¹⁶ does segment clipped ribs and supplies the per-rib numbering that a
 136 binary network cannot. Its characteristic failure is at the other end of the bone: the most
 137 medial portion of the rib — roughly the last tenth to sixth approaching the costovertebral
 138 joint — is frequently absent. A rib truncated at its medial end never reaches the vertebra it
 139 articulates with, so the rib–vertebra incidence check of Table I — the quality-control gate
 140 on the rib class labels, which matches each rib to the vertebra its number implies — has
 141 nothing to test.

142 The two were therefore unioned (Fig. S1): TotalSegmentator supplies numbering and the
 143 medial-clipped cases, Möller supplies detection where its output is more complete, and the
 144 union carries TotalSegmentator’s per-level identities.

145 The result is a pseudolabel whose review was triaged rather than exhaustive, which is
 146 worth stating plainly because a layer described only as “human-corrected” invites the reader
 147 to assume every bone was looked at. A label-based rule flagged every record in which one
 148 rib bone carried more than one identifier, or in which a rib failed to reach the vertebra it
 149 articulates with — a single bone with two numbers being a numbering error by construction,
 150 whatever the anatomy. The rule selected 152 records; 148 were reviewed and corrected, and 4
 151 were not and are identified in the release. Reviewers were medical students working through
 152 OpenSpineConsortium¹⁷, trained against a written labeling protocol, with every correction
 153 recorded against the annotator who made it. Records passing the rule were not individually
 154 inspected, so the layer’s guarantee is that it is free of the failure modes the rule detects,
 155 and no stronger. The residual rib–vertebra offsets reported in Sec. II.G are the independent
 156 check on what the rule missed.

157 II.E. Label scheme and counting anchors

158 The scheme is VerSe-native: vertebrae keep their VerSe identifiers (C1–C7 = 1–7, T1–T12
 159 = 8–19, L1–L6 = 20–25, sacrum 26), and every non-VerSe structure takes a fixed identifier
 160 above that range — S1 carved from the sacrum (29), hips (30–31), femora (32–33), ribs
 161 (34–45 left, 46–57 right), soft tissue (58–73), lumbar ribs (74–75) and surgical hardware
 162 (76–82).

163 Two anchors make the count decidable, and Fig. 2 shows both: the lowest rib-bearing
 164 vertebra above, **S1** carved from the sacrum below. Each earns its place by what it supports

rather than by being an endpoint — S1 supplies the sacral endplate landmark that sacral slope and pelvic incidence are measured from, and with both brackets present L5 and L6 are distinguishable where a spine-limited view alone would not be.

How S1 is cut. The sacrum’s outer boundary is the source annotation and is never altered; S1 is the cranial part of that same bone, separated by a plane. The plane’s normal is the sacrum’s own cranio-caudal axis, made orthogonal to a left–right axis that is measured rather than assumed — the normal of the reflection plane best mapping the sacrum onto itself, which removes the patient’s rotation within the scanner. That rotation is not negligible: median 4.5°, maximum 16.1°, and 508 of 801 records above 3°. The level of the cut preserves the sub-division volume of the preceding release, so the plane changes the shape and orientation of the S1/S2 boundary and not how much of the sacrum is called S1. Per-record geometry ships with the archive.

Neither anchor is novel: TotalSegmentator¹⁶ carries a separate S1 class and per-level ribs, and VerSe⁵ carries L6. Because this release uses the VerSe identifiers, its L6 *is* VerSe’s, interoperable rather than local to this work. L6 is retained not for novelty but because a scheme without it cannot record a six-lumbar spine at all, and must renumber the column or drop a level to fit.

A rib on a lumbar body receives its own class, and this is the one class here with no published counterpart. A lumbar rib and a stump rib are the same object under two counts, so assigning it a number would commit the label to one reading of an enumeration anomaly. A scheme that numbers every rib 1–12 has nowhere to put a thirteenth: the annotator must either call it rib 12 — asserting the vertebra beneath it is thoracic, the very question at issue — or discard it. TotalSegmentator carries ribs 1–12 per side and no such class, and VerSe’s T13 is a vertebra rather than a rib. 16 released records carry a lumbar rib and 18 carry an L6.

Figure 2 shows why the interval is the primitive and the label is derived. Among the records with four rib-free vertebrae, some reach that count through a lumbar rib and some without one; the two are indistinguishable by the count alone and are distinguished here only because rib presence is itself labeled.

Sixteen records carry a lumbar rib in the released volumes: thirteen bilateral and three unilateral, all three on the right. That laterality is reported because it is what the release contains and not as a finding — sixteen records cannot support a statement about side




FIG. 2 The two anchors and the interval between them, from the released labels, posterior view. Rostral anchor: the lowest rib-bearing vertebra and the rib that makes it one. Caudal anchor: S1. Both are found from the labels rather than assumed, and panels are colored by role rather than identifier, since coloring T12 would assert the count the figure derives. Panels (b) and (c) reach the same interval of four from opposite ends and are indistinguishable by it. All panels are at one millimeter scale, aligned at the sacral base.

197 preference, and none is made.

198 Eighteen records carry an L6. Fourteen are labeled lumbarization, two sacralization,
 199 one semi-sacralization and one normal, which is worth stating because it shows the two
 200 annotation axes disagreeing in the open: a record can carry six lumbar-type vertebrae and
 201 still be read as a sacralization, or as unremarkable, depending on where the reader started
 202 the count.

203 Both counts are derived from the released label volumes, and so are the manifest fields
 204 that report them. That is a correction rather than a design note: through v5 the `has_l6`
 205 flag was true for exactly one record, which contains no L6, and false for all eighteen that
 206 do, while `has_lumbar_rib` was false everywhere and `n_lumbar_labels` read zero in 799 of
 207 802 records. A reader selecting the six-lumbar cases on those fields would have received one
 208 record that is not one, and a reader selecting lumbar ribs an empty set, with nothing to
 209 indicate either. The fields are rebuilt for this release by counting identifiers 25 and 74–75 in
 210 each volume, and `lumbar_rib_side` is added alongside them.

211 II.F. Computational tools

212 *Access.* Every processing step described above is performed by code released under an
213 open-source licence and archived with the dataset; no step depends on software that is
214 not publicly obtainable. The release archive carries the reference loader, the label-scheme
215 definition and the quality-control scripts that generated the tables in this manuscript.

216 *Versions and parameters.* Segmentation of femora, the sacral sub-division and the per-
217 level rib numbering used TotalSegmentator¹⁶ (`total` task, release 2.x) on a single graphics
218 processing unit. The binary rib network is Möller’s¹⁵ released weights, applied through the
219 nnU-Net v2 framework¹⁸. The pelvic completion for records whose pelvis was absent used a
220 five-fold nnU-Net v2 ensemble, applied out-of-fold: each record is completed by the fold that
221 never trained on it, so no record is completed by a model that has seen it. Image handling
222 throughout used NiBabel and SciPy; no image is resampled or reoriented when a label is
223 written, so every label shares its image’s grid and affine exactly.

224 *Generative artificial intelligence.* Large-language-model assistance was used for manuscript
225 preparation — drafting and editing of text, and generation of analysis and figure code —
226 under author review. It was not used to generate, annotate or interpret any imaging data,
227 and no scientific claim in this manuscript originates from it. All numbers reported here are
228 computed by the released code from the released volumes.

229 II.G. Validation

230 Validation is layered, and the ordering matters: measurements computed on a corpus
231 that has not established internal consistency do not look broken, they look finished.

232 *Geometric invariants.* Every case is checked for label geometry agreeing with its CT in
233 shape, affine and voxel spacing; for identifiers within the published scheme; for a non-empty
234 label; and for sided structures falling on the correct sides, with the left–right axis read from
235 the affine rather than assumed. The check exists because a renumbering pass once wrote
236 a reoriented array under a canonical affine and transposed a label away from its image, a
237 failure invisible downstream.

238 The check compares *each sided pair separately*, and both simpler alternatives miss cases
239 it catches. Testing the ribs alone passed all 802 records while four of them carried `left_hip`

TABLE I Release validation. Gates run in order and a failure stops the chain.

Check	Result
Geometric invariants (all cases)	802/802 pass
label shape, affine, spacing match CT	—
no identifier outside the scheme	—
left/right structures correctly sided	—
Ribs present	11,548
not evaluable (expected vertebra outside the FOV)	5,788
no vertebra within the anchor distance	11
evaluable	5,749
matching the expected vertebra	5,747
residual offset	2 (0.035% of evaluable)
misnumbered cases	0
Spinopelvic medians within published range	6/6

on the patient’s right. Pooling all sided structures into one test recovered three of those four and missed the fourth, because hips and femora outweigh the ribs by voxel count, so a large correctly-sided structure masks a smaller transposed one. A gate that passes everything is not evidence of a clean corpus until it has been shown capable of failing.

Rib-vertebra incidence. Each rib is matched to its nearest vertebra and compared against the vertebra its number implies; the chain is given in Table I. The denominator is stated there rather than only the numerator because quoting the two residual offsets against all 11,548 ribs would halve the rate by counting ribs the check never examined — and because more than half the ribs in a screening abdominal CT being unevaluable is itself a fact about the corpus worth reporting. Both offsets are field-of-view truncations on individually characterized cases, one at +1 level and one at −1, and are reported rather than suppressed: a threshold tuned until the count reaches zero yields a cleaner number and a less informative one.

The same check carries an invariant that reads as a null result. Of the 5,749 evaluable ribs, *zero* are assigned to a lumbar vertebra — not because no lumbar ribs exist, since 16 records carry one, but because such a rib takes class 74 or 75 and never enters the numbered

TABLE II Derived measures against a published asymptomatic reference ($n = 260$, standing)¹⁹. This cohort is supine, and the two postural measures depart in opposite directions — sacral slope 4.1° lower, pelvic tilt 4.3° higher. That is one discrepancy, not two: $PI = PT + SS$ holds by construction, so with pelvic incidence at its published value any fall in one appears as an equal rise in the other.

Measure	This dataset (supine)	Reference (standing)
Pelvic incidence	54.7°	54.7 ± 10.6
Pelvic tilt	17.3°	13 ± 6
Sacral slope	36.9°	41.0 ± 8.4
Lumbar lordosis (supine)	52.7°	43 ± 11.2 (standing)
PI – LL mismatch	2.6°	centered near 0

set the test examines. A numbered rib landing on a lumbar vertebra would be a counting error of exactly the kind this dataset exists to expose.

Agreement with published values. Derived spinopelvic measures are compared against Vialle *et al.*¹⁹ (Table II). Pelvic incidence is the strongest of these checks because it is a morphological property of the pelvis rather than a posture, and so is directly comparable to a standing reference cohort.

II.H. Castellvi typing

Every record whose vertebral count departs from five carries a Castellvi grade²⁰ in the released manifest: 33 cases, typed I–IV with the *a/b* unilateral–bilateral qualifier. This is released as an annotation layer because the grade and the count are *different axes* and the dataset is built to show that they are: grade IIIb occurs here at rib-free counts of four, five and six alike, and seven of the 33 graded cases carry a perfectly normal count of five. A corpus that recorded only the count would describe those seven as unremarkable.

The grades were assigned by two radiology resident physicians. The distinction the classification turns on — bony fusion against articulation without it² — is the one carrying clinical weight.

272 II.I. Surgical hardware, and why a threshold is not enough

273 Metal is trivial to detect and hard to interpret, and the difference matters here for one rea-
274 son: **an iatrogenic fusion is indistinguishable from a congenital one to a distance**
275 **measurement**. A cage-bridged interspace reads as “no gap” exactly as a congenitally fused
276 transitional vertebra does, so an unrecorded instrumented case is a false positive on the cen-
277 tral measure — and identifiers 76–79 were declared in every prior version and populated in
278 none.

279 *II.I.0.a. Detection over-calls by a factor of seven.* A threshold at 1800 HU inside a shell
280 around the labeled skeleton flags 84 of the 802 records. Raising it to 2500 HU — the lower of
281 the two values validated in the metal-segmentation literature — leaves 52 with a component
282 above a 40-voxel floor. One of the two radiology resident physicians read all 52 and confirmed
283 11, rejecting 41 (79%) as contrast, calcification and reconstruction artefact. Instrumentation
284 prevalence in this cohort is 1.4%, not the 10.5% an unreviewed threshold reports.

285 *II.I.0.b. Saturation does not separate implant from artefact.* Both groups reach the scan-
286 ner ceiling and 4 rejected proposals *exceed* it, peaking at 11,798 HU — the signature of
287 reconstruction overshoot rather than of denser metal. What separates them is position and
288 size together: a real surgical site *and* volume, every confirmed implant at 2,586 mm³ or
289 above and every rejection at 1,768 mm³ or below.

290 *II.I.0.c. The subtype block had to be extended.* Identifiers 76–79 name spinal instrumen-
291 tation — generic, cage, screw-and-rod, plate — and cannot describe what this cohort holds,
292 so **80 arthroplasty**, **81 sacroiliac screw** and **82 osteosynthesis** were added. The dis-
293 tinction is clinical, not cosmetic: osteosynthesis holds parts of the same bone together where
294 arthroplasty replaces a joint, and without it the shape rules call a femoral stem a rod, since
295 a stem is long and thin and that is all “linear” tests. The confirmed cases comprise 8 arthro-
296 plasties, one osteosynthesis, one sacroiliac screw fixation and one pair of threaded cylindrical
297 interbody cages.

298 *II.I.0.d. Metal outranks bone, and that is a subtraction.* In every confirmed case the
299 implant was already inside a bone label, a segmenter handed a bright object against a
300 cortical surface having taken it for bone. Naming the hardware therefore reclaims voxels
301 rather than adding them: 1,538,852 voxels across the eleven records. The consequence
302 reaches past bookkeeping. Pelvic incidence and pelvic tilt are measured from the femoral

303 head center, and in 8 of these cases that head is an implant — any spinopelvic parameter
304 previously computed from them was measured on metal.

305 III. DATA FORMAT AND USAGE NOTES

306 Each record is a gzipped NIfTI image/label pair sharing an identical 4×4 affine, canon-
307 icalized to PIR, requiring no resampling. Masks are NIfTI rather than DICOM: the policy
308 prefers DICOM where applicable, and for volumetric segmentation masks NIfTI is what the
309 downstream tooling expects.

310 Splits are patient-grouped and LSTV-stratified five-fold, frozen and shipped with the
311 data as `splits_5fold.json`. The five validation folds partition all 802 records, so the
312 release provides cross-validation and *not* a held-out test set. Users reporting a single held-
313 out number should carve one out and state which records it holds, because the folds shipped
314 here will not do that job.

315 Stratification is on the transitional subtype rather than on a binary LSTV flag, and it
316 enforces a minimum of each rare subtype in every validation fold: each fold’s validation
317 set carries three sacralization and three lumbarization records, so no fold is evaluated on a
318 cohort missing the class the dataset exists for. With 15 lumbarization and 15 sacralization
319 records across 802, that is the practical limit of what stratification can guarantee, and it is
320 why a fold-level metric on the rare classes carries an error bar far wider than its own decimal
321 places.

322 Table III lists the key data types and metadata against the fraction of records for which
323 each is populated. Two entries are worth reading carefully. The Castellvi grade is present
324 for 4.1% of records because it is defined only for the transitional ones; among those it is
325 complete. A pelvic series identifier is present for 88.9% because the crosswalk records it
326 only where a pelvic annotation was actually placed, and users joining on it should expect
327 the shortfall rather than treat it as loss.

328 The archive of record is Zenodo (10.5281/zenodo.22139642). Code and documentation
329 are on GitHub at the tagged release commit; any model-hub copy is a convenience mirror,
330 not the archive.

TABLE III Key data types and metadata, and the fraction of the 802 released records for which each is populated. Counts are obtained by reading the released manifest, not by assertion. A field below 100% is not an error: age and sex are absent where the source collection did not publish them, a pelvic series identifier exists only where a pelvic annotation was placed, and a Castellvi grade exists only for the transitional records that were graded.

Field	Records	Available
Identifier, image path, label path, configuration, match type	802	100.0%
Patient position, tube potential, section thickness, kernel	802	100.0%
Transitional label and class	802	100.0%
Spine and pelvic annotation origin	802	100.0%
Instrumentation and partial-annotation flags	802	100.0%
Image/label alignment check	802	100.0%
Spine series identifier; bone-fraction check	782	97.5%
Scanner manufacturer and model	768	95.8%
Sex	749	93.4%
Age	709	88.4%
Pelvic series identifier	713	88.9%
Castellvi grade	33	4.1%

IV. DESCRIPTIVE ANALYSIS

The cohort is a colorectal cancer screening population of 802 records. 738 are recorded female (393) or male (345); 11 carry DICOM’s *other* value, a recorded value and not a missing one, and 53 carry no sex field. Age is present for 709 (median 59, range 50–89). All 802 are counted in totals, and records outside the female/male dichotomy are excluded from sex-stratified measures — a limitation of those measures rather than of the cohort, stated because folding eleven people into “unrecorded” would misdescribe the source. The lower age bound is the screening guideline rather than a selection applied here, and it is why the measures below sit closer to a reference range than a surgical series would.

Count-free measures (Fig. 3). Five rib-free vertebrae above the sacrum is typical; four and six are where transitional anatomy sits, and which of the two is present does not by




FIG. 3 Count-free measures, valid without knowing which vertebra is which. (a) Vertebrae between the lowest rib and the sacrum: an interval, not a level assignment. (b) Lowest rib length as a fraction of the rib above, showing two populations near 0.33 and 0.69 rather than one distribution with a tail. (c) Lowest lumbar transverse span against its gap to the sacral ala; cases carrying a source transitional label are marked.




FIG. 4 Spinopelvic measures against published reference bands (mean $\pm$ SD) from an asymptomatic standing cohort.¹⁹ Sacral slope is expected to sit below a standing reference because it is postural and this cohort is supine; pelvic incidence is not postural, which is why it can be compared without apology and why its agreement is the strongest check here.

342 itself determine what to call it. The lowest-rib length ratio is bimodal.

343 *Level gradients and spinopelvic measures* (Fig. 4). Vertebral dimensions increase caudally
 344 in step with the load each level carries. Pelvic incidence does not vary with age, while pelvic
 345 tilt increases — the descriptive signature of a pelvis retroverting as a spine ages.

346 *Opportunistic measures.* Every scan carries a vertebral bone density measurement at no
 347 additional dose.²¹ L1 trabecular attenuation is reported at the published standard site.

348 V. POTENTIAL APPLICATIONS AND LIMITATIONS

349 The principal use is the *spinopelvic interface*: spine and pelvis carry radiologist-sourced
350 annotation on one coordinate frame in 802 records, which makes sacral morphology measur-
351 able against the lumbar column rather than in isolation — the sacral endplate and its slope,
352 the alar corridors constraining S2-alar-iliac and iliac screw trajectories, and the L5–S1 geom-
353 etry a lumbosacral fusion must cross. Pelvic incidence is derived from the segmented sacrum
354 and femoral heads rather than from landmarks placed on a radiograph, so it is reproducible
355 from the labels themselves.

356 The transitional layer serves that use rather than competing with it: no trajectory can be
357 planned across a junction whose segments cannot be named. Level-numbering benchmarks,
358 variant studies and opportunistic screening follow from the same annotation.

359 V.A. Three uses the released measurements support

360 *Naming a level from its own geometry.* Wrong-level spine surgery runs at roughly one
361 in 3,110 spinal procedures, and the cause most often cited is the transitional anatomy this
362 cohort was assembled around.²² The failure is structural rather than careless: the count
363 is ambiguous exactly where the variant sits. Per-level body height, canal width, endplate
364 width, transverse-process span and wedge ratio are released for T11–L5 across all 802 records,
365 and because the labels were assigned against the twelfth rib rather than by enumeration, a
366 model trained on them is not learning the convention that fails.

367 *Reference morphometry with its spread.* The values a surgeon consults for vertebral
368 body height, canal dimensions and pedicle width come from cadaveric series of a few dozen
369 specimens, plotted as single values with no indication of spread.¹⁰ They cannot tell a reader
370 whether the patient in front of them is unusual, because they never show what usual looks
371 like. The same measures are given here from 802 records with the distribution attached
372 (Fig. 5), reproducing two properties of the classical figures — body height crosses over,
373 dorsal exceeding ventral at T11–T12 and reversing by L4–L5, and pedicle width widens
374 steeply into the lower lumbar spine — and adding the width of each distribution, which is
375 the part needed to judge an individual. A larger n sharpens a distribution without changing
376 the population it came from, and this one is a screening cohort aged 50 and over.




FIG. 5 Lumbar and lower thoracic morphometry by level, in the format the standard reference figures use — level on the vertical axis, measurement on the horizontal. Marker, median; bar, interquartile range; line, 5th–95th percentile; the number at the right of a row is its n . (a) Vertebral body height, ventral against dorsal. (b) Superior endplate width. (c) Canal width against depth. (d) Transverse pedicle width. (e) Disc height, drawn at the interspace it occupies. (f) Trabecular attenuation at the published opportunistic-screening site. Coverage is not uniform and is drawn rather than assumed: canal depth and pedicle width rest on 330–656 records because they need an axial extent many field-limited series do not carry, while the remaining measures rest on 720–784. No comparison between levels is tested and none is implied; the intervals are dispersion.

377 *Patient-specific models.* Planning is moving toward patient-specific biomechanical models
 378 rather than population norms.²³ Such a model needs spine, pelvis and femoral heads in
 379 one frame, which is this release’s construction; in 351 patients the annotation sits on two
 380 acquisitions in different positions, giving a within-patient change in alignment against which
 381 a predicted postural response can be checked.

382 Its limitations are part of it, and they are stated at the level of detail a reader would
 383 need to catch them independently.

384 *Coverage.* Thoracic ground truth is field-of-view limited and does not extend to T1.
385 Postural angles are supine; pelvic incidence is not postural and needs no such caveat. The
386 cohort is a colorectal screening population aged 50 and over, so its distributions should not
387 be read as representative of a surgical one.

388 *Declared but empty, and confirmed so.* The soft-tissue identifiers (58–73) and the partial-
389 annotation sentinel (255) are declared in the scheme and occur in no record — established by
390 counting identifiers across all 802 released volumes rather than asserted. Both are retained
391 so a later release can use them without renumbering, and their absence should be read as
392 absence of annotation rather than absence of the structure. The hardware identifiers were
393 in this position until v6 and are no longer: 76–82 are populated in the eleven instrumented
394 records described in Sec. [II.I](#).

395 *Strength of the labels varies by structure, and the release does not average over that.*
396 Vertebral labels derive from radiologist-supervised source annotations, corrected where those
397 sources were wrong. Pelvic labels on records that lacked one are pseudolabeled. The rib
398 layer is a pseudolabel whose human review was triaged by an automated rule rather than
399 exhaustive, so its guarantee is the absence of the failure modes that rule detects. Transitional
400 labels derive from two independent sources and are not uniformly adjudicated — which is
401 precisely why the measures reported here are count-free — and the Castellvi grades were
402 assigned by two radiology resident physicians. The sacral sub-division inherits its level from
403 an automatic method, and in 100 records that level places S1 above half the sacrum or
404 below fifteen percent of it, which no first sacral segment is; those records are flagged in the
405 released quality-control table and should be filtered before any measurement depending on
406 the sacral endplate.

407 *Structures are not universally present.* The same census shows the release is not uniformly
408 complete. Sacrum, hips and femora are present in all 802 records. One record carries no
409 S1, one lacks T12, and nine lack an L5 identifier — in every case because the structure
410 lies outside the field of view or, for L5, because the lowest lumbar segment is labeled L6 or
411 incorporated into the sacrum. Users filtering on the presence of the caudal anchor should
412 filter explicitly rather than assume it.

413 VI. DISCUSSION

414 *Comparison with related material.* Against the collections in Fig. 1, the contribution
415 is the crosswalk placing spine and pelvis on the same series, together with classes for the
416 anatomy an enumeration anomaly actually produces: where another scheme must record
417 a sixth lumbar vertebra or a lumbar-borne rib as something else, this release records it as
418 itself.

419 VII. CONCLUSION

420 CTSpinoPelvic1K places spine, pelvis, per-level ribs and femora on one coordinate frame
421 in 802 records, and gives the anatomy that makes lumbar numbering ambiguous a class of
422 its own rather than recording it as something it is not. Because both counting anchors are
423 explicit, a level can be identified from a field of view that cannot support the conventional
424 count downward from C2 — which is every record here. The release is archived under a
425 persistent identifier, ships what is needed to verify it, and states per structure how strong
426 each label is.

427 DATA AVAILABILITY

428 The release described here is v7, permanently archived at [https://doi.org/10.5281/](https://doi.org/10.5281/zenodo.22242745)
429 [zenodo.22242745](https://doi.org/10.5281/zenodo.22242745). The dataset as a whole, across versions, is cited as [https://doi.org/](https://doi.org/10.5281/zenodo.22139642)
430 [10.5281/zenodo.22139642](https://doi.org/10.5281/zenodo.22139642), which always resolves to the current version. The archive car-
431 ries the reference loader, the label-scheme definition and the quality-control tables, and the
432 analysis code is distributed with it under an open-source licence. The repository URL is
433 omitted here because it identifies the authors; it is given on the title page and will be restored
434 to this section for the camera-ready version.

435 CONFLICT OF INTEREST

436 The authors have no relevant conflicts of interest to disclose.

ETHICS

This work uses publicly available, de-identified imaging and annotations derived from it. The authors' institutional review board determined that it does not constitute human participant research under the Common Rule at 45 CFR 46, and that review and oversight are therefore not required.

REFERENCES

- ¹J. Lian, N. Levine, and W. Cho, A review of lumbosacral transitional vertebrae and associated vertebral numeration, *Eur. Spine J.* **27**, 995–1004 (2018), doi:10.1007/s00586-018-5554-8.
- ²G. P. Konin and D. M. Walz, Lumbosacral transitional vertebrae: classification, imaging findings, and clinical relevance, *AJNR Am. J. Neuroradiol.* **31**, 1778–1786 (2010), doi:10.3174/ajnr.A2036.
- ³Y. Deng et al., CTSpine1K: a large-scale dataset for spinal vertebrae segmentation in computed tomography, *Mach. Learn. Biomed. Imaging* **3**, 824–832 (2025), doi:10.59275/j.melba.2025-gf84.
- ⁴P. Liu et al., Deep learning to segment pelvic bones: large-scale CT datasets and baseline models, *Int. J. Comput. Assist. Radiol. Surg.* **16**, 749–756 (2021), doi:10.1007/s11548-021-02363-8.
- ⁵A. Sekuboyina et al., VerSe: a vertebrae labelling and segmentation benchmark for multi-detector CT images, *Med. Image Anal.* **73**, 102166 (2021), doi:10.1016/j.media.2021.102166.
- ⁶H. Liebl et al., A computed tomography vertebral segmentation dataset with anatomical variations and multi-vendor scanner data, *Sci. Data* **8**, 284 (2021), doi:10.1038/s41597-021-01060-0.
- ⁷M. R. Zindrick et al., Analysis of the morphometric characteristics of the thoracic and lumbar pedicles, *Spine (Phila. Pa. 1976)* **12**, 160–166 (1987), doi:10.1097/00007632-198703000-00012.
- ⁸M. M. Panjabi et al., Thoracic human vertebrae: quantitative three-dimensional anatomy, *Spine (Phila. Pa. 1976)* **16**, 888–901 (1991), doi:10.1097/00007632-199108000-00006.

⁹M. M. Panjabi et al., Human lumbar vertebrae: quantitative three-dimensional anatomy, Spine (Phila. Pa. 1976) **17**, 299–306 (1992), doi:10.1097/00007632-199203000-00010.

¹⁰E. C. Benzel, *Biomechanics of Spine Stabilization*, Thieme, New York, 3rd edition, 2015.

¹¹R. J. Hughes and A. Saifuddin, Numbering of lumbosacral transitional vertebrae on MRI: role of the iliolumbar ligaments, AJR Am. J. Roentgenol. **187**, W59–W65 (2006), doi:10.2214/AJR.05.0415.

¹²L. Jin et al., RibSeg v2: a large-scale benchmark for rib labeling and anatomical centerline extraction, IEEE Trans. Med. Imaging **43**, 570–581 (2024), doi:10.1109/TMI.2023.3313627.

¹³K. Smith et al., Data from CT COLONOGRAPHY, The Cancer Imaging Archive, 2015, doi:10.7937/K9/TCIA.2015.NWTESAY1.

¹⁴K. Clark et al., The Cancer Imaging Archive (TCIA): maintaining and operating a public information repository, J. Digit. Imaging **26**, 1045–1057 (2013), doi:10.1007/s10278-013-9622-7.

¹⁵H. Möller et al., Automated thoracolumbar stump rib detection and analysis in a large CT cohort, 2025, arXiv:2505.05004. doi:10.48550/arXiv.2505.05004.

¹⁶J. Wasserthal et al., TotalSegmentator: robust segmentation of 104 anatomic structures in CT images, Radiol. Artif. Intell. **5**, e230024 (2023), doi:10.1148/ryai.230024.

¹⁷A. Schehr, J. Kim, and G. Schwing, OpenSpineConsortium: an open-source framework for medical student engagement in computational spine imaging research, Cureus (2026), Published 14 July 2026. doi:10.7759/cureus.112661.

¹⁸F. Isensee, P. F. Jaeger, S. A. A. Kohl, J. Petersen, and K. H. Maier-Hein, nnU-Net: a self-configuring method for deep learning-based biomedical image segmentation, Nat. Methods **18**, 203–211 (2021), doi:10.1038/s41592-020-01008-z.

¹⁹R. Vialle, N. Levassor, L. Rillardon, A. Templier, W. Skalli, and P. Guigui, Radiographic analysis of the sagittal alignment and balance of the spine in asymptomatic subjects, J. Bone Joint Surg. Am. **87**, 260–267 (2005), doi:10.2106/JBJS.D.02043.

²⁰A. E. Castellvi, L. A. Goldstein, and D. P. K. Chan, Lumbosacral transitional vertebrae and their relationship with lumbar extradural defects, Spine (Phila. Pa. 1976) **9**, 493–495 (1984), doi:10.1097/00007632-198407000-00014.

²¹P. J. Pickhardt, B. D. Pooler, T. Lauder, A. M. del Rio, R. J. Bruce, and N. Binkley, Opportunistic screening for osteoporosis using abdominal computed tomography scans

498 obtained for other indications, *Ann. Intern. Med.* **158**, 588–595 (2013), doi:10.7326/0003-
499 4819-158-8-201304160-00003.

500 ²²N. E. Epstein, A perspective on wrong level, wrong side, and wrong site spine surgery,
501 *Surg. Neurol. Int.* **12**, 286 (2021), doi:10.25259/SNI_402_2021.

502 ²³E. Beaulieu et al., A review of finite element analysis in spine surgery decision-making, *J.*
503 *Clin. Med.* **15**, 2584 (2026), doi:10.3390/jcm15072584.

Supplementary Material

CTSpinoPelvic1K: spine, pelvis, ribs and femora in one coordinate frame, annotated for lumbosacral transitional anatomy




Figure S 1: How the dataset was built. Three public collections are combined patient by patient. Neither annotation source published the CT series its labels were drawn on, so each label set is matched to the series whose bone agrees with it. In 351 of those patients the two collections had drawn their outlines on *different scans of the same person*. Numbers marked *n* are scans, or corrections, at that step.