

CTSpinoPelvic1K: spine, pelvis, ribs and femora in one coordinate frame, annotated for lumbosacral transitional anatomy

Article type: Dataset and Software Article

Authors

1. Gregory Schwing, MD, PhD ¹
2. Ashley Schehr, BS ²
3. Annika Tekumulla, BS ²
4. Margret Khoushi, BS ²
5. Ryan Christian, BS ²
6. Dane Hubers, BS ²
7. Faris Mahjoub, BS ²
8. Hassan Saad, BS ²
9. Mia Sooch, BA ²
10. Sathya Siddapureddy, BS ²
11. Michael McLellan, BS ²
12. Jerick Kim, BS ²
13. Miraziz Ismoilov, MD ³
14. Nizar Alnabahneh, MD ³

¹Department of Surgery, Detroit Medical Center and Wayne State University, Detroit, Michigan, United States of America

²School of Medicine, Wayne State University, Detroit, Michigan, United States of America

³Department of Radiology, Detroit Medical Center and Wayne State University, Detroit, Michigan, United States of America

Corresponding author

Gregory Schwing

Department of Surgery, Detroit Medical Center and Wayne State University

Detroit, Michigan, United States of America

Email: [corresponding author email]

Author contribution statement

G.S. designed the dataset and the label scheme, implemented the processing pipeline and the quality-control procedures, and drafted the manuscript. A.S., A.T., M.K., R.C., D.H., F.M., H.S., M.S., S.S., M.M. and J.K. performed annotation correction and quality-control review of the released segmentations under supervision. The radiological reads — the Castellvi grading and the adjudication of candidate instrumentation — were performed by one radiology resident physician. [BEFORE SUBMISSION: confirm each contribution individually with the author concerned and expand to CRediT roles.] All authors read and approved the final version of the manuscript.

Conflict of interest statement

The authors have no relevant conflicts of interest to disclose.

Acknowledgements

The imaging analysed here was made available by The Cancer Imaging Archive and by the investigators who published the source annotations. [Funding sources to be listed.]

Data and code availability (unblinded)

Dataset, archived before submission: <https://doi.org/10.5281/zenodo.22139642> (concept identifier, resolves to the current version); the version described in the manuscript is v7, <https://doi.org/10.5281/zenodo.22242745>.

Code: <https://github.com/Gregory-Schwing-MD-PhD/CTSpinoPelvic1K>, tagged at the release commit.

Mirror: <https://huggingface.co/datasets/OpenSpineConsortium/CTSpinoPelvic1K> — a convenience mirror, not the archive of record.

Ethics

This work uses publicly available, de-identified imaging and annotations derived from it. A determination of not-human-participant research has been submitted to the Wayne State University Institutional Review Board.

**CTSpinoPelvic1K: spine, pelvis, ribs and femora in one coordinate frame, annotated
for lumbosacral transitional anatomy**

[Author names removed for double-anonymised review]

[Affiliations removed for double-anonymised review]

(Dated: 2 September 2026)

Purpose: Lumbar levels cannot be numbered reliably when a lumbosacral transitional vertebra (LSTV) is present. Sacralisation and lumbarisation are different developmental events and leave different traces, but those traces are not decisive in every case, and the count that would arbitrate is taken downward from C2 — the one landmark present in every spine and never duplicated¹. A field of view that begins in the thorax contains no such landmark, so the name rests on morphology alone exactly where morphology is least certain. Public collections either omit the pelvis (CTSpine1K², VerSe³) or omit the classes a transitional vertebra requires: TotalSegmentator⁴, the largest whole-body collection, provides no sixth lumbar vertebra and no class for a rib arising from a lumbar level, and VerSe excludes vertebrae fused to the sacrum by design⁵. CTSpinoPelvic1K provides 802 CT records carrying radiologist-sourced spine *and* pelvic annotation on one coordinate frame, with per-level ribs and femora. Four classes exist so that a variant need not be forced into an ordinary level: a sixth lumbar vertebra, a thirteenth thoracic vertebra, the first sacral segment held separate from the sacrum, and left and right lumbar ribs. The primary use is the lumbosacral interface — sacral morphology measured against the lumbar spine, for pelvic fixation and lumbar surgical planning — and the transitional labelling is what makes that interface trustworthy, since a trajectory cannot be planned across a junction whose segments cannot be named. Because every record follows one screening protocol, between-scan variation reflects anatomy rather than acquisition, which suits the collection to learning morphology directly.

Acquisition and Validation Methods: Records were built by patient-level crosswalk of three public collections — TCIA CT COLONOGRAPHY imaging, CTSpine1K² vertebral labels, and CTPelvic1K⁶ pelvic labels — with labels placed on the acquisition carrying the highest bone coverage and unified under a VerSe-native scheme. Validation is threefold: per-case geometric invariants (802/802 pass); rib-vertebra incidence across 5,749 evaluable ribs (2 residual offsets, 0.035%); and agreement of derived spinopelvic measures with published reference values.

Data Format and Usage Notes: Gzipped NIfTI image/label pairs sharing an identical affine, canonicalised to PIR, with frozen patient-grouped LSTV-stratified five-fold cross-validation splits and a reference loader. Archived at 10.5281/zenodo.22139642.

Potential Applications: Segmentation and level-numbering benchmarks, anatomical variant studies, surgical planning, and opportunistic screening research. *Limitations:* thoracic ground truth is field-of-view limited; postural angles are supine; the soft-tissue identifiers are declared but unpopulated; the splits are cross-validation folds and include no held-out test set; the rib layer is a reviewed pseudolabel whose review was triaged by an automated rule rather than exhaustive; and the Castellvi grades were assigned by a radiology resident physician.

6 I. INTRODUCTION

7 Labelling a lumbar vertebra requires counting, and counting requires two fixed ends.
8 The spine commonly carries seven cervical, twelve thoracic and five lumbar vertebrae, but
9 *enumeration anomalies* — eleven or thirteen thoracic, four or six lumbar — are common, and
10 the transitional phenotypes that produce them sit at the two ends of the lumbar column: the
11 lumbosacral transitional vertebra (LSTV) below and the thoracolumbar transitional vertebra
12 (TLTV) above. In the presence of an LSTV the lower end is not where it is assumed to be.
13 The two phenotypes are distinct developmental events — a lumbar segment assimilating
14 to the sacrum, or a sacral segment separating from it — and they leave different traces:
15 a sacralised L5 tends toward an anteriorly wedged body, a reduced disc beneath it and
16 hypoplastic or absent facet joints, while a lumbarised S1 tends toward a squared body in the
17 sagittal plane, well formed lumbar-type facets and a full-height disc¹. What counting settles
18 is the *name*. The difficulty is that these traces are not decisive in every case, so the name and
19 the morphology can disagree, and the count that would arbitrate is taken downward from
20 C2 — unavailable in any spine-limited acquisition. LSTV is common, reported between 4%
21 and 30% depending on definition, and wrong-level surgery is its most serious consequence.

22 Existing public collections do not make this tractable, and not because they are small
23 (Fig. 1): the spine collections omit the pelvis, the pelvic ones carry no vertebral count, and
24 the whole-body schemes have neither an L6 class nor a class for a rib borne by a lumbar
25 vertebra, so a six-lumbar spine cannot be written down in them at any segmentation quality.
26 VerSe^{3,5} is the instructive case, and Fig. 1(b) gives it: a collection can enrich for anatomical
27 variants deliberately, grade them, and still be unable to represent the ones it grades.

28 CTSpinoPelvic1K places the spine, pelvis, ribs and femora in one coordinate frame and
29 gives a class to each structure an enumeration anomaly produces: a sixth lumbar vertebra,
30 and a rib borne by a lumbar vertebra. A scheme without those classes must record such
31 anatomy as something it is not.

32 *Where the material is.* Everything described here is v7 of a single archived dataset,
33 deposited at <https://doi.org/10.5281/zenodo.22242745>. Citations should use [https:](https://doi.org/10.5281/zenodo.22139642)
34 [//doi.org/10.5281/zenodo.22139642](https://doi.org/10.5281/zenodo.22139642), the permanent identifier for the dataset across all
35 of its versions, which resolves to whichever version is current; the version identifier above
36 pins the exact release measured in this manuscript. Documentation, the label scheme, the

reference loader and the per-record quality-control tables are distributed with the archive
and are described in Sec. [III](#).

I.A. Vertebra labelling when the scan cannot settle the count

The limitation, stated plainly. No scan in this corpus contains C2 or the whole cervical spine, so the conventional count cannot be performed. That is a property of abdominal imaging rather than of the annotation, and it means a variant at the thoracolumbar junction cannot be *resolved by counting* here at all: a thirteenth thoracic vertebra, a lumbar rib, and a *stump rib* — a hypoplastic twelfth rib, one of the cardinal indicators of a thoracolumbar transitional vertebra — produce overlapping appearances.

What the release does instead. No vertebral label is asserted where it cannot be supported. Every quantity is expressed against the two landmarks present in essentially every abdominal scan — the sacrum below and the lowest rib-bearing vertebra above — so the description is positional rather than nominal: a vertebra by how many rib-free vertebrae separate it from the sacrum, a lowest rib by its length as a fraction of the rib above it, a transverse process by its span and its distance from the ala. None of these changes according to whether a reader would have said L5 or S1, so cases that disagree about the name remain comparable. Where the junction is in view the count is determined and the conventional name follows, and the two descriptions agree.

And the morphology may be local rather than global. Counting is not the only route to an identity. A thoracic vertebra differs in shape from a lumbar one — in the span of its transverse processes relative to its body, in the proportions of its endplates and canal — and if that difference holds at the boundary itself, the phenotype is decidable from the vertebra in front of the reader without any global count. On this corpus it does hold: separating T12 from L1 on shape alone, with each patient’s own size divided out and cross-validation grouped by case, gives an area under the curve of 0.990 and 97.3% accuracy over 1485 vertebrae. Size is deliberately removed, because a classifier given raw millimetres separates thoracic from lumbar immediately and has learned nothing that helps at the junction, where the two are nearly the same size.

That is reported as a property of the released measurements rather than as a method: the information needed to settle these variants is present locally.




FIG. 1 Public CT collections at the lumbosacral junction. (a) Size, stated first because it is the claim this dataset does *not* make: three of the five comparators carry more scans, and RibSeg v2⁷ labels more ribs. (b) What each published scheme can *represent*, a property of its class list rather than of its segmentation quality. VerSe is graded-but-excluding: it applies the Castellvi classification and uses it to decide what to omit, leaving vertebrae partially fused to the sacrum unsegmented⁵ — 25 of the 33 graded records here.

II. ACQUISITION AND VALIDATION METHODS

II.A. Sources, and why a crosswalk was necessary

Both CTSpine1K² and CTPelvic1K⁶ draw part of their imaging from the TCIA CT colonography collection^{8,9}, and the vertebral and pelvic annotations in both were produced under board-certified radiologist supervision. Neither released the mapping from an annotation to the specific CT series it was drawn on.

That mapping is not a bookkeeping detail. Every patient in CT COLONOGRAPHY was scanned *twice*, prone and supine, often with more than one kernel, so one patient identifier resolves to several volumes differing in posture, position, field of view and noise — and an annotation carries a patient identifier and nothing finer.

Placing it on the wrong volume of the same patient is not merely a displacement. The spine is mildly deformable, and turning a patient from supine to prone alters lumbar lordosis and the alignment of each segment against the next, so a vertebra moves by a different amount, and in a different direction, from the one above it. Outlines drawn on one series are the wrong *shape* for the other, and no rigid realignment recovers them. In isolation such

82 a mask looks like competent work; it reveals itself only when overlaid on the image.

83 The crosswalk (Fig. S1) therefore resolves each annotation first to a patient and then,
84 within that patient’s volumes only, to the series whose bone *agrees* with it: candidates are
85 thresholded at bone attenuation and scored by overlap with the annotation mask. Bone
86 rather than image similarity, because bone is what the annotation describes — two acqui-
87 sitions of one abdomen resemble each other everywhere except in where the skeleton sits.

88 II.B. Fused and split records

89 The two annotation sets were placed independently, so they do not always land on the
90 same series. Of 802 records, 342 carry spine and pelvic annotations on the same volume
91 (*fused*), 351 are *separate*, 89 are *spine-only* with no pelvic annotation for that patient, and
92 20 are *pelvic-native*. The split is the crosswalk’s main finding rather than bookkeeping: in
93 those separate records — nearly half of all patients — the two public collections annotated
94 *different volumes of the same person*, and neither recorded which. A pelvic annotation
95 placed on the other acquisition of a prone/supine pair cannot be combined with the spine
96 annotation without resampling across a change of posture. Those records are retained rather
97 than discarded, and held out of the postural analyses, because two acquisitions of one patient
98 differing in spinal alignment are a resource for a mobility study rather than a defect.

99 II.C. Densification, and why it is asymmetric

100 A corpus in which half the records lack a pelvis is not usable as a spinopelvic dataset,
101 so the missing structures were pseudolabelled — asymmetrically (Fig. S1), because the two
102 directions are not equivalent.

103 Pelvis onto spine-only records is the safe direction. The sacrum and innominate bones
104 are single objects with no enumeration: no count to get wrong, no naming convention to
105 choose, no analogue of a transitional vertebra. Quality is reported as held-out performance
106 on the *pelvic-native* records, withheld from training for exactly this purpose.

107 The reverse direction carries the whole problem this dataset exists to address, since a
108 pseudolabelled spine must commit to a vertebral count and on a transitional vertebra will
109 commit to the commoner reading. Those 20 records were therefore corrected by hand rather

than densified automatically — tractable at twenty cases, which is the practical reason the asymmetry runs the way it does.

II.D. Ribs

No public rib annotation covers this cohort, and the two available automatic sources fail in complementary ways.

Möller’s binary rib network¹⁰ is a detection model of high reported accuracy, but it segments ribs that lie *entirely within* the field of view; a rib clipped by the edge of an abdominal acquisition is liable to be missed altogether. Since this is a colonography cohort, the lower ribs are routinely clipped, and those are precisely the ribs the count depends on.

TotalSegmentator⁴ does segment clipped ribs and supplies the per-rib numbering that a binary network cannot. Its characteristic failure is at the other end of the bone: the most medial portion of the rib — roughly the last tenth to sixth approaching the costovertebral joint — is frequently absent. A rib truncated at its medial end never reaches the vertebra it articulates with, so the incidence test that assigns a rib to its level has nothing to test.

The two were therefore unioned (Fig. S1): TotalSegmentator supplies numbering and the medial-clipped cases, Möller supplies detection where its output is more complete, and the union carries TotalSegmentator’s per-level identities.

The result is a pseudolabel whose review was triaged rather than exhaustive, which is worth stating plainly because a layer described only as “human-corrected” invites the reader to assume every bone was looked at. A label-based rule flagged every record in which one rib bone carried more than one identifier, or in which a rib failed to reach the vertebra it articulates with — a single bone with two numbers being a numbering error by construction, whatever the anatomy. The rule selected 152 records; 148 were reviewed and corrected, and 4 were not and are identified in the release. Reviewers were medical students working through OpenSpineConsortium¹¹, trained against a written labelling protocol, with every correction recorded against the annotator who made it. Records passing the rule were not individually inspected, so the layer’s guarantee is that it is free of the failure modes the rule detects, and no stronger. The residual rib–vertebra offsets reported in Sec. II.G are the independent check on what the rule missed.

139 II.E. Label scheme and counting anchors

140 The scheme is VerSe-native: vertebrae keep their VerSe identifiers (C1–C7 = 1–7, T1–T12
141 = 8–19, L1–L6 = 20–25, sacrum 26), and every non-VerSe structure takes a fixed identifier
142 above that range — S1 carved from the sacrum (29), hips (30–31), femora (32–33), ribs
143 (34–45 left, 46–57 right), soft tissue (58–73), lumbar ribs (74–75) and surgical hardware
144 (76–82).

145 Two anchors make the count decidable, and Fig. 2 shows both: the lowest rib-bearing
146 vertebra above, **S1** carved from the sacrum below. Each earns its place by what it supports
147 rather than by being an endpoint — S1 supplies the sacral endplate landmark that sacral
148 slope and pelvic incidence are measured from, and with both brackets present L5 and L6
149 are distinguishable where a spine-limited view alone would not be.

150 *How S1 is cut.* The sacrum’s outer boundary is the source annotation and is never al-
151 tered; S1 is the cranial part of that same bone, separated by a plane. The plane’s normal
152 is the sacrum’s own cranio-caudal axis, made orthogonal to a left–right axis that is mea-
153 sured rather than assumed — the normal of the reflection plane best mapping the sacrum
154 onto itself, which removes the patient’s rotation within the scanner. That rotation is not
155 negligible: median 4.5°, maximum 16.1°, and 508 of 801 records above 3°. The level of the
156 cut preserves the sub-division volume of the preceding release, so the plane changes the
157 shape and orientation of the S1/S2 boundary and not how much of the sacrum is called S1.
158 Per-record geometry ships with the archive.

159 Neither anchor is novel: TotalSegmentator⁴ carries a separate S1 class and per-level ribs,
160 and VerSe³ carries L6. Because this release uses the VerSe identifiers, its L6 *is* VerSe’s,
161 interoperable rather than local to this work. L6 is retained not for novelty but because a
162 scheme without it cannot record a six-lumbar spine at all, and must renumber the column
163 or drop a level to fit.

164 **A rib on a lumbar body receives its own class**, and this is the one class here with no
165 published counterpart. A lumbar rib and a stump rib are the same object under two counts,
166 so assigning it a number would commit the label to one reading of an enumeration anomaly.
167 A scheme that numbers every rib 1–12 has nowhere to put a thirteenth: the annotator must
168 either call it rib 12 — asserting the vertebra beneath it is thoracic, the very question at
169 issue — or discard it. TotalSegmentator carries ribs 1–12 per side and no such class, and




FIG. 2 The two anchors and the interval between them, from the released labels, posterior view. Rostral anchor: the lowest rib-bearing vertebra and the rib that makes it one. Caudal anchor: S1. Both are found from the labels rather than assumed, and panels are coloured by role rather than identifier, since colouring T12 would assert the count the figure derives. Panels (b) and (c) reach the same interval of four from opposite ends and are indistinguishable by it. All panels are at one millimetre scale, aligned at the sacral base.

170 VerSe's T13 is a vertebra rather than a rib. 16 released records carry a lumbar rib and 18
171 carry an L6.

172 Figure 2 shows why the interval is the primitive and the label is derived. Among the
173 records with four rib-free vertebrae, some reach that count through a lumbar rib and some
174 without one; the two are indistinguishable by the count alone and are distinguished here
175 only because rib presence is itself labelled.

176 Sixteen records carry a lumbar rib in the released volumes: thirteen bilateral and three
177 unilateral, all three on the right. That laterality is reported because it is what the release
178 contains and not as a finding — sixteen records cannot support a statement about side
179 preference, and none is made.

180 Eighteen records carry an L6. Fourteen are labelled lumbarisation, two sacralisation,
181 one semi-sacralisation and one normal, which is worth stating because it shows the two
182 annotation axes disagreeing in the open: a record can carry six lumbar-type vertebrae and
183 still be read as a sacralisation, or as unremarkable, depending on where the reader started
184 the count.

185 Both counts are derived from the released label volumes, and so are the manifest fields
186 that report them. That is a correction rather than a design note: through v5 the `has_16`

187 flag was true for exactly one record, which contains no L6, and false for all eighteen that
188 do, while `has_lumbar_rib` was false everywhere and `n_lumbar_labels` read zero in 799 of
189 802 records. A reader selecting the six-lumbar cases on those fields would have received one
190 record that is not one, and a reader selecting lumbar ribs an empty set, with nothing to
191 indicate either. The fields are rebuilt for this release by counting identifiers 25 and 74–75 in
192 each volume, and `lumbar_rib_side` is added alongside them.

193 II.F. Computational tools

194 *Access.* Every processing step described above is performed by code released under an
195 open-source licence and archived with the dataset; no step depends on software that is
196 not publicly obtainable. The release archive carries the reference loader, the label-scheme
197 definition and the quality-control scripts that generated the tables in this manuscript.

198 *Versions and parameters.* Segmentation of femora, the sacral sub-division and the per-
199 level rib numbering used TotalSegmentator⁴ (`total` task, release 2.x) on a single graphics
200 processing unit. The binary rib network is Möller’s¹⁰ released weights, applied through the
201 nnU-Net v2 framework¹². The pelvic completion for records whose pelvis was absent used a
202 five-fold nnU-Net v2 ensemble, applied out-of-fold: each record is completed by the fold that
203 never trained on it, so no record is completed by a model that has seen it. Image handling
204 throughout used NiBabel and SciPy; no image is resampled or reoriented when a label is
205 written, so every label shares its image’s grid and affine exactly.

206 *Generative artificial intelligence.* Large-language-model assistance was used for manuscript
207 preparation — drafting and editing of text, and generation of analysis and figure code —
208 under author review. It was not used to generate, annotate or interpret any imaging data,
209 and no scientific claim in this manuscript originates from it. All numbers reported here are
210 computed by the released code from the released volumes.

211 II.G. Validation

212 Validation is layered, and the ordering matters: measurements computed on a corpus
213 that has not established internal consistency do not look broken, they look finished.

214 *Geometric invariants.* Every case is checked for label geometry agreeing with its CT in

TABLE I Release validation. Gates run in order and a failure stops the chain.

Check	Result
Geometric invariants (all cases)	802/802 pass
label shape, affine, spacing match CT	—
no identifier outside the scheme	—
left/right structures correctly sided	—
Ribs present	11,548
not evaluable (expected vertebra outside the FOV)	5,788
no vertebra within the anchor distance	11
evaluable	5,749
matching the expected vertebra	5,747
residual offset	2 (0.035% of evaluable)
misnumbered cases	0
Spinopelvic medians within published range	6/6

shape, affine and voxel spacing; for identifiers within the published scheme; for a non-empty label; and for sided structures falling on the correct sides, with the left–right axis read from the affine rather than assumed. The check exists because a renumbering pass once wrote a reoriented array under a canonical affine and transposed a label away from its image, a failure invisible downstream.

It compares *each sided pair separately*, and neither alternative finds what it finds. Ribs alone passed all 802 while four records carried `left_hip` on the patient’s right; pooling the pairs recovered three and missed one, because hips and femora outweigh the ribs by voxel count and a large correctly-sided structure masks a smaller transposed one. A gate that passes everything is not evidence of a clean corpus until it has been shown capable of failing.

Rib–vertebra incidence. Each rib is matched to its nearest vertebra and compared against the vertebra its number implies; the chain is given in Table I. The denominator is stated there rather than only the numerator because quoting the two residual offsets against all 11,548 ribs would halve the rate by counting ribs the check never examined — and because more than half the ribs in a screening abdominal CT being unevaluable is itself a fact about the corpus worth reporting. Both offsets are field-of-view truncations on individually

231 characterised cases, one at +1 level and one at −1, and are reported rather than suppressed:
232 a threshold tuned until the count reaches zero yields a cleaner number and a less informative
233 one.

234 The same check carries an invariant that reads as a null result. Of the 5,749 evaluable
235 ribs, *zero* are assigned to a lumbar vertebra — not because no lumbar ribs exist, since 16
236 records carry one, but because such a rib takes class 74 or 75 and never enters the numbered
237 set the test examines. A numbered rib landing on a lumbar vertebra would be a counting
238 error of exactly the kind this dataset exists to expose.

239 *Agreement with published values.* Derived spinopelvic measures are compared against
240 Vialle *et al.*¹³ (Table II). Pelvic incidence is the strongest of these checks because it is a
241 morphological property of the pelvis rather than a posture, and so is directly comparable to
242 a standing reference cohort.

243 II.H. Castellvi typing

244 Every record whose vertebral count departs from five carries a Castellvi grade¹⁴ in the
245 released manifest: 33 cases, typed I–IV with the *a/b* unilateral–bilateral qualifier. This is
246 released as an annotation layer because the grade and the count are *different axes* and the
247 dataset is built to show that they are: grade IIIb occurs here at rib-free counts of four, five
248 and six alike, and seven of the 33 graded cases carry a perfectly normal count of five. A
249 corpus that recorded only the count would describe those seven as unremarkable.

250 The grades were assigned by a radiology resident physician. The distinction the classifi-
251 cation turns on — bony fusion against articulation without it¹ — is the one carrying clinical
252 weight.

253 II.I. Surgical hardware, and why a threshold is not enough

254 Metal is trivial to detect and hard to interpret, and the difference matters here for one rea-
255 son: **an iatrogenic fusion is indistinguishable from a congenital one to a distance**
256 **measurement.** A cage-bridged interspace reads as “no gap” exactly as a congenitally fused
257 transitional vertebra does, so an unrecorded instrumented case is a false positive on the cen-
258 tral measure — and identifiers 76–79 were declared in every prior version and populated in

TABLE II Derived measures against a published asymptomatic reference ($n = 260$, standing)¹³. This cohort is supine, and the two postural measures depart in opposite directions — sacral slope 4.1° lower, pelvic tilt 4.3° higher. That is one discrepancy, not two: $PI = PT + SS$ holds by construction, so with pelvic incidence at its published value any fall in one appears as an equal rise in the other.

Measure	This dataset (supine)	Reference (standing)
Pelvic incidence	54.7°	54.7 ± 10.6
Pelvic tilt	17.3°	13 ± 6
Sacral slope	36.9°	41.0 ± 8.4
Lumbar lordosis (supine)	52.7°	43 ± 11.2 (standing)
PI – LL mismatch	2.6°	centred near 0

259 none.

260 *II.I.0.a. Detection over-calls by a factor of seven.* A threshold at 1800 HU inside a shell
261 around the labelled skeleton flags 84 of the 802 records. Raising it to 2500 HU — the lower of
262 the two values validated in the metal-segmentation literature — leaves 52 with a component
263 above a 40-voxel floor. The same radiology resident physician read all 52 and confirmed 11,
264 rejecting 41 (79%) as contrast, calcification and reconstruction artefact. Instrumentation
265 prevalence in this cohort is 1.4%, not the 10.5% an unreviewed threshold reports.

266 *II.I.0.b. Saturation does not separate implant from artefact.* Both groups reach the scan-
267 ner ceiling and 4 rejected proposals *exceed* it, peaking at 11,798 HU — the signature of
268 reconstruction overshoot rather than of denser metal. What separates them is position and
269 size together: a real surgical site *and* volume, every confirmed implant at $2,586 \text{ mm}^3$ or
270 above and every rejection at $1,768 \text{ mm}^3$ or below.

271 *II.I.0.c. The subtype block had to be extended.* Identifiers 76–79 name spinal instrumen-
272 tation — generic, cage, screw-and-rod, plate — and cannot describe what this cohort holds,
273 so **80 arthroplasty**, **81 sacroiliac screw** and **82 osteosynthesis** were added. The dis-
274 tinction is clinical, not cosmetic: osteosynthesis holds parts of the same bone together where
275 arthroplasty replaces a joint, and without it the shape rules call a femoral stem a rod, since
276 a stem is long and thin and that is all “linear” tests. The confirmed cases comprise 8 arthro-
277 plasties, one osteosynthesis, one sacroiliac screw fixation and one pair of threaded cylindrical

278 interbody cages.

279 *II.I.0.d. Metal outranks bone, and that is a subtraction.* In every confirmed case the
280 implant was already inside a bone label, a segmenter handed a bright object against a
281 cortical surface having taken it for bone. Naming the hardware therefore reclaims voxels
282 rather than adding them: 1,538,852 voxels across the eleven records. The consequence
283 reaches past bookkeeping. Pelvic incidence and pelvic tilt are measured from the femoral
284 head centre, and in 8 of these cases that head is an implant — any spinopelvic parameter
285 previously computed from them was measured on metal.

286 III. DATA FORMAT AND USAGE NOTES

287 Each record is a gzipped NIfTI image/label pair sharing an identical 4×4 affine, canon-
288 icalised to PIR, requiring no resampling. Masks are NIfTI rather than DICOM: the policy
289 prefers DICOM where applicable, and for volumetric segmentation masks NIfTI is what the
290 downstream tooling expects.

291 Splits are patient-grouped and LSTV-stratified five-fold, frozen and shipped with the
292 data as `splits_5fold.json`. The five validation folds partition all 802 records, so the
293 release provides cross-validation and *not* a held-out test set. Anyone reporting a single held-
294 out number should carve one out and say which records it holds, because the folds shipped
295 here will not do that job.

296 Stratification is on the transitional subtype rather than on a binary LSTV flag, and it
297 enforces a minimum of each rare subtype in every validation fold: each fold’s validation
298 set carries three sacralisation and three lumbarisation records, so no fold is evaluated on a
299 cohort missing the class the dataset exists for. With 15 lumbarisation and 15 sacralisation
300 records across 802, that is the practical limit of what stratification can guarantee, and it is
301 why a fold-level metric on the rare classes carries an error bar far wider than its own decimal
302 places.

303 Table III lists the key data types and metadata against the fraction of records for which
304 each is populated. Two entries are worth reading carefully. The Castellvi grade is present
305 for 4.1% of records because it is defined only for the transitional ones; among those it is
306 complete. A pelvic series identifier is present for 88.9% because the crosswalk records it only
307 where a pelvic annotation was actually placed, and a user joining on it should expect the

TABLE III Key data types and metadata, and the fraction of the 802 released records for which each is populated. Counts are obtained by reading the released manifest, not by assertion. A field below 100% is not an error: age and sex are absent where the source collection did not publish them, a pelvic series identifier exists only where a pelvic annotation was placed, and a Castellvi grade exists only for the transitional records that were graded.

Field	Records	Available
Identifier, image path, label path, configuration, match type	802	100.0%
Patient position, tube potential, section thickness, kernel	802	100.0%
Transitional label and class	802	100.0%
Spine and pelvic annotation origin	802	100.0%
Instrumentation and partial-annotation flags	802	100.0%
Image/label alignment check	802	100.0%
Spine series identifier; bone-fraction check	782	97.5%
Scanner manufacturer and model	768	95.8%
Sex	749	93.4%
Age	709	88.4%
Pelvic series identifier	713	88.9%
Castellvi grade	33	4.1%

308 shortfall rather than treat it as loss.

309 The archive of record is Zenodo (10.5281/zenodo.22139642). Code and documentation
310 are on GitHub at the tagged release commit; any model-hub copy is a convenience mirror,
311 not the archive.

312 IV. DESCRIPTIVE ANALYSIS

313 The cohort is a colorectal cancer screening population of 802 records. 738 are recorded
314 female (393) or male (345); 11 carry DICOM’s *other* value, a recorded value and not a
315 missing one, and 53 carry no sex field. Age is present for 709 (median 59, range 50–89). All
316 802 are counted in totals, and records outside the female/male dichotomy are excluded from
317 sex-stratified measures — a limitation of those measures rather than of the cohort, stated




FIG. 3 Count-free measures, valid without knowing which vertebra is which. (a) Vertebrae between the lowest rib and the sacrum: an interval, not a level assignment. (b) Lowest rib length as a fraction of the rib above, showing two populations near 0.33 and 0.69 rather than one distribution with a tail. (c) Lowest lumbar transverse span against its gap to the sacral ala; cases carrying a source transitional label are marked.

because folding eleven people into “unrecorded” would misdescribe the source. The lower age bound is the screening guideline rather than a selection applied here, and it is why the measures below sit closer to a reference range than a surgical series would.

Count-free measures (Fig. 3). Five rib-free vertebrae above the sacrum is typical; four and six are where transitional anatomy sits, and which of the two is present does not by itself determine what to call it. The lowest-rib length ratio is bimodal.

Level gradients and spinopelvic measures (Fig. 4). Vertebral dimensions increase caudally in step with the load each level carries. Pelvic incidence does not vary with age, while pelvic tilt increases — the descriptive signature of a pelvis retroverting as a spine ages.

Opportunistic measures. Every scan carries a vertebral bone density measurement at no additional dose.¹⁵ L1 trabecular attenuation is reported at the published standard site.

V. POTENTIAL APPLICATIONS AND LIMITATIONS

The principal use is the *spinopelvic interface*: spine and pelvis carry radiologist-sourced annotation on one coordinate frame in 802 records, which makes sacral morphology measurable against the lumbar column rather than in isolation — the sacral endplate and its slope, the alar corridors constraining S2-alar-iliac and iliac screw trajectories, and the L5–S1 geom-




FIG. 4 Spinopelvic measures against published reference bands (mean $\pm$ SD) from an asymptomatic standing cohort.¹³ Sacral slope is expected to sit below a standing reference because it is postural and this cohort is supine; pelvic incidence is not postural, which is why it can be compared without apology and why its agreement is the strongest check here.

etry a lumbosacral fusion must cross. Pelvic incidence is derived from the segmented sacrum and femoral heads rather than from landmarks placed on a radiograph, so it is reproducible from the labels themselves.

The transitional layer serves that use rather than competing with it: no trajectory can be planned across a junction whose segments cannot be named. Level-numbering benchmarks, variant studies and opportunistic screening follow from the same annotation.

V.A. Three uses the released measurements support

Naming a level from its own geometry. Wrong-level spine surgery runs at roughly one in 3,110 spinal procedures, and the cause most often cited is the transitional anatomy this cohort was assembled around.¹⁶ The failure is structural rather than careless: the count is ambiguous exactly where the variant sits. Per-level body height, canal width, endplate width, transverse-process span and wedge ratio are released for T11–L5 across all 802 records, and because the labels were assigned against the twelfth rib rather than by enumeration, a model trained on them is not learning the convention that fails.

Reference morphometry with its spread. The values a surgeon consults for vertebral body height, canal dimensions and pedicle width come from cadaveric series of a few dozen specimens, plotted as single values with no indication of spread.¹⁷ They cannot tell a reader whether the patient in front of them is unusual, because they never show what usual looks




FIG. 5 Lumbar and lower thoracic morphometry by level, in the format the standard reference figures use — level on the vertical axis, measurement on the horizontal. Marker, median; bar, interquartile range; line, 5th–95th percentile; the number at the right of a row is its n . (a) Vertebral body height, ventral against dorsal. (b) Superior endplate width. (c) Canal width against depth. (d) Transverse pedicle width. (e) Disc height, drawn at the interspace it occupies. (f) Trabecular attenuation at the published opportunistic-screening site. Coverage is not uniform and is drawn rather than assumed: canal depth and pedicle width rest on 330–656 records because they need an axial extent many field-limited series do not carry, while the remaining measures rest on 720–784. No comparison between levels is tested and none is implied; the intervals are dispersion.

like. The same measures are given here from 802 records with the distribution attached (Fig. 5), reproducing two properties of the classical figures — body height crosses over, dorsal exceeding ventral at T11–T12 and reversing by L4–L5, and pedicle width widens steeply into the lower lumbar spine — and adding the width of each distribution, which is the part needed to judge an individual. A larger n sharpens a distribution without changing the population it came from, and this one is a screening cohort aged 50 and over.

Patient-specific models. Planning is moving toward patient-specific biomechanical models

rather than population norms.¹⁸ Such a model needs spine, pelvis and femoral heads in one frame, which is this release’s construction; in 351 patients the annotation sits on two acquisitions in different positions, giving a within-patient change in alignment against which a predicted postural response can be checked.

Its limitations are part of it, and they are stated at the level of detail a reader would need to catch them independently.

Coverage. Thoracic ground truth is field-of-view limited and does not extend to T1. Postural angles are supine; pelvic incidence is not postural and needs no such caveat. The cohort is a colorectal screening population aged 50 and over, so its distributions should not be read as representative of a surgical one.

Declared but empty, and confirmed so. The soft-tissue identifiers (58–73) and the partial-annotation sentinel (255) are declared in the scheme and occur in no record — established by counting identifiers across all 802 released volumes rather than asserted. Both are retained so a later release can use them without renumbering, and their absence should be read as absence of annotation rather than absence of the structure. The hardware identifiers were in this position until v6 and are no longer: 76–82 are populated in the eleven instrumented records described in Sec. II.I.

Strength of the labels varies by structure, and the release does not average over that. Vertebral labels derive from radiologist-supervised source annotations, corrected where those sources were wrong. Pelvic labels on records that lacked one are pseudolabelled. The rib layer is a pseudolabel whose human review was triaged by an automated rule rather than exhaustive, so its guarantee is the absence of the failure modes that rule detects. Transitional labels derive from two independent sources and are not uniformly adjudicated — which is precisely why the measures reported here are count-free — and the Castellvi grades were assigned by a radiology resident physician. The sacral sub-division inherits its level from an automatic method, and in 100 records that level places S1 above half the sacrum or below fifteen percent of it, which no first sacral segment is; those records are flagged in the released quality-control table and should be filtered before any measurement depending on the sacral endplate.

Structures are not universally present. The same census shows the release is not uniformly complete. Sacrum, hips and femora are present in all 802 records. One record carries no S1, one lacks T12, and nine lack an L5 identifier — in every case because the structure lies

391 outside the field of view or, for L5, because the lowest lumbar segment is labelled L6 or
392 incorporated into the sacrum. A user filtering on the presence of the caudal anchor should
393 filter explicitly rather than assume it.

394 VI. DISCUSSION

395 *Comparison with related material.* Against the collections in Fig. 1, the contribution
396 is the crosswalk placing spine and pelvis on the same series, together with classes for the
397 anatomy an enumeration anomaly actually produces: where another scheme must record
398 a sixth lumbar vertebra or a lumbar-borne rib as something else, this release records it as
399 itself.

400 VII. CONCLUSION

401 CTSpinoPelvic1K places spine, pelvis, per-level ribs and femora on one coordinate frame
402 in 802 records, and gives the anatomy that makes lumbar numbering ambiguous a class of
403 its own rather than recording it as something it is not. Because both counting anchors are
404 explicit, a level can be identified from a field of view that cannot support the conventional
405 count downward from C2 — which is every record here. The release is archived under a
406 persistent identifier, ships what is needed to verify it, and states per structure how strong
407 each label is.

408 DATA AVAILABILITY

409 The release described here is v7, permanently archived at [https://doi.org/10.5281/](https://doi.org/10.5281/zenodo.22242745)
410 [zenodo.22242745](https://doi.org/10.5281/zenodo.22242745). The dataset as a whole, across versions, is cited as [https://doi.org/](https://doi.org/10.5281/zenodo.22139642)
411 [10.5281/zenodo.22139642](https://doi.org/10.5281/zenodo.22139642), which always resolves to the current version. The archive car-
412 ries the reference loader, the label-scheme definition and the quality-control tables, and the
413 analysis code is distributed with it under an open-source licence. The repository URL is
414 omitted here because it identifies the authors; it is given on the title page and will be restored
415 to this section for the camera-ready version.

CONFLICT OF INTEREST

The authors have no relevant conflicts of interest to disclose.

ETHICS

This work uses publicly available, de-identified imaging from The Cancer Imaging Archive^{8,9} and annotations derived from it. The investigators obtain no identifiable private information and have no interaction with the individuals imaged, so the activity does not meet the definition of human participant research under 45 CFR 46.102. A determination to that effect has been submitted to the authors' institutional Review Board.

REFERENCES

- ¹G. P. Konin and D. M. Walz, Lumbosacral transitional vertebrae: classification, imaging findings, and clinical relevance, *AJNR Am. J. Neuroradiol.* **31**, 1778–1786 (2010), doi:10.3174/ajnr.A2036.
- ²Y. Deng et al., CTSpine1K: a large-scale dataset for spinal vertebrae segmentation in computed tomography, *Mach. Learn. Biomed. Imaging* **3**, 824–832 (2025), doi:10.59275/j.melba.2025-gf84.
- ³A. Sekuboyina et al., VerSe: a vertebrae labelling and segmentation benchmark for multi-detector CT images, *Med. Image Anal.* **73**, 102166 (2021), doi:10.1016/j.media.2021.102166.
- ⁴J. Wasserthal et al., TotalSegmentator: robust segmentation of 104 anatomic structures in CT images, *Radiol. Artif. Intell.* **5**, e230024 (2023), doi:10.1148/ryai.230024.
- ⁵H. Liebl et al., A computed tomography vertebral segmentation dataset with anatomical variations and multi-vendor scanner data, *Sci. Data* **8**, 284 (2021), doi:10.1038/s41597-021-01060-0.
- ⁶P. Liu et al., Deep learning to segment pelvic bones: large-scale CT datasets and baseline models, *Int. J. Comput. Assist. Radiol. Surg.* **16**, 749–756 (2021), doi:10.1007/s11548-021-02363-8.

⁷L. Jin et al., RibSeg v2: a large-scale benchmark for rib labeling and anatomical centerline extraction, *IEEE Trans. Med. Imaging* **43**, 570–581 (2024), doi:10.1109/TMI.2023.3313627.

⁸K. Smith et al., Data from CT COLONOGRAPHY, The Cancer Imaging Archive, 2015, doi:10.7937/K9/TCIA.2015.NWTESAY1.

⁹K. Clark et al., The Cancer Imaging Archive (TCIA): maintaining and operating a public information repository, *J. Digit. Imaging* **26**, 1045–1057 (2013), doi:10.1007/s10278-013-9622-7.

¹⁰H. Möller et al., Automated thoracolumbar stump rib detection and analysis in a large CT cohort, 2025, arXiv:2505.05004. doi:10.48550/arXiv.2505.05004.

¹¹A. Schehr, J. Kim, and G. Schwing, OpenSpineConsortium: an open-source framework for medical student engagement in computational spine imaging research, *Cureus* (2026), Published 14 July 2026. doi:10.7759/cureus.112661.

¹²F. Isensee, P. F. Jaeger, S. A. A. Kohl, J. Petersen, and K. H. Maier-Hein, nnU-Net: a self-configuring method for deep learning-based biomedical image segmentation, *Nat. Methods* **18**, 203–211 (2021), doi:10.1038/s41592-020-01008-z.

¹³R. Vialle, N. Levassor, L. Rillardon, A. Templier, W. Skalli, and P. Guigui, Radiographic analysis of the sagittal alignment and balance of the spine in asymptomatic subjects, *J. Bone Joint Surg. Am.* **87**, 260–267 (2005), doi:10.2106/JBJS.D.02043.

¹⁴A. E. Castellvi, L. A. Goldstein, and D. P. K. Chan, Lumbosacral transitional vertebrae and their relationship with lumbar extradural defects, *Spine (Phila. Pa. 1976)* **9**, 493–495 (1984), doi:10.1097/00007632-198407000-00014.

¹⁵P. J. Pickhardt, B. D. Pooler, T. Lauder, A. M. del Rio, R. J. Bruce, and N. Binkley, Opportunistic screening for osteoporosis using abdominal computed tomography scans obtained for other indications, *Ann. Intern. Med.* **158**, 588–595 (2013), doi:10.7326/0003-4819-158-8-201304160-00003.

¹⁶N. E. Epstein, A perspective on wrong level, wrong side, and wrong site spine surgery, *Surg. Neurol. Int.* **12**, 286 (2021), doi:10.25259/SNI.402.2021.

¹⁷E. C. Benzel, *Biomechanics of Spine Stabilization*, Thieme, New York, 3rd edition, 2015.

¹⁸E. Beaulieu et al., A review of finite element analysis in spine surgery decision-making, *J. Clin. Med.* **15**, 2584 (2026), doi:10.3390/jcm15072584.

Supplementary Material

CTSpinoPelvic1K: spine, pelvis, ribs and femora in one coordinate frame, annotated for lumbosacral transitional anatomy




Figure S 1: How the dataset was built. Three public collections are combined patient by patient. Neither annotation source published the CT series its labels were drawn on, so each label set is matched to the series whose bone agrees with it. In 351 of those patients the two collections had drawn their outlines on *different scans of the same person*. Numbers marked n are scans, or corrections, at that step.