

**CTSpinoPelvic1K: spine, pelvis, ribs and femora in one coordinate frame, annotated
for lumbosacral transitional anatomy**

[Author names removed for double-anonymized review]

[Affiliations removed for double-anonymized review]

(Dated: 3 September 2026)

Purpose: Lumbar levels cannot be numbered reliably when a lumbosacral transitional vertebra (LSTV) is present. Levels are conventionally established by counting downward from C2, the reference standard for numeration¹. An abdominal field of view lacks that landmark, leaving a specific ambiguity: four rib-free lumbar vertebrae may mean an L1 bearing a lumbar rib or an L5 assimilated to the sacrum, and six may mean a true sixth lumbar vertebra or a T12 whose ribs are aplastic². CT-SpinoPelvic1K is built to test whether local morphology can resolve them without counting from the C2 anchor. It provides 802 CT records anchoring each lumbar level on the lowest rib-bearing vertebra above and the first sacral segment below, with radiologist-sourced spine *and* pelvic annotation on one frame, per-level ribs and femora. Classes for a sixth lumbar vertebra, a thirteenth thoracic vertebra, a separate first sacral segment, and lumbar ribs make the set expressive enough for anatomical anomalies.

Acquisition and Validation Methods: Records pair TCIA CT COLONOGRAPHY imaging with CTSpine1K³ vertebral and CTPelvic1K⁴ pelvic labels on the acquisition with highest bone coverage, under a VerSe-native scheme. Validation: geometric invariants (802/802 pass); rib-vertebra incidence across 5,749 evaluable ribs (2 offsets, 0.035%); and agreement of spinopelvic measures with published values.

Data Format and Usage Notes: Gzipped NIfTI image/label pairs sharing one affine, canonicalized to PIR, with frozen patient-grouped LSTV-stratified five-fold splits and a reference loader; archived at 10.5281/zenodo.22139642.

Potential Applications: Testing whether a network can classify a vertebra from local features alone; updating textbook morphometry drawn from small cadaveric series; spinopelvic assessment at the lumbosacral interface; variant studies; opportunistic screening. *Limitations:* thoracic ground truth is field-of-view limited; postural angles supine; no held-out test set; the rib layer is a triaged-review pseudolabel; Castellvi grades assigned by two radiology resident physicians.

6 I. INTRODUCTION

7 The spine commonly carries seven cervical, twelve thoracic and five lumbar vertebrae, and
8 a vertebra is identified by where it falls in that sequence. Transitional variants sit at the two
9 ends of the lumbar column — the thoracolumbar transitional vertebra (TLTV) above, the
10 lumbosacral transitional vertebra (LSTV) below — and they disturb identification in two
11 ways at once. They change what the border vertebra and its ribs *look like*: a thoracolumbar
12 border vertebra may carry a full rib, a hypoplastic stump, or none at all, and a lumbosacral
13 one may be partly or completely assimilated to the sacrum, with an anteriorly wedged body,
14 a reduced disc beneath it and hypoplastic facet joints, or conversely separated from it with
15 a squared body, lumbar-type facets and a full-height disc². They also change *how many*
16 vertebrae each region contains — eleven or thirteen thoracic, four or six lumbar.

17 Either change alone would be tractable. Together they make it difficult to state defini-
18 tively what any vertebra at these borders *is*. Four rib-free vertebrae above the sacrum may
19 mean an L1 bearing a lumbar rib or an L5 assimilated to it; six may mean a true sixth
20 lumbar vertebra or a T12 whose ribs are aplastic. Each pair yields the same count, and
21 the morphology that would separate the two readings is not decisive in every case, so what
22 remains is the sequence — conventionally established by counting downward from C2, the
23 reference standard for numeration¹, and unavailable in any spine-limited acquisition. LSTV
24 is common, reported between 4% and 30% depending on definition, and wrong-level surgery
25 is its most serious consequence.

26 Existing public collections do not make this tractable, and not because they are small
27 (Fig. 1): the spine collections omit the pelvis, the pelvic ones carry no vertebral count, and
28 the whole-body schemes have neither an L6 class nor a class for a rib borne by a lumbar
29 vertebra, so a six-lumbar spine cannot be written down in them at any segmentation quality.
30 VerSe^{5,6} is the instructive case, and Fig. 1(b) gives it: a collection can enrich for anatomical
31 variants deliberately, grade them, and still be unable to represent the ones it grades.

32 CTSpinoPelvic1K places the spine, pelvis, ribs and femora in one coordinate frame and
33 gives a class to each structure an enumeration anomaly produces: a sixth lumbar vertebra,
34 and a rib borne by a lumbar vertebra. A scheme without those classes must record such
35 anatomy as something it is not.

36 *Where the material is.* Everything described here is v7 of a single archived dataset,

deposited at <https://doi.org/10.5281/zenodo.22242745>. Citations should use <https://doi.org/10.5281/zenodo.22139642>, the permanent identifier for the dataset across all of its versions, which resolves to whichever version is current; the version identifier above pins the exact release measured in this manuscript. Documentation, the label scheme, the reference loader and the per-record quality-control tables are distributed with the archive and are described in Sec. III.

I.A. Vertebra labeling when the scan cannot settle the count

The limitation, stated plainly. No scan in this corpus contains C2 or the whole cervical spine, so the conventional count cannot be performed. That is a property of abdominal imaging rather than of the annotation, and it means a variant at the thoracolumbar junction cannot be *resolved by counting* here at all: a thirteenth thoracic vertebra, a lumbar rib, and a *stump rib* — a hypoplastic twelfth rib, one of the cardinal indicators of a thoracolumbar transitional vertebra — produce overlapping appearances.

What the release does instead. No vertebral label is asserted where it cannot be supported. Every quantity is expressed against the two landmarks present in essentially every abdominal scan — the sacrum below and the lowest rib-bearing vertebra above — so the description is positional rather than nominal: a vertebra by how many rib-free vertebrae separate it from the sacrum, a lowest rib by its length as a fraction of the rib above it, a transverse process by its span and its distance from the ala. None of these changes according to whether a reader would have said L5 or S1, so cases that disagree about the name remain comparable. Where the junction is in view the count is determined and the conventional name follows, and the two descriptions agree.

And the morphology may be local rather than global. Counting is not the only route to an identity. A thoracic vertebra differs in shape from a lumbar one — in the span of its transverse processes relative to its body, in the proportions of its endplates and canal — and if that difference holds at the boundary itself, the phenotype is decidable from the vertebra in front of the reader without any global count. On this corpus it does hold: separating T12 from L1 on shape alone, with each patient’s own size divided out and cross-validation grouped by case, gives an area under the curve of 0.990 and 97.3% accuracy over 1485 vertebrae. Size is deliberately removed, because a classifier given raw millimeters separates




FIG. 1 Public CT collections at the lumbosacral junction. (a) Size, stated first because it is the claim this dataset does *not* make: three of the five comparators carry more scans, and RibSeg v2⁷ labels more ribs. (b) What each published scheme can *represent*, a property of its class list rather than of its segmentation quality. VerSe is graded-but-excluding: it applies the Castellvi classification and uses it to decide what to omit, leaving vertebrae partially fused to the sacrum unsegmented⁶ — 25 of the 33 graded records here.

67 thoracic from lumbar immediately and has learned nothing that helps at the junction, where
68 the two are nearly the same size.

69 That is reported as a property of the released measurements rather than as a method:
70 the information needed to settle these variants is present locally.

71 II. ACQUISITION AND VALIDATION METHODS

72 II.A. Sources, and why a crosswalk was necessary

73 Both CTSpine1K³ and CTPelvic1K⁴ draw part of their imaging from the TCIA CT
74 colonography collection^{8,9}, and the vertebral and pelvic annotations in both were produced
75 under board-certified radiologist supervision. Neither released the mapping from an anno-
76 tation to the specific CT series it was drawn on.

77 That mapping is not a bookkeeping detail. Every patient in CT COLONOGRAPHY was
78 scanned *twice*, prone and supine, often with more than one kernel, so one patient identifier
79 resolves to several volumes differing in posture, position, field of view and noise — and an
80 annotation carries a patient identifier and nothing finer.

81 Placing it on the wrong volume of the same patient is not merely a displacement. The
82 spine is mildly deformable, and turning a patient from supine to prone alters lumbar lordosis
83 and the alignment of each segment against the next, so a vertebra moves by a different
84 amount, and in a different direction, from the one above it. Outlines drawn on one series
85 are the wrong *shape* for the other, and no rigid realignment recovers them. In isolation such
86 a mask looks like competent work; it reveals itself only when overlaid on the image.

87 The crosswalk (Fig. S1) therefore resolves each annotation first to a patient and then,
88 within that patient’s volumes only, to the series whose bone *agrees* with it: candidates are
89 thresholded at bone attenuation and scored by overlap with the annotation mask. Bone
90 rather than image similarity, because bone is what the annotation describes — two acqui-
91 sitions of one abdomen resemble each other everywhere except in where the skeleton sits.

92 II.B. Fused and split records

93 The two annotation sets were placed independently, so they do not always land on the
94 same series. Of 802 records, 342 carry spine and pelvic annotations on the same volume
95 (*fused*), 351 are *separate*, 89 are *spine_only* with no pelvic annotation for that patient, and
96 20 are *pelvic_only*. These four are the values of the manifest’s `match_type` field; the coarser
97 `config` field folds *separate* into *spine_only*, so a user wanting the four types should read
98 `match_type`. The split is the crosswalk’s main finding rather than bookkeeping: in those
99 separate records — nearly half of all patients — the two public collections annotated *different*
100 *volumes of the same person*, and neither recorded which. A pelvic annotation placed on the
101 other acquisition of a prone/supine pair cannot be combined with the spine annotation
102 without resampling across a change of posture. Those records are retained rather than
103 discarded, and held out of the postural analyses, because two acquisitions of one patient
104 differing in spinal alignment are a resource for a mobility study rather than a defect.

105 II.C. Densification, and why it is asymmetric

106 A corpus in which half the records lack a pelvis is not usable as a spinopelvic dataset,
107 so the missing structures were pseudolabeled — asymmetrically (Fig. S1), because the two
108 directions are not equivalent.

109 Pelvis onto spine-only records is the safe direction. The sacrum and innominate bones
110 are single objects with no enumeration: no count to get wrong, no naming convention to
111 choose, no analog of a transitional vertebra. Quality is reported as held-out performance on
112 the *pelvic_only* records, withheld from training for exactly this purpose.

113 The reverse direction carries the whole problem this dataset exists to address, since a
114 pseudolabeled spine must commit to a vertebral count and on a transitional vertebra will
115 commit to the commoner reading. Those 20 records were therefore corrected by hand rather
116 than densified automatically — tractable at twenty cases, which is the practical reason the
117 asymmetry runs the way it does.

118 II.D. Ribs

119 No public rib annotation covers this cohort, and the two available automatic sources fail
120 in complementary ways.

121 Möller’s binary rib network¹⁰ is a detection model of high reported accuracy, but it
122 segments ribs that lie *entirely within* the field of view; a rib clipped by the edge of an
123 abdominal acquisition is liable to be missed altogether. Since this is a colonography cohort,
124 the lower ribs are routinely clipped, and those are precisely the ribs the count depends on.

125 TotalSegmentator¹¹ does segment clipped ribs and supplies the per-rib numbering that a
126 binary network cannot. Its characteristic failure is at the other end of the bone: the most
127 medial portion of the rib — roughly the last tenth to sixth approaching the costovertebral
128 joint — is frequently absent. A rib truncated at its medial end never reaches the vertebra
129 it articulates with, so the incidence test that assigns a rib to its level has nothing to test.

130 The two were therefore unioned (Fig. S1): TotalSegmentator supplies numbering and the
131 medial-clipped cases, Möller supplies detection where its output is more complete, and the
132 union carries TotalSegmentator’s per-level identities.

133 The result is a pseudolabel whose review was triaged rather than exhaustive, which is
134 worth stating plainly because a layer described only as “human-corrected” invites the reader
135 to assume every bone was looked at. A label-based rule flagged every record in which one
136 rib bone carried more than one identifier, or in which a rib failed to reach the vertebra it
137 articulates with — a single bone with two numbers being a numbering error by construction,
138 whatever the anatomy. The rule selected 152 records; 148 were reviewed and corrected, and 4

were not and are identified in the release. Reviewers were medical students working through OpenSpineConsortium¹², trained against a written labeling protocol, with every correction recorded against the annotator who made it. Records passing the rule were not individually inspected, so the layer’s guarantee is that it is free of the failure modes the rule detects, and no stronger. The residual rib–vertebra offsets reported in Sec. II.G are the independent check on what the rule missed.

II.E. Label scheme and counting anchors

The scheme is VerSe-native: vertebrae keep their VerSe identifiers (C1–C7 = 1–7, T1–T12 = 8–19, L1–L6 = 20–25, sacrum 26), and every non-VerSe structure takes a fixed identifier above that range — S1 carved from the sacrum (29), hips (30–31), femora (32–33), ribs (34–45 left, 46–57 right), soft tissue (58–73), lumbar ribs (74–75) and surgical hardware (76–82).

Two anchors make the count decidable, and Fig. 2 shows both: the lowest rib-bearing vertebra above, **S1** carved from the sacrum below. Each earns its place by what it supports rather than by being an endpoint — S1 supplies the sacral endplate landmark that sacral slope and pelvic incidence are measured from, and with both brackets present L5 and L6 are distinguishable where a spine-limited view alone would not be.

How S1 is cut. The sacrum’s outer boundary is the source annotation and is never altered; S1 is the cranial part of that same bone, separated by a plane. The plane’s normal is the sacrum’s own cranio-caudal axis, made orthogonal to a left–right axis that is measured rather than assumed — the normal of the reflection plane best mapping the sacrum onto itself, which removes the patient’s rotation within the scanner. That rotation is not negligible: median 4.5°, maximum 16.1°, and 508 of 801 records above 3°. The level of the cut preserves the sub-division volume of the preceding release, so the plane changes the shape and orientation of the S1/S2 boundary and not how much of the sacrum is called S1. Per-record geometry ships with the archive.

Neither anchor is novel: TotalSegmentator¹¹ carries a separate S1 class and per-level ribs, and VerSe⁵ carries L6. Because this release uses the VerSe identifiers, its L6 *is* VerSe’s, interoperable rather than local to this work. L6 is retained not for novelty but because a scheme without it cannot record a six-lumbar spine at all, and must renumber the column




FIG. 2 The two anchors and the interval between them, from the released labels, posterior view. Rostral anchor: the lowest rib-bearing vertebra and the rib that makes it one. Caudal anchor: S1. Both are found from the labels rather than assumed, and panels are colored by role rather than identifier, since coloring T12 would assert the count the figure derives. Panels (b) and (c) reach the same interval of four from opposite ends and are indistinguishable by it. All panels are at one millimeter scale, aligned at the sacral base.

169 or drop a level to fit.

170 **A rib on a lumbar body receives its own class**, and this is the one class here with no
 171 published counterpart. A lumbar rib and a stump rib are the same object under two counts,
 172 so assigning it a number would commit the label to one reading of an enumeration anomaly.
 173 A scheme that numbers every rib 1–12 has nowhere to put a thirteenth: the annotator must
 174 either call it rib 12 — asserting the vertebra beneath it is thoracic, the very question at
 175 issue — or discard it. TotalSegmentator carries ribs 1–12 per side and no such class, and
 176 VerSe’s T13 is a vertebra rather than a rib. 16 released records carry a lumbar rib and 18
 177 carry an L6.

178 Figure 2 shows why the interval is the primitive and the label is derived. Among the
 179 records with four rib-free vertebrae, some reach that count through a lumbar rib and some
 180 without one; the two are indistinguishable by the count alone and are distinguished here
 181 only because rib presence is itself labeled.

182 Sixteen records carry a lumbar rib in the released volumes: thirteen bilateral and three
 183 unilateral, all three on the right. That laterality is reported because it is what the release
 184 contains and not as a finding — sixteen records cannot support a statement about side
 185 preference, and none is made.

186 Eighteen records carry an L6. Fourteen are labeled lumbarization, two sacralization,
187 one semi-sacralization and one normal, which is worth stating because it shows the two
188 annotation axes disagreeing in the open: a record can carry six lumbar-type vertebrae and
189 still be read as a sacralization, or as unremarkable, depending on where the reader started
190 the count.

191 Both counts are derived from the released label volumes, and so are the manifest fields
192 that report them. That is a correction rather than a design note: through v5 the `has_l6`
193 flag was true for exactly one record, which contains no L6, and false for all eighteen that
194 do, while `has_lumbar_rib` was false everywhere and `n_lumbar_labels` read zero in 799 of
195 802 records. A reader selecting the six-lumbar cases on those fields would have received one
196 record that is not one, and a reader selecting lumbar ribs an empty set, with nothing to
197 indicate either. The fields are rebuilt for this release by counting identifiers 25 and 74–75 in
198 each volume, and `lumbar_rib_side` is added alongside them.

199 II.F. Computational tools

200 *Access.* Every processing step described above is performed by code released under an
201 open-source licence and archived with the dataset; no step depends on software that is
202 not publicly obtainable. The release archive carries the reference loader, the label-scheme
203 definition and the quality-control scripts that generated the tables in this manuscript.

204 *Versions and parameters.* Segmentation of femora, the sacral sub-division and the per-
205 level rib numbering used TotalSegmentator¹¹ (`total` task, release 2.x) on a single graphics
206 processing unit. The binary rib network is Möller’s¹⁰ released weights, applied through the
207 nnU-Net v2 framework¹³. The pelvic completion for records whose pelvis was absent used a
208 five-fold nnU-Net v2 ensemble, applied out-of-fold: each record is completed by the fold that
209 never trained on it, so no record is completed by a model that has seen it. Image handling
210 throughout used NiBabel and SciPy; no image is resampled or reoriented when a label is
211 written, so every label shares its image’s grid and affine exactly.

212 *Generative artificial intelligence.* Large-language-model assistance was used for manuscript
213 preparation — drafting and editing of text, and generation of analysis and figure code —
214 under author review. It was not used to generate, annotate or interpret any imaging data,
215 and no scientific claim in this manuscript originates from it. All numbers reported here are

TABLE I Release validation. Gates run in order and a failure stops the chain.

Check	Result
Geometric invariants (all cases)	802/802 pass
label shape, affine, spacing match CT	—
no identifier outside the scheme	—
left/right structures correctly sided	—
Ribs present	11,548
not evaluable (expected vertebra outside the FOV)	5,788
no vertebra within the anchor distance	11
evaluable	5,749
matching the expected vertebra	5,747
residual offset	2 (0.035% of evaluable)
misnumbered cases	0
Spinopelvic medians within published range	6/6

216 computed by the released code from the released volumes.

217 II.G. Validation

218 Validation is layered, and the ordering matters: measurements computed on a corpus
 219 that has not established internal consistency do not look broken, they look finished.

220 *Geometric invariants.* Every case is checked for label geometry agreeing with its CT in
 221 shape, affine and voxel spacing; for identifiers within the published scheme; for a non-empty
 222 label; and for sided structures falling on the correct sides, with the left–right axis read from
 223 the affine rather than assumed. The check exists because a renumbering pass once wrote
 224 a reoriented array under a canonical affine and transposed a label away from its image, a
 225 failure invisible downstream.

226 The check compares *each sided pair separately*, and both simpler alternatives miss cases
 227 it catches. Testing the ribs alone passed all 802 records while four of them carried **left_hip**
 228 on the patient’s right. Pooling all sided structures into one test recovered three of those four
 229 and missed the fourth, because hips and femora outweigh the ribs by voxel count, so a large

230 correctly-sided structure masks a smaller transposed one. A gate that passes everything is
231 not evidence of a clean corpus until it has been shown capable of failing.

232 *Rib-vertebra incidence.* Each rib is matched to its nearest vertebra and compared against
233 the vertebra its number implies; the chain is given in Table I. The denominator is stated
234 there rather than only the numerator because quoting the two residual offsets against all
235 11,548 ribs would halve the rate by counting ribs the check never examined — and because
236 more than half the ribs in a screening abdominal CT being unevaluable is itself a fact
237 about the corpus worth reporting. Both offsets are field-of-view truncations on individually
238 characterized cases, one at +1 level and one at −1, and are reported rather than suppressed:
239 a threshold tuned until the count reaches zero yields a cleaner number and a less informative
240 one.

241 The same check carries an invariant that reads as a null result. Of the 5,749 evaluable
242 ribs, *zero* are assigned to a lumbar vertebra — not because no lumbar ribs exist, since 16
243 records carry one, but because such a rib takes class 74 or 75 and never enters the numbered
244 set the test examines. A numbered rib landing on a lumbar vertebra would be a counting
245 error of exactly the kind this dataset exists to expose.

246 *Agreement with published values.* Derived spinopelvic measures are compared against
247 Vialle *et al.*¹⁴ (Table II). Pelvic incidence is the strongest of these checks because it is a
248 morphological property of the pelvis rather than a posture, and so is directly comparable to
249 a standing reference cohort.

250 II.H. Castellvi typing

251 Every record whose vertebral count departs from five carries a Castellvi grade¹⁵ in the
252 released manifest: 33 cases, typed I–IV with the *a/b* unilateral–bilateral qualifier. This is
253 released as an annotation layer because the grade and the count are *different axes* and the
254 dataset is built to show that they are: grade IIIb occurs here at rib-free counts of four, five
255 and six alike, and seven of the 33 graded cases carry a perfectly normal count of five. A
256 corpus that recorded only the count would describe those seven as unremarkable.

257 The grades were assigned by two radiology resident physicians. The distinction the
258 classification turns on — bony fusion against articulation without it² — is the one carrying
259 clinical weight.

TABLE II Derived measures against a published asymptomatic reference ($n = 260$, standing)¹⁴. This cohort is supine, and the two postural measures depart in opposite directions — sacral slope 4.1° lower, pelvic tilt 4.3° higher. That is one discrepancy, not two: $PI = PT + SS$ holds by construction, so with pelvic incidence at its published value any fall in one appears as an equal rise in the other.

Measure	This dataset (supine)	Reference (standing)
Pelvic incidence	54.7°	54.7 ± 10.6
Pelvic tilt	17.3°	13 ± 6
Sacral slope	36.9°	41.0 ± 8.4
Lumbar lordosis (supine)	52.7°	43 ± 11.2 (standing)
PI – LL mismatch	2.6°	centered near 0

260 II.I. Surgical hardware, and why a threshold is not enough

261 Metal is trivial to detect and hard to interpret, and the difference matters here for one rea-
262 son: **an iatrogenic fusion is indistinguishable from a congenital one to a distance**
263 **measurement**. A cage-bridged interspace reads as “no gap” exactly as a congenitally fused
264 transitional vertebra does, so an unrecorded instrumented case is a false positive on the cen-
265 tral measure — and identifiers 76–79 were declared in every prior version and populated in
266 none.

267 *II.I.0.a. Detection over-calls by a factor of seven.* A threshold at 1800 HU inside a shell
268 around the labeled skeleton flags 84 of the 802 records. Raising it to 2500 HU — the lower of
269 the two values validated in the metal-segmentation literature — leaves 52 with a component
270 above a 40-voxel floor. One of the two radiology resident physicians read all 52 and confirmed
271 11, rejecting 41 (79%) as contrast, calcification and reconstruction artefact. Instrumentation
272 prevalence in this cohort is 1.4%, not the 10.5% an unreviewed threshold reports.

273 *II.I.0.b. Saturation does not separate implant from artefact.* Both groups reach the scan-
274 ner ceiling and 4 rejected proposals *exceed* it, peaking at 11,798 HU — the signature of
275 reconstruction overshoot rather than of denser metal. What separates them is position and
276 size together: a real surgical site *and* volume, every confirmed implant at $2,586 \text{ mm}^3$ or
277 above and every rejection at $1,768 \text{ mm}^3$ or below.

278 *II.I.0.c. The subtype block had to be extended.* Identifiers 76–79 name spinal instrumen-
 279 tation — generic, cage, screw-and-rod, plate — and cannot describe what this cohort holds,
 280 so **80 arthroplasty**, **81 sacroiliac screw** and **82 osteosynthesis** were added. The dis-
 281 tinction is clinical, not cosmetic: osteosynthesis holds parts of the same bone together where
 282 arthroplasty replaces a joint, and without it the shape rules call a femoral stem a rod, since
 283 a stem is long and thin and that is all “linear” tests. The confirmed cases comprise 8 arthro-
 284 plasties, one osteosynthesis, one sacroiliac screw fixation and one pair of threaded cylindrical
 285 interbody cages.

286 *II.I.0.d. Metal outranks bone, and that is a subtraction.* In every confirmed case the
 287 implant was already inside a bone label, a segmenter handed a bright object against a
 288 cortical surface having taken it for bone. Naming the hardware therefore reclaims voxels
 289 rather than adding them: 1,538,852 voxels across the eleven records. The consequence
 290 reaches past bookkeeping. Pelvic incidence and pelvic tilt are measured from the femoral
 291 head center, and in 8 of these cases that head is an implant — any spinopelvic parameter
 292 previously computed from them was measured on metal.

293 III. DATA FORMAT AND USAGE NOTES

294 Each record is a gzipped NIfTI image/label pair sharing an identical 4×4 affine, canon-
 295 icalized to PIR, requiring no resampling. Masks are NIfTI rather than DICOM: the policy
 296 prefers DICOM where applicable, and for volumetric segmentation masks NIfTI is what the
 297 downstream tooling expects.

298 Splits are patient-grouped and LSTV-stratified five-fold, frozen and shipped with the
 299 data as `splits_5fold.json`. The five validation folds partition all 802 records, so the
 300 release provides cross-validation and *not* a held-out test set. Users reporting a single held-
 301 out number should carve one out and state which records it holds, because the folds shipped
 302 here will not do that job.

303 Stratification is on the transitional subtype rather than on a binary LSTV flag, and it
 304 enforces a minimum of each rare subtype in every validation fold: each fold’s validation
 305 set carries three sacralization and three lumbarization records, so no fold is evaluated on a
 306 cohort missing the class the dataset exists for. With 15 lumbarization and 15 sacralization
 307 records across 802, that is the practical limit of what stratification can guarantee, and it is

TABLE III Key data types and metadata, and the fraction of the 802 released records for which each is populated. Counts are obtained by reading the released manifest, not by assertion. A field below 100% is not an error: age and sex are absent where the source collection did not publish them, a pelvic series identifier exists only where a pelvic annotation was placed, and a Castellvi grade exists only for the transitional records that were graded.

Field	Records	Available
Identifier, image path, label path, configuration, match type	802	100.0%
Patient position, tube potential, section thickness, kernel	802	100.0%
Transitional label and class	802	100.0%
Spine and pelvic annotation origin	802	100.0%
Instrumentation and partial-annotation flags	802	100.0%
Image/label alignment check	802	100.0%
Spine series identifier; bone-fraction check	782	97.5%
Scanner manufacturer and model	768	95.8%
Sex	749	93.4%
Age	709	88.4%
Pelvic series identifier	713	88.9%
Castellvi grade	33	4.1%

308 why a fold-level metric on the rare classes carries an error bar far wider than its own decimal
309 places.

310 Table III lists the key data types and metadata against the fraction of records for which
311 each is populated. Two entries are worth reading carefully. The Castellvi grade is present
312 for 4.1% of records because it is defined only for the transitional ones; among those it is
313 complete. A pelvic series identifier is present for 88.9% because the crosswalk records it
314 only where a pelvic annotation was actually placed, and users joining on it should expect
315 the shortfall rather than treat it as loss.

316 The archive of record is Zenodo (10.5281/zenodo.22139642). Code and documentation
317 are on GitHub at the tagged release commit; any model-hub copy is a convenience mirror,
318 not the archive.




FIG. 3 Count-free measures, valid without knowing which vertebra is which. (a) Vertebrae between the lowest rib and the sacrum: an interval, not a level assignment. (b) Lowest rib length as a fraction of the rib above, showing two populations near 0.33 and 0.69 rather than one distribution with a tail. (c) Lowest lumbar transverse span against its gap to the sacral ala; cases carrying a source transitional label are marked.

IV. DESCRIPTIVE ANALYSIS

The cohort is a colorectal cancer screening population of 802 records. 738 are recorded female (393) or male (345); 11 carry DICOM’s *other* value, a recorded value and not a missing one, and 53 carry no sex field. Age is present for 709 (median 59, range 50–89). All 802 are counted in totals, and records outside the female/male dichotomy are excluded from sex-stratified measures — a limitation of those measures rather than of the cohort, stated because folding eleven people into “unrecorded” would misdescribe the source. The lower age bound is the screening guideline rather than a selection applied here, and it is why the measures below sit closer to a reference range than a surgical series would.

Count-free measures (Fig. 3). Five rib-free vertebrae above the sacrum is typical; four and six are where transitional anatomy sits, and which of the two is present does not by itself determine what to call it. The lowest-rib length ratio is bimodal.

Level gradients and spinopelvic measures (Fig. 4). Vertebral dimensions increase caudally in step with the load each level carries. Pelvic incidence does not vary with age, while pelvic tilt increases — the descriptive signature of a pelvis retroverting as a spine ages.

Opportunistic measures. Every scan carries a vertebral bone density measurement at no additional dose.¹⁶ L1 trabecular attenuation is reported at the published standard site.




FIG. 4 Spinopelvic measures against published reference bands (mean $\pm$ SD) from an asymptomatic standing cohort.¹⁴ Sacral slope is expected to sit below a standing reference because it is postural and this cohort is supine; pelvic incidence is not postural, which is why it can be compared without apology and why its agreement is the strongest check here.

V. POTENTIAL APPLICATIONS AND LIMITATIONS

The principal use is the *spinopelvic interface*: spine and pelvis carry radiologist-sourced annotation on one coordinate frame in 802 records, which makes sacral morphology measurable against the lumbar column rather than in isolation — the sacral endplate and its slope, the alar corridors constraining S2-alar-iliac and iliac screw trajectories, and the L5–S1 geometry a lumbosacral fusion must cross. Pelvic incidence is derived from the segmented sacrum and femoral heads rather than from landmarks placed on a radiograph, so it is reproducible from the labels themselves.

The transitional layer serves that use rather than competing with it: no trajectory can be planned across a junction whose segments cannot be named. Level-numbering benchmarks, variant studies and opportunistic screening follow from the same annotation.

V.A. Three uses the released measurements support

Naming a level from its own geometry. Wrong-level spine surgery runs at roughly one in 3,110 spinal procedures, and the cause most often cited is the transitional anatomy this cohort was assembled around.¹⁷ The failure is structural rather than careless: the count is ambiguous exactly where the variant sits. Per-level body height, canal width, endplate width, transverse-process span and wedge ratio are released for T11–L5 across all 802 records, and because the labels were assigned against the twelfth rib rather than by enumeration, a




FIG. 5 Lumbar and lower thoracic morphometry by level, in the format the standard reference figures use — level on the vertical axis, measurement on the horizontal. Marker, median; bar, interquartile range; line, 5th–95th percentile; the number at the right of a row is its n . (a) Vertebral body height, ventral against dorsal. (b) Superior endplate width. (c) Canal width against depth. (d) Transverse pedicle width. (e) Disc height, drawn at the interspace it occupies. (f) Trabecular attenuation at the published opportunistic-screening site. Coverage is not uniform and is drawn rather than assumed: canal depth and pedicle width rest on 330–656 records because they need an axial extent many field-limited series do not carry, while the remaining measures rest on 720–784. No comparison between levels is tested and none is implied; the intervals are dispersion.

354 model trained on them is not learning the convention that fails.

355 *Reference morphometry with its spread.* The values a surgeon consults for vertebral
 356 body height, canal dimensions and pedicle width come from cadaveric series of a few dozen
 357 specimens, plotted as single values with no indication of spread.¹⁸ They cannot tell a reader
 358 whether the patient in front of them is unusual, because they never show what usual looks
 359 like. The same measures are given here from 802 records with the distribution attached
 360 (Fig. 5), reproducing two properties of the classical figures — body height crosses over,

dorsal exceeding ventral at T11–T12 and reversing by L4–L5, and pedicle width widens steeply into the lower lumbar spine — and adding the width of each distribution, which is the part needed to judge an individual. A larger n sharpens a distribution without changing the population it came from, and this one is a screening cohort aged 50 and over.

Patient-specific models. Planning is moving toward patient-specific biomechanical models rather than population norms.¹⁹ Such a model needs spine, pelvis and femoral heads in one frame, which is this release’s construction; in 351 patients the annotation sits on two acquisitions in different positions, giving a within-patient change in alignment against which a predicted postural response can be checked.

Its limitations are part of it, and they are stated at the level of detail a reader would need to catch them independently.

Coverage. Thoracic ground truth is field-of-view limited and does not extend to T1. Postural angles are supine; pelvic incidence is not postural and needs no such caveat. The cohort is a colorectal screening population aged 50 and over, so its distributions should not be read as representative of a surgical one.

Declared but empty, and confirmed so. The soft-tissue identifiers (58–73) and the partial-annotation sentinel (255) are declared in the scheme and occur in no record — established by counting identifiers across all 802 released volumes rather than asserted. Both are retained so a later release can use them without renumbering, and their absence should be read as absence of annotation rather than absence of the structure. The hardware identifiers were in this position until v6 and are no longer: 76–82 are populated in the eleven instrumented records described in Sec. II.I.

Strength of the labels varies by structure, and the release does not average over that. Vertebral labels derive from radiologist-supervised source annotations, corrected where those sources were wrong. Pelvic labels on records that lacked one are pseudolabeled. The rib layer is a pseudolabel whose human review was triaged by an automated rule rather than exhaustive, so its guarantee is the absence of the failure modes that rule detects. Transitional labels derive from two independent sources and are not uniformly adjudicated — which is precisely why the measures reported here are count-free — and the Castellvi grades were assigned by two radiology resident physicians. The sacral sub-division inherits its level from an automatic method, and in 100 records that level places S1 above half the sacrum or below fifteen percent of it, which no first sacral segment is; those records are flagged in the

393 released quality-control table and should be filtered before any measurement depending on
394 the sacral endplate.

395 *Structures are not universally present.* The same census shows the release is not uniformly
396 complete. Sacrum, hips and femora are present in all 802 records. One record carries no
397 S1, one lacks T12, and nine lack an L5 identifier — in every case because the structure
398 lies outside the field of view or, for L5, because the lowest lumbar segment is labeled L6 or
399 incorporated into the sacrum. Users filtering on the presence of the caudal anchor should
400 filter explicitly rather than assume it.

401 VI. DISCUSSION

402 *Comparison with related material.* Against the collections in Fig. 1, the contribution
403 is the crosswalk placing spine and pelvis on the same series, together with classes for the
404 anatomy an enumeration anomaly actually produces: where another scheme must record
405 a sixth lumbar vertebra or a lumbar-borne rib as something else, this release records it as
406 itself.

407 VII. CONCLUSION

408 CTSpinoPelvic1K places spine, pelvis, per-level ribs and femora on one coordinate frame
409 in 802 records, and gives the anatomy that makes lumbar numbering ambiguous a class of
410 its own rather than recording it as something it is not. Because both counting anchors are
411 explicit, a level can be identified from a field of view that cannot support the conventional
412 count downward from C2 — which is every record here. The release is archived under a
413 persistent identifier, ships what is needed to verify it, and states per structure how strong
414 each label is.

415 DATA AVAILABILITY

416 The release described here is v7, permanently archived at [https://doi.org/10.5281/](https://doi.org/10.5281/zenodo.22242745)
417 [zenodo.22242745](https://doi.org/10.5281/zenodo.22242745). The dataset as a whole, across versions, is cited as [https://doi.org/](https://doi.org/10.5281/zenodo.22139642)
418 [10.5281/zenodo.22139642](https://doi.org/10.5281/zenodo.22139642), which always resolves to the current version. The archive car-
419 ries the reference loader, the label-scheme definition and the quality-control tables, and the

analysis code is distributed with it under an open-source licence. The repository URL is omitted here because it identifies the authors; it is given on the title page and will be restored to this section for the camera-ready version.

CONFLICT OF INTEREST

The authors have no relevant conflicts of interest to disclose.

ETHICS

This work uses publicly available, de-identified imaging from The Cancer Imaging Archive^{8,9} and annotations derived from it. The investigators obtain no identifiable private information and have no interaction with the individuals imaged, so the activity does not meet the definition of human participant research under 45 CFR 46.102. A determination to that effect has been submitted to the authors' institutional Review Board.

REFERENCES

- ¹J. Lian, N. Levine, and W. Cho, A review of lumbosacral transitional vertebrae and associated vertebral numeration, *Eur. Spine J.* **27**, 995–1004 (2018), doi:10.1007/s00586-018-5554-8.
- ²G. P. Konin and D. M. Walz, Lumbosacral transitional vertebrae: classification, imaging findings, and clinical relevance, *AJNR Am. J. Neuroradiol.* **31**, 1778–1786 (2010), doi:10.3174/ajnr.A2036.
- ³Y. Deng et al., CTSpine1K: a large-scale dataset for spinal vertebrae segmentation in computed tomography, *Mach. Learn. Biomed. Imaging* **3**, 824–832 (2025), doi:10.59275/j.melba.2025-gf84.
- ⁴P. Liu et al., Deep learning to segment pelvic bones: large-scale CT datasets and baseline models, *Int. J. Comput. Assist. Radiol. Surg.* **16**, 749–756 (2021), doi:10.1007/s11548-021-02363-8.
- ⁵A. Sekuboyina et al., VerSe: a vertebrae labelling and segmentation benchmark for multi-detector CT images, *Med. Image Anal.* **73**, 102166 (2021), doi:10.1016/j.media.2021.102166.

⁶H. Liebl et al., A computed tomography vertebral segmentation dataset with anatomical variations and multi-vendor scanner data, *Sci. Data* **8**, 284 (2021), doi:10.1038/s41597-021-01060-0.

⁷L. Jin et al., RibSeg v2: a large-scale benchmark for rib labeling and anatomical centerline extraction, *IEEE Trans. Med. Imaging* **43**, 570–581 (2024), doi:10.1109/TMI.2023.3313627.

⁸K. Smith et al., Data from CT COLONOGRAPHY, The Cancer Imaging Archive, 2015, doi:10.7937/K9/TCIA.2015.NWTESAY1.

⁹K. Clark et al., The Cancer Imaging Archive (TCIA): maintaining and operating a public information repository, *J. Digit. Imaging* **26**, 1045–1057 (2013), doi:10.1007/s10278-013-9622-7.

¹⁰H. Möller et al., Automated thoracolumbar stump rib detection and analysis in a large CT cohort, 2025, arXiv:2505.05004. doi:10.48550/arXiv.2505.05004.

¹¹J. Wasserthal et al., TotalSegmentator: robust segmentation of 104 anatomic structures in CT images, *Radiol. Artif. Intell.* **5**, e230024 (2023), doi:10.1148/ryai.230024.

¹²A. Schehr, J. Kim, and G. Schwing, OpenSpineConsortium: an open-source framework for medical student engagement in computational spine imaging research, *Cureus* (2026), Published 14 July 2026. doi:10.7759/cureus.112661.

¹³F. Isensee, P. F. Jaeger, S. A. A. Kohl, J. Petersen, and K. H. Maier-Hein, nnU-Net: a self-configuring method for deep learning-based biomedical image segmentation, *Nat. Methods* **18**, 203–211 (2021), doi:10.1038/s41592-020-01008-z.

¹⁴R. Vialle, N. Levassor, L. Rillardon, A. Templier, W. Skalli, and P. Guigui, Radiographic analysis of the sagittal alignment and balance of the spine in asymptomatic subjects, *J. Bone Joint Surg. Am.* **87**, 260–267 (2005), doi:10.2106/JBJS.D.02043.

¹⁵A. E. Castellvi, L. A. Goldstein, and D. P. K. Chan, Lumbosacral transitional vertebrae and their relationship with lumbar extradural defects, *Spine (Phila. Pa. 1976)* **9**, 493–495 (1984), doi:10.1097/00007632-198407000-00014.

¹⁶P. J. Pickhardt, B. D. Pooler, T. Lauder, A. M. del Rio, R. J. Bruce, and N. Binkley, Opportunistic screening for osteoporosis using abdominal computed tomography scans obtained for other indications, *Ann. Intern. Med.* **158**, 588–595 (2013), doi:10.7326/0003-4819-158-8-201304160-00003.

- ⁴⁷⁸ ¹⁷N. E. Epstein, A perspective on wrong level, wrong side, and wrong site spine surgery,
⁴⁷⁹ Surg. Neurol. Int. **12**, 286 (2021), doi:10.25259/SNI_402_2021.
- ⁴⁸⁰ ¹⁸E. C. Benzel, *Biomechanics of Spine Stabilization*, Thieme, New York, 3rd edition, 2015.
- ⁴⁸¹ ¹⁹E. Beaulieu et al., A review of finite element analysis in spine surgery decision-making, J.
⁴⁸² Clin. Med. **15**, 2584 (2026), doi:10.3390/jcm15072584.