

**CTSpinoPelvic1K: spine, pelvis, ribs and femora in one coordinate frame, annotated
for lumbosacral transitional anatomy**

[Author names removed for double-anonymized review]

[Affiliations removed for double-anonymized review]

(Dated: 6 September 2026)

Purpose: A lumbar vertebra is named by counting down from C2, the reference standard for numeration¹. Abdominal CT does not contain C2, and where a lumbosacral transitional vertebra (LSTV) alters the count the local anatomy is ambiguous: four rib-free lumbar vertebrae may mean an L1 bearing a lumbar rib or an L5 assimilated to the sacrum, and six may mean a sixth lumbar vertebra, a T12 whose ribs are aplastic, or an S1 lumbarized from the sacrum². CTSpinoPelvic1K is built so that this can be asked of the data: whether local morphology resolves them without counting from C2. It provides 802 CT records anchoring each lumbar level on the lowest rib-bearing vertebra above and the first sacral segment below, with radiologist-sourced spine *and* pelvic annotation on one frame, per-level ribs and femora. Classes for a sixth lumbar vertebra, a thirteenth thoracic vertebra, a separate first sacral segment and lumbar ribs cover the anomalies.

Acquisition and Validation Methods: Records pair TCIA CT COLONOGRAPHY imaging with CTSpine1K³ vertebral and CTPelvic1K⁴ pelvic labels on the acquisition with highest bone coverage, under a VerSe-native scheme. Validation: geometric invariants (802/802 pass); rib-vertebra incidence across 5,749 evaluable ribs (2 offsets, 0.035%); and agreement of spinopelvic measures with published values.

Data Format and Usage Notes: Gzipped NIfTI image/label pairs sharing one affine, canonicalized to PIR, with frozen patient-grouped LSTV-stratified five-fold splits and a reference loader; archived under a persistent identifier, <https://doi.org/10.5281/zenodo.22139642>.

Potential Applications: Testing whether a network can classify a vertebra from local features alone; updating textbook morphometry drawn from small cadaveric series; spinopelvic assessment at the lumbosacral interface; variant studies; opportunistic screening. *Limitations:* thoracic ground truth is field-of-view limited; postural angles supine; no held-out test set; the rib layer is a triaged-review pseudolabel; Castellvi grades are a two-reader consensus.

6 I. INTRODUCTION

7 The spine commonly carries seven cervical, twelve thoracic and five lumbar vertebrae, and
8 a vertebra is identified by where it falls in that sequence. Transitional variants sit at the
9 two ends of the lumbar column, the thoracolumbar transitional vertebra (TLTV) above and
10 the lumbosacral transitional vertebra (LSTV) below, and they disturb identification in two
11 ways at once. They change what the border vertebra and its ribs *look like*: a thoracolumbar
12 border vertebra may carry a full rib, a hypoplastic stump, or none at all, and a lumbosacral
13 one may be partly or completely assimilated to the sacrum (*sacralization*), with an anteriorly
14 wedged body, a reduced disc beneath it and hypoplastic facet joints, or conversely separated
15 from it (*lumbarization*), with a squared body, lumbar-type facets and a full-height disc².
16 They also change *how many* vertebrae each region contains: eleven or thirteen thoracic,
17 four or six lumbar.

18 Either change alone would be tractable. Together they make it difficult to state defini-
19 tively what any vertebra at these borders *is*. Four rib-free vertebrae above the sacrum may
20 mean an L1 bearing a lumbar rib or an L5 assimilated to it; six may mean a true sixth
21 lumbar vertebra, a T12 whose ribs are aplastic, or an S1 lumbarized away from the sacrum.
22 Each set yields the same count, and the morphology that would separate the readings is not
23 decisive to a human reader in every case. Whether it differs enough for a learned classifier
24 to separate them, or at least to assign each reading a probability, is the question this release
25 is built to support. Absent that, what remains is the sequence, conventionally established
26 by counting downward from C2, the reference standard for numeration¹, and unavailable in
27 any spine-limited acquisition. LSTV is common, reported between 4% and 30% depending
28 on definition, and wrong-level surgery is its most serious consequence.

29 Existing public collections cannot address this, and size is not the reason (Table I). Spine
30 datasets omit the pelvis; pelvic datasets do not number the vertebrae; and TotalSegmentator,
31 the whole-body scheme, has no class for a sixth lumbar vertebra or for a rib on a lumbar
32 vertebra. A six-lumbar spine therefore cannot be recorded in any of them, however good the
33 segmentation, and a segmenter trained on them inherits the gap: faced with an enumeration
34 anomaly it must shift the whole column by a level or absorb a vertebra into its neighbor.
35 The size of that problem has now been measured on CT, where prior labeling methods
36 assigned every vertebra correctly in 77% of subjects and an anomaly-aware method raised

that to 99%⁵. VerSe^{6,7} makes the point most clearly (Table I): it deliberately enriched for anatomical variants and graded them by Castellvi, yet left the graded vertebrae fused to the sacrum unsegmented.

The intraoperative side of the problem has prior art of its own. LevelCheck registers the intraoperative radiograph to the preoperative CT and projects the CT’s vertebral labels onto it, localizing the target level in a limited field of view to within a few percent of cases in clinical use^{8–10}. The labels it projects are placed by hand on the CT, by a skilled operator and verified by the surgeon⁸: the method assumes a correctly labeled CT and does not supply one. Automating that step, and getting it right precisely where the anatomy is anomalous, is what a dataset carrying the classes those anomalies produce is for.

CTSpinoPelvic1K places the spine, pelvis, ribs and femora in one coordinate frame and gives a class to each structure an enumeration anomaly produces: a sixth lumbar vertebra, and a rib borne by a lumbar vertebra. A scheme without those classes lacks a class vocabulary expressive enough to describe anatomic anomalies.

Where the material is. Everything described here is v8 of a single archived dataset, deposited at <https://doi.org/10.5281/zenodo.22293878>. Citations should use <https://doi.org/10.5281/zenodo.22139642>, the permanent identifier for the dataset across all of its versions, which resolves to whichever version is current; the version identifier above pins the exact release measured in this manuscript. Documentation, the label scheme, the reference loader and the per-record quality-control tables are distributed with the archive and are described in Sec. III.

I.A. Vertebra labeling when the scan cannot settle the count

The limitation, stated plainly. No scan in this corpus contains C2, so the conventional count cannot be performed: a property of abdominal imaging, not of the annotation. A thirteenth thoracic vertebra, a lumbar rib and a *stump rib* (a hypoplastic twelfth rib, a cardinal indicator of a thoracolumbar transitional vertebra) leave the same number of rib-free vertebrae between the two borders of the lumbar spine, the last rib-bearing vertebra and the sacrum.

What the release does in addition. Every vertebra carries the identifier its radiologist-sourced annotation assigned, corrected where the source was wrong; those labels are the

TABLE I Public CT collections at the lumbosacral junction. Black, class present; outlined, graded only to exclude it⁷.

Collection	Scans	Count	L6	Sacrum	S1	Pelvis	Ribs	L. rib	Femora	Grade
CTSpine1K ³	1,005									
CTPelvic1K ⁴	1,184									
VerSe ⁶	374									
RibSeg v2 ¹⁶	660									
TotalSegmentator ²⁰	1,204									
CTSpinoPelvic1K	802									

67 ground truth. Alongside them, every released measurement is also expressed against the
68 two landmarks present in essentially every abdominal scan, the sacrum below and the lowest
69 rib-bearing vertebra above, so that the measurement is positional rather than nominal: a
70 vertebra by how many rib-free vertebrae separate it from the sacrum, a lowest rib by its
71 length as a fraction of the rib above it, a transverse process by its span and its distance from
72 the ala. None of these changes according to whether a reader would have said L5 or S1, so
73 cases that disagree about the name remain comparable, and where the junction is in view
74 the count is determined and the two descriptions agree.

75 *But the vertebrae themselves are known to differ.* Counting is not the only route to an
76 identity. Level-specific morphometry is long established: pedicle width and height change
77 systematically from the thoracic to the lumbar spine,¹¹ vertebral body, endplate and canal
78 dimensions differ by level,^{12–14} and the transverse process and the iliolumbar ligament mark
79 the last lumbar vertebra independently of any count.^{2,15} If those differences hold at the
80 boundary itself, the phenotype is decidable from the vertebra in front of the reader without
81 any global count. On this corpus it does hold: separating T12 from L1 on shape alone, with
82 each patient’s own size divided out and cross-validation grouped by case, gives an area under
83 the curve of 0.990 and 97.3% accuracy over 1485 vertebrae. Size is removed because raw
84 millimeters separate thoracic from lumbar trivially and teach nothing about the junction,
85 where the two are nearly the same size.

86 That is reported as a property of the released measurements rather than as a method:
87 the information needed to settle these variants is present locally.




FIG. 1 How the dataset was built. n counts scans, or corrections, at each step.

II. ACQUISITION AND VALIDATION METHODS

II.A. Sources, and why a crosswalk was necessary

Both CTSpine1K³ and CTPelvic1K⁴ draw part of their imaging from the TCIA CT colonography collection^{17,18}, and the vertebral and pelvic annotations in both were produced under board-certified radiologist supervision. Neither released the mapping from an annotation to the specific CT series it was drawn on.

That mapping matters because every patient in CT COLONOGRAPHY was scanned *twice*, prone and supine, and an annotation carries a patient identifier and nothing finer.

96 Turning a patient from supine to prone alters lumbar lordosis and the alignment of each
97 segment against the next, so outlines drawn on one series are the wrong *shape* for the other
98 and no rigid realignment recovers them. Such a mask looks competent in isolation and
99 reveals itself only overlaid on the image. The crosswalk (Fig. 1) therefore resolves each
100 annotation to a patient and then, within that patient’s volumes, to the series whose bone
101 *agrees* with it: candidates are thresholded at bone attenuation and scored by overlap with
102 the mask.

103 II.B. Fused and split records

104 CTSpine1K and CTPelvic1K each annotated a patient independently, so their labels are
105 not guaranteed to come from the same CT series (Fig. 1). This turned out to matter for
106 351 patients, nearly half the cohort: the two collections had labeled different acquisitions of
107 the same person, one prone and one supine, with neither recording which. Because posture
108 changes spinal alignment, a spine label from one acquisition cannot simply be paired with a
109 pelvic label from the other. These 351 records are released but held out of any analysis that
110 assumes a single posture. Two labeled acquisitions of the same patient in different postures
111 are, instead, a resource for studying postural alignment, not a defect.

112 II.C. Densification, and why it is asymmetric

113 A corpus in which half the records lack a pelvis is not a spinopelvic dataset, so the missing
114 structures were pseudolabeled, asymmetrically (Fig. 1). Pelvis onto spine-only records is the
115 safe direction: the sacrum and innominate bones carry no enumeration to get wrong, and
116 quality is reported as held-out performance on the *pelvic_only* records, withheld from training
117 for that purpose. The reverse direction carries the whole problem this dataset addresses,
118 since a pseudolabeled spine must commit to a count and on a transitional vertebra commits
119 to the commoner reading; those 20 records were corrected by hand, tractable at that scale.

120 II.D. Ribs

121 No public rib annotation covers this cohort, and the two available automatic sources fail
122 in complementary ways.

123 Möller’s binary rib network¹⁹ reports high accuracy, and its authors show that stump ribs
124 can be classified from a partial field of view; that result concerns the morphological features
125 computed on a rib once segmented. On this corpus the segmentation itself failed on ribs
126 that were *partially outside* the field of view: a rib crossing the boundary of an abdominal
127 acquisition was liable to be missed altogether rather than segmented to the edge. This is
128 not a limitation reported by its authors: their evaluation excluded subjects in which the last
129 rib was not visible, and their failure analysis attributes errors to costal-process fusion, so it
130 is recorded here as an observation from this cohort. Since this is a colonography cohort, the
131 lower ribs are routinely clipped, and those are precisely the ribs the count depends on.

132 TotalSegmentator²⁰ does segment clipped ribs and supplies the per-rib numbering that a
133 binary network cannot. Its characteristic failure is at the other end of the bone: the most
134 medial portion of the rib (roughly the last tenth to sixth approaching the costovertebral
135 joint) is frequently absent. A rib truncated at its medial end never reaches the vertebra
136 it articulates with, so the rib–vertebra incidence check (the quality-control gate on the rib
137 class labels, which matches each rib to the vertebra its number implies) has nothing to test.

138 The two were therefore unioned (Fig. 1): TotalSegmentator supplies numbering and the
139 medial-clipped cases, Möller supplies detection where its output is more complete, and the
140 union carries TotalSegmentator’s per-level identities.

141 The result is a pseudolabel whose review was triaged rather than exhaustive, which
142 is worth stating plainly because a layer described only as “human-corrected” invites the
143 reader to assume every bone was looked at. A label-based rule flagged every record in
144 which one rib bone carried more than one identifier, or in which a rib failed to reach the
145 vertebra it articulates with, since a single bone carrying two numbers is a numbering error by
146 construction, whatever the anatomy. The rule selected 152 records; 148 were reviewed and
147 corrected, and 4 were not and are identified in the release. Reviewers were medical students
148 working through a student research consortium²¹, trained against a written labeling protocol,
149 with every correction recorded against the annotator who made it. Records passing the rule
150 were not individually inspected, so the layer only guarantees freedom from the failure modes
151 the rule detects, no more. The residual rib–vertebra offsets reported in Sec. II.G are the
152 independent check on what the rule missed.

153 II.E. Label scheme and counting anchors

154 The scheme is VerSe-native: vertebrae keep their VerSe identifiers (C1–C7 = 1–7, T1–T12
155 = 8–19, L1–L6 = 20–25, sacrum 26), and every non-VerSe structure takes a fixed identifier
156 above that range: S1 carved from the sacrum (29), hips (30–31), femora (32–33), ribs (34–45
157 left, 46–57 right), lumbar ribs (74–75) and surgical hardware (76–82). Identifiers 58–73 are
158 unassigned: a soft-tissue block once reserved there was never populated and is retired.

159 Two anchors make the count decidable, and Fig. 2 shows both: the lowest rib-bearing
160 vertebra above, **S1** carved from the sacrum below. Each earns its place by what it supports
161 rather than by being an endpoint. S1 supplies the sacral endplate landmark that sacral
162 slope and pelvic incidence are measured from, and with both brackets present L5 and L6
163 are distinguishable where a spine-limited view alone would not be.

164 *How S1 is cut.* The sacrum’s outer boundary is the source annotation and is never altered;
165 S1 is the cranial part of that same bone, separated by a plane. The plane’s normal is the
166 sacrum’s own cranio-caudal axis, made orthogonal to a left–right axis measured directly
167 rather than assumed, so as to remove the patient’s rotation within the scanner. That
168 rotation is not negligible: median 4.5°, maximum 16.1°, and 508 of 801 records above 3°.
169 The level of the cut preserves the sub-division volume of the preceding release, so the plane
170 changes the shape and orientation of the S1/S2 boundary and not how much of the sacrum
171 is called S1. Per-record geometry ships with the archive.

172 Neither anchor is novel on its own: TotalSegmentator²⁰ carries a separate S1 class and
173 per-level ribs, and VerSe⁶ carries L6, whose identifier this release keeps, so its L6 *is* VerSe’s.
174 L6 is retained because a scheme without it cannot record a six-lumbar spine at all.

175 **A rib on a lumbar body receives its own class**, and this is the one class here with no
176 published counterpart. A lumbar rib and a stump rib are the same object under two counts,
177 so assigning it a number would commit the label to one reading of an enumeration anomaly.
178 A scheme that numbers every rib 1–12 has nowhere to put a thirteenth: the annotator must
179 either call it rib 12, asserting the vertebra beneath it is thoracic, the very question at issue,
180 or discard it. TotalSegmentator carries ribs 1–12 per side and no such class, and VerSe’s
181 T13 is a vertebra rather than a rib. This scheme can record both: VerSe’s T13 (28) for a
182 thoracic-type thirteenth vertebra with a true rib, and 74/75 for a lumbar-type body bearing
183 a rudimentary one, so either reading is recorded as itself. No record here carries a T13; the




FIG. 2 The two anchors, from the released labels. Red, lowest rib-bearing vertebra and its rib; blue, S1; yellow, the rib-free vertebrae between. (b) and (c) both count four.

184 sixteen lumbar-rib cases were read as lumbar-type bodies. Where a border vertebra could
 185 not be settled by its readers the label stands as read, and reclassifying those records with
 186 a classifier trained on local vertebral morphology (Sec. V) is the intended path for a later
 187 release.

188 Figure 2 shows why the interval is the primitive and the label is derived. Among the
 189 records with four rib-free vertebrae, some reach that count through a lumbar rib and some
 190 without one; the two are indistinguishable by the count alone and are distinguished here
 191 only because rib presence is itself labeled.

192 Sixteen records carry a lumbar rib: thirteen bilateral and three unilateral, all on the
 193 right, a count too small to support any statement about side preference.

194 Eighteen records carry an L6. Fourteen are labeled lumbarization, two sacralization,
 195 one semi-sacralization and one normal, which is worth stating because it shows the two
 196 annotation axes disagreeing in the open: a record can carry six lumbar-type vertebrae and
 197 still be read as a sacralization, or as unremarkable, depending on where the reader started
 198 the count.

199 Both counts are derived from the released label volumes, and so are the manifest fields
 200 that report them. That is a correction: through v5, `has_l6` was true for one record that
 201 contains no L6 and false for all eighteen that do, `has_lumbar_rib` was false everywhere, and
 202 `n_lumbar_labels` read zero in 799 of 802 records. The fields are rebuilt for this release by
 203 counting identifiers 25 and 74–75 in each volume, with `lumbar_rib_side` added alongside
 204 them.

205 II.F. Computational tools

206 *Access.* Every processing step described above is performed by code released under an
207 open-source licence and archived with the dataset; no step depends on software that is
208 not publicly obtainable. The release archive carries the reference loader, the label-scheme
209 definition and the quality-control scripts that generated the tables in this manuscript.

210 *Versions and parameters.* Segmentation of femora, the sacral sub-division and the per-
211 level rib numbering used TotalSegmentator²⁰ (`total` task, release 2.x) on a single graphics
212 processing unit. The binary rib network is Möller’s¹⁹ released weights, applied through the
213 nnU-Net v2 framework²². The pelvic completion for records whose pelvis was absent used
214 a five-fold nnU-Net v2 ensemble, applied out-of-fold: each record is completed by the fold
215 that never trained on it, so no record is completed by a model that has seen it. The five
216 fold checkpoints, with their nnU-Net plans and dataset descriptor, are released alongside the
217 data (<https://huggingface.co/anonymous-mlhc/spinopelvic-seg-checkpoints>). Im-
218 age handling throughout used NiBabel and SciPy; no image is resampled or reoriented when
219 a label is written, so every label shares its image’s grid and affine exactly.

220 *Generative artificial intelligence.* Large language models assisted with drafting and edit-
221 ing text and with analysis and figure code, under author review. They were not used to
222 generate, annotate or interpret imaging data; every number reported here is computed by
223 the released code from the released volumes.

224 II.G. Validation

225 Validation is layered, and the ordering matters: measurements computed on a corpus
226 that has not established internal consistency do not look broken, they look finished.

227 *Geometric invariants.* Every case is checked for label geometry agreeing with its CT in
228 shape, affine and voxel spacing; for identifiers within the published scheme; for a non-empty
229 label; and for sided structures falling on the correct sides, with the left–right axis read from
230 the affine rather than assumed. The check exists because a renumbering pass once wrote
231 a reoriented array under a canonical affine and transposed a label away from its image, a
232 failure invisible downstream.

233 The check compares *each sided pair separately*: testing the ribs alone passed all 802

records while four carried `left_hip` on the patient’s right, and pooling all sided structures into one test still missed one of those four, because a large correctly-sided structure masks a smaller transposed one. A gate that passes everything is not evidence of a clean corpus until it has been shown capable of failing.

Rib-vertebra incidence. Each rib is matched to its nearest vertebra and compared against the vertebra its number implies. Of 11,548 ribs, 5,788 are not evaluable because the expected vertebra lies outside the field of view and 11 have no vertebra within the anchor distance; of the 5,749 evaluable, 5,747 match, two (0.035%) are offset by one level, and no case is misnumbered. The denominator is stated because quoting the two offsets against all 11,548 ribs would halve the rate by counting ribs the check never examined. Both offsets are field-of-view truncations on individually characterized cases, one at +1 level and one at −1, reported rather than suppressed.

The same check carries an invariant that reads as a null result. Of the 5,749 evaluable ribs, *zero* are assigned to a lumbar vertebra, not because no lumbar ribs exist, since 16 records carry one, but because such a rib takes class 74 or 75 and never enters the numbered set the test examines. A numbered rib landing on a lumbar vertebra would be a counting error of exactly the kind this dataset exists to expose.

Agreement with published values. Derived spinopelvic measures are compared against Vialle *et al.*²³ (Table II). Pelvic incidence is the strongest of these checks because it is a morphological property of the pelvis rather than a posture, and so is directly comparable to a standing reference cohort.

II.H. Castellvi typing

Every record whose vertebral count departs from five carries a Castellvi grade²⁴ in the released manifest: 33 cases, typed I–IV with the *a/b* unilateral–bilateral qualifier. This is released as an annotation layer because the grade and the count are *different axes* and the dataset is built to show that they are: grade IIb occurs here at rib-free counts of four, five and six alike, and seven of the 33 graded cases carry a perfectly normal count of five. A corpus that recorded only the count would describe those seven as unremarkable.

Two radiology resident physicians graded every case independently and resolved their six disagreements by consensus; both reads and the consensus are in the manifest. The

TABLE II Derived measures against a standing reference ($n = 260$)²³; this cohort is supine.

Measure	This dataset (supine)	Reference (standing)
Pelvic incidence	54.7°	54.7 ± 10.6
Pelvic tilt	17.3°	13 ± 6
Sacral slope	36.9°	41.0 ± 8.4
Lumbar lordosis (supine)	52.7°	43 ± 11.2 (standing)
PI – LL mismatch	2.6°	centered near 0

distinction the classification turns on, bony fusion against articulation without it², is the one carrying clinical weight, and it is where the disagreements lay: four of the six were III against II.

II.I. Surgical hardware, and why a threshold is not enough

Metal is trivial to detect and hard to interpret, and the difference matters here because **an iatrogenic fusion is indistinguishable from a congenital one to a distance measurement**: a cage-bridged interspace reads as “no gap” exactly as a fused transitional vertebra does. An unrecorded instrumented case is therefore a false positive on the central measure.

II.I.0.a. Detection over-calls by a factor of five. A 2500 HU threshold inside a shell around the labeled skeleton flags 52 of the 802 records. One of the two radiology resident physicians read all 52 and confirmed 11, rejecting 41 (79%) as contrast, calcification and reconstruction artefact. Attenuation does not separate the two groups (both saturate); site and volume do, every confirmed implant at 2,586 mm³ or above and every rejection at 1,768 mm³ or below. Instrumentation prevalence is 1.4%, not the 6.5% the threshold alone reports.

II.I.0.b. The subtype block had to be extended. Identifiers 76–79 name spinal instrumentation (generic, cage, screw-and-rod, plate) and cannot describe what this cohort holds, so **80 arthroplasty**, **81 sacroiliac screw** and **82 osteosynthesis** were added: osteosynthesis holds parts of one bone together where arthroplasty replaces a joint, and a shape rule calls a femoral stem a rod. The confirmed cases are 8 arthroplasties, one osteosynthesis, one sacroiliac fixation and one pair of interbody cages (Fig. 3).




FIG. 3 Instrumentation in the release, one exemplar each, drawn through bone. (a) Hip arthroplasty, 80, $n = 8$. (b) Femoral-neck screws, 82, $n = 1$. (c) Sacroiliac screws, 81, $n = 1$. (d) Interbody cages, 77, $n = 1$. Bar, 5 cm.

II.I.0.c. Metal outranks bone, and that is a subtraction. In every confirmed case the implant was already inside a bone label, so naming it reclaims 1,538,852 voxels rather than adding them. Pelvic incidence and tilt are measured from the femoral head center, and in 8 of these cases that head is an implant: any spinopelvic parameter previously computed from them was measured on metal.

III. DATA FORMAT AND USAGE NOTES

Each record is a gzipped NifTI image/label pair sharing an identical 4×4 affine, canonicalized to PIR, requiring no resampling. Masks are NifTI rather than DICOM: the policy prefers DICOM where applicable, and for volumetric segmentation masks NifTI is what the downstream tooling expects.

Splits are patient-grouped and LSTV-stratified five-fold, frozen and shipped with the data as `splits_5fold.json`. The five validation folds partition all 802 records, so the release provides cross-validation and *not* a held-out test set; users reporting a single held-out number should carve one out and state which records it holds.

Stratification is on the transitional subtype rather than a binary LSTV flag: every validation fold carries three sacralization and three lumbarization records. With 15 of each across 802 that is the most stratification can guarantee, and it is why a fold-level metric on the rare classes carries an error bar far wider than its own decimal places.

Table III lists the key data types and metadata against the fraction of records for which each is populated.

TABLE III Data types and metadata, and the records populating each, from the released manifest.

Field	Records	Available
Identifier, image path, label path, configuration, match type	802	100.0%
Patient position, tube potential, section thickness, kernel	802	100.0%
Transitional label and class	802	100.0%
Spine and pelvic annotation origin	802	100.0%
Instrumentation and partial-annotation flags	802	100.0%
Image/label alignment check	802	100.0%
Spine series identifier; bone-fraction check	782	97.5%
Scanner manufacturer and model	768	95.8%
Sex	749	93.4%
Age	709	88.4%
Pelvic series identifier	713	88.9%
Castellvi grade (consensus of two readers)	33	4.1%

306 The archive of record is Zenodo (<https://doi.org/10.5281/zenodo.22139642>). Code
307 and documentation are on GitHub at the tagged release commit; any model-hub copy is a
308 convenience mirror, not the archive.

309 IV. DESCRIPTIVE ANALYSIS

310 The cohort is a colorectal cancer screening population of 802 records. 738 are recorded
311 female (393) or male (345); 11 carry DICOM’s *other* value, a recorded value and not a
312 missing one, and 53 carry no sex field. Age is present for 709 (median 59, range 50–89).
313 All 802 are counted in totals; the eleven records outside the female/male dichotomy are
314 excluded from sex-stratified measures rather than folded into “unrecorded”. The lower age
315 bound is the screening guideline, not a selection applied here.

316 *Count-free measures* (Fig. 4). Five rib-free vertebrae above the sacrum is typical; four
317 and six are where transitional anatomy sits, and which of the two is present does not by
318 itself determine what to call it. The lowest-rib length ratio is bimodal.

319 *Level gradients and spinopelvic measures* (Fig. 5). Vertebral dimensions increase caudally




FIG. 4 Count-free measures. (a) Vertebrae between the lowest rib and the sacrum. (b) Lowest rib length over the rib above. (c) Lowest lumbar transverse span against its gap to the ala; transitional labels marked.




FIG. 5 Spinopelvic measures against standing reference bands (mean $\pm$ SD)²³.

in step with the load each level carries. Pelvic incidence does not vary with age, while pelvic tilt increases, the descriptive signature of a pelvis retroverting as a spine ages.

Opportunistic measures. Every scan carries a vertebral bone density measurement at no additional dose.²⁵ L1 trabecular attenuation is reported at the published standard site.

V. POTENTIAL APPLICATIONS AND LIMITATIONS

The principal use is the *spinopelvic interface*. Spine and pelvis carry radiologist-sourced annotation on one coordinate frame in 802 records, making sacral morphology measurable against the lumbar column rather than in isolation: the sacral endplate and its slope, the alar corridors constraining S2-alar-iliac and iliac screw trajectories, and the L5–S1 geometry a lumbosacral fusion must cross. Pelvic incidence is derived from the segmented sacrum and




FIG. 6 Morphometry by level: median, interquartile range, 5th–95th percentile, n at right. (a) Body height. (b) Endplate width. (c) Canal. (d) Pedicle width. (e) Disc height. (f) Trabecular attenuation.

330 femoral heads rather than from landmarks placed on a radiograph, so it is reproducible from
 331 the labels themselves.

332 The transitional layer serves that use rather than competing with it: no trajectory can be
 333 planned across a junction whose segments cannot be named. Level-numbering benchmarks,
 334 variant studies and opportunistic screening follow from the same annotation.

335 V.A. Three uses the released measurements support

336 *Naming a level from its own geometry.* Wrong-level spine surgery runs at roughly one
 337 in 3,110 spinal procedures, and the cause most often cited is the transitional anatomy this
 338 cohort was assembled around.²⁶ The failure is structural rather than careless: the count
 339 is ambiguous exactly where the variant sits. Per-level body height, canal width, endplate
 340 width, transverse-process span and wedge ratio are released for T11–L5 across all 802 records,
 341 and because the labels were assigned against the twelfth rib rather than by enumeration, a
 342 model trained on them is not learning the convention that fails.




FIG. 7 Records carrying each vertebral level, cranial to caudal. Counts per identifier in Table S1.

343 *Reference morphometry with its spread.* The values a surgeon consults for vertebral
 344 body height, canal dimensions and pedicle width come from cadaveric series of a few dozen
 345 specimens, plotted as single values with no indication of spread.¹⁴ They cannot tell a reader
 346 whether the patient in front of them is unusual, because they never show what usual looks
 347 like. The same measures are given here from 802 records with the distribution attached
 348 (Fig. 6), reproducing two properties of the classical figures (body height crosses over, dorsal
 349 exceeding ventral at T11–T12 and reversing by L4–L5, and pedicle width widens steeply
 350 into the lower lumbar spine) and adding the width of each distribution, which is the part
 351 needed to judge an individual. A larger n sharpens a distribution without changing the
 352 population it came from, and this one is a screening cohort aged 50 and over.

353 *Patient-specific models.* Planning is moving toward patient-specific biomechanical models
 354 rather than population norms.²⁷ Such a model needs spine, pelvis and femoral heads in
 355 one frame, which is this release’s construction. In 351 patients the annotation sits on two
 356 acquisitions in different positions, giving a within-patient change in alignment against which
 357 a predicted postural response can be checked.

358 Its limitations are stated at the level of detail a reader would need to catch them inde-
 359 pendently.

360 *Coverage.* Thoracic ground truth is field-of-view limited (Fig. 7) and does not extend to
361 T1. Postural angles are supine; pelvic incidence is not postural and needs no such caveat.
362 The cohort is a colorectal screening population aged 50 and over, so its distributions should
363 not be read as representative of a surgical one.

364 *Declared but empty, and confirmed so.* The partial-annotation sentinel (255) is declared
365 in the scheme and occurs in no record, established by counting identifiers across all 802
366 released volumes rather than asserted. It is retained because the protocol is part of the
367 scheme and a later release may ship a partially traced record. The hardware identifiers were
368 in this position until v6 and are no longer: 76–82 are populated in the eleven instrumented
369 records described in Sec. II.I.

370 *Strength of the labels varies by structure, and the release does not average over that.*
371 Vertebral labels derive from radiologist-supervised source annotations, corrected where those
372 sources were wrong. Pelvic labels on records that lacked one are pseudolabeled. The rib
373 layer is a pseudolabel whose human review was triaged by an automated rule rather than
374 exhaustive, so its guarantee is the absence of the failure modes that rule detects. Transitional
375 labels derive from two independent sources and are not uniformly adjudicated, which is
376 precisely why the measures reported here are count-free, and the Castellvi grades are a
377 consensus of two radiology resident physicians. The sacral sub-division inherits its level
378 from an automatic method, and in 100 records that level places S1 above half the sacrum
379 or below fifteen percent of it, which no first sacral segment is; those records are flagged in
380 the released quality-control table and should be filtered before any measurement depending
381 on the sacral endplate.

382 *Structures are not universally present.* The same census shows the release is not uniformly
383 complete. Sacrum, hips and femora are present in all 802 records. One record carries no S1,
384 one lacks T12, and nine lack an L5 identifier, in every case because the structure lies outside
385 the field of view or, for L5, because the lowest lumbar segment is labeled L6 or incorporated
386 into the sacrum. Users filtering on the presence of the caudal anchor should filter explicitly
387 rather than assume it.

388 VI. DISCUSSION

389 *Comparison with related material.* Against the collections in Table I, the contribution
390 is the crosswalk placing spine and pelvis on the same series, together with classes for the
391 anatomy an enumeration anomaly actually produces. VerSe⁶ comes closest in intent: it
392 also enriched its cohort for transitional anatomy and graded each case by Castellvi. The
393 difference is what a graded case leaves behind. VerSe names the anomaly but does not
394 segment the sacrum, so the fused or articulating structure the grade describes has no mask
395 to go with it. This release segments that structure directly, as L6 or as a separately carved
396 S1, so a user has the anatomy itself to measure rather than only a grade that describes it.
397 TotalSegmentator²⁰ invites a different comparison, since it segments far more of the body
398 than any single source used here. But its label scheme simply has no class for a sixth lumbar
399 vertebra or for a rib on a lumbar body. A six-lumbar spine cannot be recorded in it: the
400 column must be renumbered or a level dropped to fit. What this release adds, then, is not
401 more anatomy per scan (TotalSegmentator already covers more), but the specific classes an
402 enumeration anomaly needs in order to be recorded accurately.

403 *Limitations.* The cohort comes from one acquisition protocol of a colorectal-screening
404 population aged 50 and over, so how well it generalizes to a pathologic population is untested.
405 The rib layer’s human review was triaged by an automated rule rather than performed on
406 every bone, so it guarantees only the absence of the failure modes that rule detects.

407 *Future directions.* The version history shows the scheme being extended more than once:
408 femora and the S1 anchor, then per-level ribs, then lumbar ribs and surgical hardware, were
409 each folded into the existing release rather than requiring a new dataset. The same path is
410 open going forward: imaging from beyond this cohort’s single protocol, and the C2 anchor
411 this acquisition cannot supply, are the natural next additions.

412 VII. CONCLUSION

413 CTSpinoPelvic1K places spine, pelvis, per-level ribs and femora on one coordinate frame
414 in 802 records, and gives the anatomy that makes lumbar numbering ambiguous a class of
415 its own rather than recording it as something it is not. Because both counting anchors are
416 explicit, a level can be identified from a field of view that cannot support the conventional

417 count downward from C2, which is every record here.

418 DATA AVAILABILITY

419 The release described here is v8, permanently archived at [https://doi.org/10.5281/](https://doi.org/10.5281/zenodo.22293878)
420 [zenodo.22293878](https://doi.org/10.5281/zenodo.22293878); the concept DOI <https://doi.org/10.5281/zenodo.22139642> resolves
421 to the current version. The archive carries the loader, the label scheme, the quality-control
422 tables and the analysis code under an open-source licence. The repository URL identifies
423 the authors; it is on the title page and will be restored here for the camera-ready version.

424 ACKNOWLEDGEMENTS

425 Acknowledgements are given on the title page, per the journal’s double-anonymized re-
426 view policy. This work received no funding.

427 CONFLICT OF INTEREST

428 The authors have no relevant conflicts of interest to disclose.

429 ETHICS

430 This work uses publicly available, de-identified imaging and annotations derived from it.
431 The authors’ institutional review board determined that it is not human participant research
432 under 45 CFR 46 and requires no oversight.

433 REFERENCES

- 434 ¹J. Lian, N. Levine, and W. Cho, A review of lumbosacral transitional vertebrae and
435 associated vertebral numeration, *Eur. Spine J.* **27**, 995–1004 (2018), doi:10.1007/s00586-
436 018-5554-8.
- 437 ²G. P. Konin and D. M. Walz, Lumbosacral transitional vertebrae: classification, imag-
438 ing findings, and clinical relevance, *AJNR Am. J. Neuroradiol.* **31**, 1778–1786 (2010),
439 doi:10.3174/ajnr.A2036.

- ³Y. Deng et al., CTSpine1K: a large-scale dataset for spinal vertebrae segmentation in computed tomography, *Mach. Learn. Biomed. Imaging* **3**, 824–832 (2025), doi:10.59275/j.melba.2025-gf84.
- ⁴P. Liu et al., Deep learning to segment pelvic bones: large-scale CT datasets and baseline models, *Int. J. Comput. Assist. Radiol. Surg.* **16**, 749–756 (2021), doi:10.1007/s11548-021-02363-8.
- ⁵H. Möller, H. Schoen, R. Graf, et al., VERIDAH: solving enumeration anomaly aware vertebra labeling across imaging sequences, 2026, arXiv:2601.14066. doi:10.48550/arXiv.2601.14066.
- ⁶A. Sekuboyina et al., VerSe: a vertebrae labelling and segmentation benchmark for multi-detector CT images, *Med. Image Anal.* **73**, 102166 (2021), doi:10.1016/j.media.2021.102166.
- ⁷H. Liebl et al., A computed tomography vertebral segmentation dataset with anatomical variations and multi-vendor scanner data, *Sci. Data* **8**, 284 (2021), doi:10.1038/s41597-021-01060-0.
- ⁸Y. Otake, S. Schafer, J. W. Stayman, et al., Automatic localization of vertebral levels in x-ray fluoroscopy using 3D-2D registration: a tool to reduce wrong-site surgery, *Phys. Med. Biol.* **57**, 5485–5508 (2012), doi:10.1088/0031-9155/57/17/5485.
- ⁹S.-F. L. Lo, Y. Otake, V. Puvanesarajah, et al., Automatic localization of target vertebrae in spine surgery: clinical evaluation of the LevelCheck registration algorithm, *Spine (Phila. Pa. 1976)* **40**, E476–E483 (2015), doi:10.1097/BRS.0000000000000814.
- ¹⁰T. De Silva, S.-F. L. Lo, N. Aygun, et al., Utility of the LevelCheck algorithm for decision support in vertebral localization, *Spine (Phila. Pa. 1976)* **41**, E1249–E1256 (2016), doi:10.1097/BRS.0000000000001589.
- ¹¹M. R. Zindrick et al., Analysis of the morphometric characteristics of the thoracic and lumbar pedicles, *Spine (Phila. Pa. 1976)* **12**, 160–166 (1987), doi:10.1097/00007632-198703000-00012.
- ¹²M. M. Panjabi et al., Thoracic human vertebrae: quantitative three-dimensional anatomy, *Spine (Phila. Pa. 1976)* **16**, 888–901 (1991), doi:10.1097/00007632-199108000-00006.
- ¹³M. M. Panjabi et al., Human lumbar vertebrae: quantitative three-dimensional anatomy, *Spine (Phila. Pa. 1976)* **17**, 299–306 (1992), doi:10.1097/00007632-199203000-00010.
- ¹⁴E. C. Benzal, *Biomechanics of Spine Stabilization*, Thieme, New York, 3rd edition, 2015.

- ¹⁵R. J. Hughes and A. Saifuddin, Numbering of lumbosacral transitional vertebrae on MRI: role of the iliolumbar ligaments, *AJR Am. J. Roentgenol.* **187**, W59–W65 (2006), doi:10.2214/AJR.05.0415.
- ¹⁶L. Jin et al., RibSeg v2: a large-scale benchmark for rib labeling and anatomical centerline extraction, *IEEE Trans. Med. Imaging* **43**, 570–581 (2024), doi:10.1109/TMI.2023.3313627.
- ¹⁷K. Smith et al., Data from CT COLONOGRAPHY, The Cancer Imaging Archive, 2015, doi:10.7937/K9/TCIA.2015.NWTESAY1.
- ¹⁸K. Clark et al., The Cancer Imaging Archive (TCIA): maintaining and operating a public information repository, *J. Digit. Imaging* **26**, 1045–1057 (2013), doi:10.1007/s10278-013-9622-7.
- ¹⁹H. Möller et al., Automated thoracolumbar stump rib detection and analysis in a large CT cohort, 2025, arXiv:2505.05004. doi:10.48550/arXiv.2505.05004.
- ²⁰J. Wasserthal et al., TotalSegmentator: robust segmentation of 104 anatomic structures in CT images, *Radiol. Artif. Intell.* **5**, e230024 (2023), doi:10.1148/ryai.230024.
- ²¹[Reference to the authors’ own work, removed for double-anonymized review.]
- ²²F. Isensee, P. F. Jaeger, S. A. A. Kohl, J. Petersen, and K. H. Maier-Hein, nnU-Net: a self-configuring method for deep learning-based biomedical image segmentation, *Nat. Methods* **18**, 203–211 (2021), doi:10.1038/s41592-020-01008-z.
- ²³R. Vialle, N. Levassor, L. Rillardon, A. Templier, W. Skalli, and P. Guigui, Radiographic analysis of the sagittal alignment and balance of the spine in asymptomatic subjects, *J. Bone Joint Surg. Am.* **87**, 260–267 (2005), doi:10.2106/JBJS.D.02043.
- ²⁴A. E. Castellvi, L. A. Goldstein, and D. P. K. Chan, Lumbosacral transitional vertebrae and their relationship with lumbar extradural defects, *Spine (Phila. Pa. 1976)* **9**, 493–495 (1984), doi:10.1097/00007632-198407000-00014.
- ²⁵P. J. Pickhardt, B. D. Pooler, T. Lauder, A. M. del Rio, R. J. Bruce, and N. Binkley, Opportunistic screening for osteoporosis using abdominal computed tomography scans obtained for other indications, *Ann. Intern. Med.* **158**, 588–595 (2013), doi:10.7326/0003-4819-158-8-201304160-00003.
- ²⁶N. E. Epstein, A perspective on wrong level, wrong side, and wrong site spine surgery, *Surg. Neurol. Int.* **12**, 286 (2021), doi:10.25259/SNI.402_2021.

⁵⁰³ ²⁷E. Beaulieu et al., A review of finite element analysis in spine surgery decision-making, J.
⁵⁰⁴ Clin. Med. **15**, 2584 (2026), doi:10.3390/jcm15072584.