

**CTSpinoPelvic1K: spine, pelvis, ribs and femora in one coordinate frame, annotated
for lumbosacral transitional anatomy**

[Author names removed for double-anonymized review]

[Affiliations removed for double-anonymized review]

(Dated: 6 September 2026)

Purpose: A vertebra at the lumbosacral junction is named by counting. The gold standard counts caudally from C2 on whole-spine imaging¹; the imaging a lumbar case is planned on is regional², and abdominal computed tomography (CT) never contains C2. Where a lumbosacral transitional vertebra (LSTV) alters the count, the local anatomy is ambiguous: four rib-free vertebrae may be an L1 bearing a lumbar rib or an L5 assimilated to the sacrum, and six may be a sixth lumbar vertebra, a T12 whose ribs are aplastic, or an S1 lumbarized from the sacrum³. CTSpinoPelvic1K exists to ask whether local morphology resolves that ambiguity without the count. It provides 802 CT records with radiologist-sourced spine and pelvic annotation on one frame, per-level ribs and femora, each lumbar level anchored on the lowest rib-bearing vertebra and the first sacral segment, and classes for a sixth lumbar vertebra, a thirteenth thoracic vertebra, a separate first sacral segment and lumbar ribs, so that the anomalies are recorded as themselves.

Acquisition and Validation Methods: Records pair TCIA CT COLONOGRAPHY imaging with CTSpine1K⁴ vertebral and CTPelvic1K⁵ pelvic labels on the acquisition with highest bone coverage, under a VerSe-native scheme. Validation: geometric invariants (802/802 pass); rib-vertebra incidence across 5,749 evaluable ribs (two offsets, 0.035%); and agreement of spinopelvic measures with published values.

Data Format and Usage Notes: Gzipped NIfTI image/label pairs sharing one affine, canonicalized to PIR, with frozen patient-grouped LSTV-stratified five-fold splits and a reference loader; archived at <https://doi.org/10.5281/zenodo.22139642>.

Potential Applications: Classifying a vertebra from local features alone; updating textbook morphometry from small cadaveric series; spinopelvic assessment; variant studies; opportunistic screening. *Limitations:* thoracic ground truth is field-of-view limited; postural angles supine; no held-out test set; the rib layer is a triaged-review pseudolabel; Castellvi grades are a two-reader consensus.

6 I. INTRODUCTION

7 The spine commonly carries seven cervical, twelve thoracic and five lumbar vertebrae, and
8 a vertebra is identified by where it falls in that sequence. Transitional variants sit at the
9 two ends of the lumbar column, the thoracolumbar transitional vertebra (TLTV) above and
10 the lumbosacral transitional vertebra (LSTV) below, and they disturb identification in two
11 ways at once. They change what the border vertebra and its ribs *look like*: a thoracolumbar
12 border vertebra may carry a full rib, a hypoplastic stump, or none at all, and a lumbosacral
13 one may be partly or completely assimilated to the sacrum (*sacralization*), with an anteriorly
14 wedged body, a reduced disc beneath it and hypoplastic facet joints, or conversely separated
15 from it (*lumbarization*), with a squared body, lumbar-type facets and a full-height disc³.
16 They also change *how many* vertebrae each region contains: eleven or thirteen thoracic,
17 four or six lumbar.

18 Either change alone would be tractable. Together they make it difficult to state defini-
19 tively what any vertebra at these borders *is*. Four rib-free vertebrae above the sacrum may
20 mean an L1 bearing a lumbar rib or an L5 assimilated to it; six may mean a true sixth
21 lumbar vertebra, a T12 whose ribs are aplastic, or an S1 lumbarized away from the sacrum.
22 Each set yields the same count, and the morphology that would separate the readings is not
23 decisive to a human reader in every case. Whether it differs enough for a learned classifier
24 to separate them, or at least to assign each reading a probability, is the question this release
25 is built to support. Absent that, what remains is the count itself: the gold standard for
26 numeration in transitional anatomy is whole-spine imaging, counting caudally from C2¹,
27 which fixes how many vertebrae and how many rib pairs the column holds, and no spine-
28 limited acquisition can supply either count. LSTV is common, reported between 4% and
29 30% depending on definition, and wrong-level surgery is its most serious consequence.

30 Existing public collections cannot address this, and size is not the reason (Table I). Spine
31 datasets omit the pelvis; pelvic datasets do not number the vertebrae; TotalSegmentator,
32 the whole-body scheme, has no class for a sixth lumbar vertebra or for a rib on a lumbar
33 vertebra; and VerSe, the one collection with an L6 class, stops at the sacrum and carries
34 no ribs. No public collection can therefore record a six-lumbar spine together with both
35 of its borders, however good the segmentation, and a segmenter trained on them inherits
36 the gap: faced with an enumeration anomaly it must shift the whole column by a level or

absorb a vertebra into its neighbor. The size of that problem has now been measured on CT, where prior labeling methods assigned every vertebra correctly in 77% of subjects and an anomaly-aware extension of SPINEPS⁶ raised that to 99%⁷. Its weights label L6 and T13 on CT, but its 1,536 training scans are in-house and unreleased: the capability exists as weights, not as data. VerSe^{8,9} makes the point most clearly (Table I): it deliberately enriched for anatomical variants and graded them by Castellvi, yet left the graded vertebrae fused to the sacrum unsegmented.

The intraoperative side of the problem has prior art of its own. LevelCheck registers the intraoperative radiograph to the preoperative CT and projects the CT’s vertebral labels onto it, localizing the target level in a limited field of view to within a few percent of cases in clinical use^{10–12}. The labels it projects are placed by hand on the CT, by a skilled operator and verified by the surgeon¹⁰: the method assumes a correctly labeled CT and does not supply one. Automating that step, and getting it right precisely where the anatomy is anomalous, is what a dataset carrying the classes those anomalies produce is for.

CTSpinoPelvic1K places the spine, pelvis, ribs and femora in one coordinate frame and gives a class to each structure an enumeration anomaly produces: a sixth lumbar vertebra, and a rib borne by a lumbar vertebra. A scheme without those classes lacks a class vocabulary expressive enough to describe anatomic anomalies.

The project began as a surgeon’s question: whether a vertebra can be named from the anatomy in front of the reader, on the images a lumbar case is actually planned on: by guideline those are regional lumbar studies^{2,13}, and whole-spine coverage is reserved for trauma with an identified injury¹⁴ and for deformity radiographs. The cohort suits that question. Every record comes from one prospective trial protocol, ACRIN 6664^{15,16}: 802 patients aged 50 and over, scanned supine and prone on multidetector CT at 15 centers, on five scanner vendors and nine models, with manufacturer, model and reconstruction kernel recorded per record. To our knowledge it is the largest spine-and-pelvis annotated CT cohort acquired under a single protocol; the comparators pool several collections (CTSpine1K four, CTPelvic1K seven), are multi-site by design (VerSe), or are routine clinical scans under many protocols (TotalSegmentator). With the protocol held fixed, scanner effects on a model become measurable rather than confounded. The price is the field of view: abdominal, so the thoracic spine enters from below and C2 is never present (Sec. V).

Where the material is. Everything described here is v8 of a single archived dataset,

deposited at <https://doi.org/10.5281/zenodo.22293878>. Citations should use <https://doi.org/10.5281/zenodo.22139642>, the permanent identifier for the dataset across all of its versions, which resolves to whichever version is current; the version identifier above pins the exact release measured in this manuscript (Sec. III).

I.A. Vertebra labeling when the scan cannot settle the count

The limitation, stated plainly. No scan in this corpus contains C2, so the conventional count cannot be performed: a property of abdominal imaging, not of the annotation. A thirteenth thoracic vertebra, a lumbar rib and a *stump rib* (a hypoplastic twelfth rib, a cardinal indicator of a thoracolumbar transitional vertebra) leave the same number of rib-free vertebrae between the two borders of the lumbar spine, the last rib-bearing vertebra and the sacrum.

What the release does in addition. Every vertebra carries the identifier its radiologist-sourced annotation assigned, corrected where the source was wrong; those labels are the ground truth. Alongside them, every released measurement is also expressed against the two landmarks present in essentially every abdominal scan, the sacrum below and the lowest rib-bearing vertebra above, so that the measurement is positional rather than nominal: a vertebra by how many rib-free vertebrae separate it from the sacrum, a lowest rib by its length as a fraction of the rib above it, a transverse process by its span and its distance from the ala. None of these changes according to whether a reader would have said L5 or S1, so cases that disagree about the name remain comparable, and where the junction is in view the count is determined and the two descriptions agree.

But the vertebrae themselves are known to differ. Counting is not the only route to an identity. Level-specific morphometry is long established: pedicle width and height change systematically from the thoracic to the lumbar spine,¹⁷ vertebral body, endplate and canal dimensions differ by level,^{18–20} and the transverse process and the iliolumbar ligament mark the last lumbar vertebra independently of any count.^{3,21} If those differences hold at the boundary itself, the phenotype is decidable from the vertebra in front of the reader without any global count. On this corpus it does hold: separating T12 from L1 on shape alone, with each patient’s own size divided out and cross-validation grouped by case, gives an area under the curve of 0.990 and 97.3% accuracy over 1485 vertebrae. Size is removed because raw

TABLE I Public CT collections at the lumbosacral junction. Black, class present; outlined, graded only to exclude it⁹.

Collection	Scans	Count	L6	Sacrum	S1	Pelvis	Ribs	L. rib	Femora	Grade
CTSpine1K ⁴	1,005									
CTPelvic1K ⁵	1,184									
VerSe ⁸	374									
RibSeg v2 ²²	660									
TotalSegmentator ²³	1,204									
CTSpinoPelvic1K	802									

99 millimeters separate thoracic from lumbar trivially and teach nothing about the junction,
100 where the two are nearly the same size.

101 That is reported as a property of the released measurements rather than as a method:
102 the information needed to settle these variants is present locally.

103 II. ACQUISITION AND VALIDATION METHODS

104 II.A. Sources, and why a crosswalk was necessary

105 Both CTSpine1K⁴ and CTPelvic1K⁵ draw part of their imaging from the TCIA CT
106 colonography collection^{15,16}, and the vertebral and pelvic annotations in both were produced
107 under board-certified radiologist supervision. CTSpine1K’s 1,005 volumes come from four
108 collections and CTPelvic1K’s 1,184 from seven; COLONOG is the only collection common
109 to both and the only one acquired under a single protocol, so this release is its COLONOG
110 subset: 782 patients with a CTSpine1K vertebral annotation and 20 with a CTPelvic1K
111 pelvic annotation only, 802 in all. Neither source released the mapping from an annotation
112 to the specific CT series it was drawn on.

113 That mapping matters because every patient in CT COLONOGRAPHY was scanned
114 *twice*, prone and supine, and an annotation carries a patient identifier and nothing finer.
115 Turning a patient from supine to prone alters lumbar lordosis and the alignment of each
116 segment against the next, so outlines drawn on one series are the wrong *shape* for the other
117 and no rigid realignment recovers them. Such a mask looks competent in isolation and




FIG. 1 How the dataset was built. n counts scans, or corrections, at each step.

reveals itself only overlaid on the image. The crosswalk (Fig. 1) therefore resolves each annotation to a patient and then, within that patient's volumes, to the series whose bone agrees with it: candidates are thresholded at bone attenuation and scored by overlap with the mask.

II.B. Fused and split records

CTSpine1K and CTPelvic1K each annotated a patient independently, so their labels are not guaranteed to come from the same CT series (Fig. 1). This turned out to matter for 351 patients, nearly half the cohort: the two collections had labeled different acquisitions of the same person, one prone and one supine, with neither recording which. Because posture

changes spinal alignment, a spine label from one acquisition cannot simply be paired with a pelvic label from the other. These 351 records are released but held out of any analysis that assumes a single posture; two labeled postures of one patient are a resource for studying alignment, not a defect.

II.C. Densification, and why it is asymmetric

A corpus in which half the records lack a pelvis is not a spinopelvic dataset, so the missing structures were pseudolabeled, asymmetrically (Fig. 1). Pelvis onto spine-only records is the safe direction: the sacrum and innominate bones carry no enumeration to get wrong, and quality is reported as held-out performance on the *pelvic-only* records, withheld from training for that purpose. The reverse direction carries the whole problem this dataset addresses, since a pseudolabeled spine must commit to a count and on a transitional vertebra commits to the commoner reading; those 20 records were corrected by hand, tractable at that scale.

II.D. Ribs

No public rib annotation covers this cohort, and the two available automatic sources fail in complementary ways.

Möller’s binary rib network²⁴ reports high accuracy, and its authors show that stump ribs can be classified from a partial field of view; that result concerns the morphological features computed on a rib once segmented. On this corpus the segmentation itself failed on ribs that were *partially outside* the field of view: a rib crossing the boundary of an abdominal acquisition was liable to be missed altogether rather than segmented to the edge. Its authors do not report this: their evaluation excluded subjects whose last rib was not visible, so it is recorded here as an observation from this cohort. In a colonography cohort the lower ribs are routinely clipped, and those are the ribs the count depends on.

TotalSegmentator²³ does segment clipped ribs and supplies the per-rib numbering that a binary network cannot. Its characteristic failure is at the other end of the bone: the most medial portion of the rib (roughly the last tenth to sixth approaching the costovertebral joint) is frequently absent. A rib truncated at its medial end never reaches the vertebra it articulates with, so the rib–vertebra incidence check (the quality-control gate on the rib

class labels, which matches each rib to the vertebra its number implies) has nothing to test.

The two were therefore unioned (Fig. 1), TotalSegmentator supplying the numbering and the medial-clipped cases, Möller the detection where its output is more complete.

The result is a pseudolabel whose review was triaged rather than exhaustive. A label-based rule flagged every record in which one rib bone carried more than one identifier, or in which a rib failed to reach the vertebra it articulates with, since a single bone carrying two numbers is a numbering error by construction, whatever the anatomy. The rule selected 152 records; 148 were reviewed and corrected, and four were not and are identified in the release. Reviewers were medical students working through a student research consortium²⁵, trained against a written labeling protocol, with every correction recorded against the annotator who made it. Records passing the rule were not individually inspected, so the layer only guarantees freedom from the failure modes the rule detects, no more. The residual rib-vertebra offsets reported in Sec. II.G are the independent check on what the rule missed.

II.E. Label scheme and counting anchors

The scheme is VerSe-native: vertebrae keep their VerSe identifiers (C1–C7 = 1–7, T1–T12 = 8–19, L1–L6 = 20–25, sacrum 26), and every non-VerSe structure takes a fixed identifier above that range: S1 carved from the sacrum (29), hips (30–31), femora (32–33), ribs (34–45 left, 46–57 right), lumbar ribs (74–75) and surgical hardware (76–82). Identifiers 58–73 are unassigned: a soft-tissue block once reserved there was never populated and is retired.

Two anchors make the count decidable, and Fig. 2 shows both: the lowest rib-bearing vertebra above, **S1** carved from the sacrum below. Each earns its place by what it supports rather than by being an endpoint. S1 supplies the sacral endplate landmark that sacral slope and pelvic incidence are measured from, and with both brackets present L5 and L6 are distinguishable where a spine-limited view alone would not be.

How S1 is cut. The sacrum’s outer boundary is the source annotation and is never altered; S1 is the cranial part of that same bone, separated by a plane. The plane’s normal is the sacrum’s own cranio-caudal axis, made orthogonal to a left–right axis measured directly rather than assumed, so as to remove the patient’s rotation within the scanner. That rotation is not negligible: median 4.5°, maximum 16.1°, and 508 of 801 records above 3°. The cut preserves the sub-division volume of the preceding release, changing the orientation




FIG. 2 The two anchors, from the released labels. Red, lowest rib-bearing vertebra and its rib; blue, S1; yellow, the rib-free vertebrae between. (b) and (c) both count four.

185 of the S1/S2 boundary and not how much sacrum is called S1; per-record geometry ships
186 with the archive.

187 Neither anchor is novel on its own: TotalSegmentator²³ carries a separate S1 class and
188 per-level ribs, and VerSe⁸ carries L6, whose identifier this release keeps, so its L6 *is* VerSe's.
189 L6 is retained because a scheme without it cannot record a six-lumbar spine at all.

190 **A rib on a lumbar body receives its own class**, and this is the one class here with no
191 published counterpart. A lumbar rib and a stump rib are the same object under two counts,
192 so assigning it a number would commit the label to one reading of an enumeration anomaly.
193 A scheme that numbers every rib 1–12 has nowhere to put a thirteenth: the annotator must
194 either call it rib 12, asserting the vertebra beneath it is thoracic, the very question at issue,
195 or discard it. TotalSegmentator carries ribs 1–12 per side and no such class, and VerSe's
196 T13 is a vertebra rather than a rib. This scheme can record both: VerSe's T13 (28) for a
197 thoracic-type thirteenth vertebra with a true rib, and 74/75 for a lumbar-type body bearing
198 a rudimentary one, so either reading is recorded as itself. No record here carries a T13; the
199 sixteen lumbar-rib cases were read as lumbar-type bodies. Where a border vertebra could
200 not be settled by its readers the label stands as read, and reclassifying those records with
201 a classifier trained on local vertebral morphology (Sec. V) is the intended path for a later
202 release.

203 Figure 2 shows why the interval is the primitive and the label is derived. Among the
204 records with four rib-free vertebrae, some reach that count through a lumbar rib and some
205 without one; the two are indistinguishable by the count alone and are distinguished here
206 only because rib presence is itself labeled.

207 Sixteen records carry a lumbar rib: thirteen bilateral and three unilateral, all on the
208 right, a count too small to support any statement about side preference.

209 Eighteen records carry an L6. Fourteen are labeled lumbarization, two sacralization,
210 one semi-sacralization and one normal, which is worth stating because it shows the two
211 annotation axes disagreeing in the open: a record can carry six lumbar-type vertebrae and
212 still be read as a sacralization, or as unremarkable, depending on where the reader started
213 the count.

214 Both counts are derived from the released label volumes, and so are the manifest fields
215 that report them. That is a correction: through v5, `has_l6` was true for one record that
216 contains no L6 and false for all eighteen that do, `has_lumbar_rib` was false everywhere, and
217 `n_lumbar_labels` read zero in 799 of 802 records. The fields are rebuilt for this release by
218 counting identifiers 25 and 74–75 in each volume, with `lumbar_rib_side` added alongside
219 them.

220 II.F. Computational tools

221 *Access.* Every step is performed by open-source code archived with the dataset, together
222 with the reference loader, the label scheme and the quality-control scripts that generated
223 the tables in this manuscript.

224 *Versions and parameters.* Segmentation of femora, the sacral sub-division and the per-
225 level rib numbering used TotalSegmentator²³ (`total` task, release 2.x) on a single graphics
226 processing unit. The binary rib network is Möller’s²⁴ released weights, applied through the
227 nnU-Net v2 framework²⁶. The pelvic completion for records whose pelvis was absent used
228 a five-fold nnU-Net v2 ensemble, applied out-of-fold: each record is completed by the fold
229 that never trained on it, so no record is completed by a model that has seen it. The five
230 fold checkpoints, with their nnU-Net plans and dataset descriptor, are released alongside the
231 data (<https://huggingface.co/anonymous-mlhc/spinopelvic-seg-checkpoints>). Im-
232 age handling throughout used NiBabel and SciPy; no image is resampled or reoriented when
233 a label is written, so every label shares its image’s grid and affine exactly.

234 *Generative artificial intelligence.* A large language model (Claude, Anthropic) assisted
235 with drafting and editing the manuscript text and with writing the analysis and figure code,
236 all under author review. It was not used to generate, annotate or interpret imaging data;

every number reported here is computed by the released code from the released volumes.

II.G. Validation

Validation is layered: measurements on a corpus that has not established internal consistency do not look broken, they look finished.

Geometric invariants. Every case is checked for label geometry agreeing with its CT in shape, affine and voxel spacing; for identifiers within the published scheme; for a non-empty label; and for sided structures falling on the correct sides, with the left–right axis read from the affine rather than assumed. The check exists because a renumbering pass once wrote a reoriented array under a canonical affine and transposed a label away from its image, a failure invisible downstream.

The check compares *each sided pair separately*: testing the ribs alone passed all 802 records while four carried `left_hip` on the patient’s right, and pooling all sided structures into one test still missed one of those four, because a large correctly-sided structure masks a smaller transposed one. A gate that passes everything is not evidence of a clean corpus until it has been shown capable of failing.

Rib–vertebra incidence. Each rib is matched to its nearest vertebra and compared against the vertebra its number implies. Of 11,548 ribs, 5,788 are not evaluable because the expected vertebra lies outside the field of view and 11 have no vertebra within the anchor distance; of the 5,749 evaluable, 5,747 match, two (0.035%) are offset by one level, and no case is misnumbered. The denominator is stated because quoting the two offsets against all 11,548 ribs would halve the rate by counting ribs the check never examined. Both offsets are field-of-view truncations on individually characterized cases, one at +1 level and one at −1, reported rather than suppressed.

The same check carries an invariant that reads as a null result. Of the 5,749 evaluable ribs, *zero* are assigned to a lumbar vertebra, not because no lumbar ribs exist, since 16 records carry one, but because such a rib takes class 74 or 75 and never enters the numbered set the test examines. A numbered rib landing on a lumbar vertebra would be a counting error of exactly the kind this dataset exists to expose.

Agreement with published values. Derived spinopelvic measures are compared against Vialle *et al.*²⁷ (Table II). Pelvic incidence is the strongest of these checks because it is a

TABLE II Derived measures against a standing reference ($n = 260$)²⁷; this cohort is supine.

Measure	This dataset (supine)	Reference (standing)
Pelvic incidence	54.7°	54.7 ± 10.6
Pelvic tilt	17.3°	13 ± 6
Sacral slope	36.9°	41.0 ± 8.4
Lumbar lordosis (supine)	52.7°	43 ± 11.2 (standing)
PI – LL mismatch	2.6°	centered near 0

267 morphological property of the pelvis rather than a posture, and so is directly comparable to
268 a standing reference cohort.

269 II.H. Castellvi typing

270 Every record whose vertebral count departs from five carries a Castellvi grade²⁸ in the
271 released manifest: 33 cases, typed I–IV with the a/b unilateral–bilateral qualifier. The grade
272 and the count are *different axes*: grade IIIb occurs here at rib-free counts of four, five and
273 six alike, and seven of the 33 graded cases carry a normal count of five.

274 Two radiology resident physicians graded every case independently and resolved their
275 six disagreements by consensus; both reads and the consensus are in the manifest. The
276 distinction the classification turns on, bony fusion against articulation without it³, is the
277 one carrying clinical weight, and it is where the disagreements lay: four of the six were III
278 against II.

279 II.I. Surgical hardware, and why a threshold is not enough

280 Metal is trivial to detect and hard to interpret, and the difference matters here because
281 **an iatrogenic fusion is indistinguishable from a congenital one to a distance**
282 **measurement**: a cage-bridged interspace reads as “no gap” exactly as a fused transitional
283 vertebra does.

284 *II.I.0.a. Detection over-calls by a factor of five.* A 2500 HU threshold inside a shell
285 around the labeled skeleton flags 52 of the 802 records. One of the two radiology resident
286 physicians read all 52 and confirmed 11, rejecting 41 (79%) as contrast, calcification and




FIG. 3 Instrumentation in the release, one exemplar each, drawn through bone. (a) Hip arthroplasty, 80, $n = 8$. (b) Femoral-neck screws, 82, $n = 1$. (c) Sacroiliac screws, 81, $n = 1$. (d) Interbody cages, 77, $n = 1$. Bar, 5 cm.

reconstruction artefact. Attenuation does not separate the two groups (both saturate); site and volume do, every confirmed implant at $2,586 \text{ mm}^3$ or above and every rejection at $1,768 \text{ mm}^3$ or below.

II.I.0.b. The subtype block had to be extended. Identifiers 76–79 name spinal instrumentation (generic, cage, screw-and-rod, plate) and cannot describe what this cohort holds, so **80 arthroplasty**, **81 sacroiliac screw** and **82 osteosynthesis** were added: osteosynthesis holds parts of one bone together where arthroplasty replaces a joint, and a shape rule calls a femoral stem a rod. The confirmed cases are eight arthroplasties, one osteosynthesis, one sacroiliac fixation and one pair of interbody cages (Fig. 3).

II.I.0.c. Metal outranks bone, and that is a subtraction. In every confirmed case the implant was already inside a bone label, so naming it reclaims $1,538,852$ voxels rather than adding them. In eight of these cases the femoral head that pelvic incidence and tilt are measured from is an implant.

III. DATA FORMAT AND USAGE NOTES

Each record is a gzipped NifTI image/label pair sharing an identical 4×4 affine, canonicalized to PIR, requiring no resampling. Masks are NifTI rather than DICOM, which is what volumetric segmentation tooling expects.

Splits are patient-grouped and LSTV-stratified five-fold, frozen and shipped with the data as `splits_5fold.json`. The five validation folds partition all 802 records, so the release provides cross-validation and *not* a held-out test set; users reporting a single held-out number should carve one out and state which records it holds.

TABLE III Data types and metadata, and the records populating each, from the released manifest.

Field	Records	Available
Identifier, image path, label path, configuration, match type	802	100.0%
Patient position, tube potential, section thickness, kernel	802	100.0%
Transitional label and class	802	100.0%
Spine and pelvic annotation origin	802	100.0%
Instrumentation and partial-annotation flags	802	100.0%
Image/label alignment check	802	100.0%
Spine series identifier; bone-fraction check	782	97.5%
Scanner manufacturer and model	768	95.8%
Sex	749	93.4%
Age	709	88.4%
Pelvic series identifier	713	88.9%
Castellvi grade (consensus of two readers)	33	4.1%

Stratification is on the transitional subtype rather than a binary LSTV flag: every validation fold carries three sacralization and three lumbarization records. With 15 of each across 802 that is the most stratification can guarantee, and it is why a fold-level metric on the rare classes carries an error bar far wider than its own decimal places.

Table III lists the key data types and metadata against the fraction of records for which each is populated.

The archive of record is Zenodo (<https://doi.org/10.5281/zenodo.22139642>). Code and documentation are on GitHub at the tagged release commit; any model-hub copy is a convenience mirror, not the archive.

IV. DESCRIPTIVE ANALYSIS

The cohort is a colorectal cancer screening population of 802 records. 738 are recorded female (393) or male (345); 11 carry DICOM’s *other* value and 53 no sex field. Age is present for 709 (median 59, range 50–89). All 802 are counted in totals; the eleven records outside the female/male dichotomy are excluded from sex-stratified measures rather than




FIG. 4 Count-free measures. (a) Vertebrae between the lowest rib and the sacrum. (b) Lowest rib length over the rib above. (c) Lowest lumbar transverse span against its gap to the ala; transitional labels marked.




FIG. 5 Spinopelvic measures against standing reference bands (mean $\pm$ SD)²⁷.

folded into “unrecorded”. The lower age bound is the screening guideline, not a selection applied here.

Count-free measures (Fig. 4). Five rib-free vertebrae above the sacrum is typical; four and six are where transitional anatomy sits, and which of the two is present does not by itself determine what to call it. The lowest-rib length ratio is bimodal.

Level gradients and spinopelvic measures (Fig. 5). Vertebral dimensions increase caudally; pelvic incidence does not vary with age while pelvic tilt does, the signature of a pelvis retroverting as a spine ages.

Opportunistic measures. Every scan carries a vertebral bone density measurement at no additional dose.²⁹ L1 trabecular attenuation is reported at the published standard site.

V. POTENTIAL APPLICATIONS AND LIMITATIONS

The principal use is the *spinopelvic interface*. Spine and pelvis carry radiologist-sourced annotation on one coordinate frame in 802 records, making sacral morphology measurable against the lumbar column rather than in isolation: the sacral endplate and its slope, the alar corridors constraining S2-alar-iliac and iliac screw trajectories, and the L5–S1 geometry a lumbosacral fusion must cross. Pelvic incidence is derived from the segmented sacrum and femoral heads rather than from landmarks placed on a radiograph, so it is reproducible from the labels themselves.

The transitional layer serves that use rather than competing with it: no trajectory can be planned across a junction whose segments cannot be named.

V.A. Three uses the released measurements support

Naming a level from its own geometry. Wrong-level spine surgery runs at roughly one in 3,110 spinal procedures, and the cause most often cited is the transitional anatomy this cohort was assembled around.³⁰ Per-level body height, canal width, endplate width, transverse process span and wedge ratio are released for T11–L5 across all 802 records, and because the labels were assigned against the twelfth rib rather than by enumeration, a model trained on them is not learning the convention that fails.

Reference morphometry with its spread. The values a surgeon consults for vertebral body height, canal dimensions and pedicle width come from cadaveric series of a few dozen specimens, plotted as single values with no indication of spread.²⁰ They cannot tell a reader whether the patient in front of them is unusual, because they never show what usual looks like. The same measures are given here from 802 records with the distribution attached (Fig. 6), reproducing two properties of the classical figures (body height crosses over, dorsal exceeding ventral at T11–T12 and reversing by L4–L5, and pedicle width widens steeply into the lower lumbar spine) and adding the width of each distribution, which is the part needed to judge an individual.

Patient-specific models. Planning is moving toward patient-specific biomechanical models rather than population norms.³¹ Such a model needs spine, pelvis and femoral heads in one frame, which is this release’s construction; in 351 patients two postures give a within-patient




FIG. 6 Morphometry by level: median, interquartile range, 5th–95th percentile, n at right. (a) Body height. (b) Endplate width. (c) Canal. (d) Pedicle width. (e) Disc height. (f) Trabecular attenuation.




FIG. 7 Records carrying each vertebral level, cranial to caudal. Counts per identifier in Table S1.

change in alignment against which a predicted postural response can be checked.

Its limitations are stated at the level of detail a reader would need to catch them independently.

364 *Coverage.* Thoracic ground truth is field-of-view limited (Fig. 7) and does not extend to
365 T1. Postural angles are supine; pelvic incidence is not postural and needs no such caveat.
366 The cohort is a colorectal screening population aged 50 and over, so its distributions should
367 not be read as representative of a surgical one.

368 *Declared but empty, and confirmed so.* The partial-annotation sentinel (255) is declared
369 in the scheme and occurs in no record, established by counting identifiers across all 802
370 released volumes rather than asserted. It is retained because the protocol is part of the
371 scheme and a later release may ship a partially traced record. The hardware identifiers were
372 in this position until v6 and are no longer: 76–82 are populated in the eleven instrumented
373 records described in Sec. II.I.

374 *Strength of the labels varies by structure, and the release does not average over that.*
375 Vertebral labels derive from radiologist-supervised source annotations, corrected where those
376 sources were wrong. Pelvic labels on records that lacked one are pseudolabeled. The rib
377 layer is a pseudolabel whose human review was triaged by an automated rule rather than
378 exhaustive, so its guarantee is the absence of the failure modes that rule detects. Transitional
379 labels derive from two independent sources and are not uniformly adjudicated, which is
380 precisely why the measures reported here are count-free, and the Castellvi grades are a
381 consensus of two radiology resident physicians. The sacral sub-division inherits its level
382 from an automatic method, and in 100 records that level places S1 above half the sacrum
383 or below fifteen percent of it, which no first sacral segment is; those records are flagged in
384 the released quality-control table and should be filtered before any measurement depending
385 on the sacral endplate.

386 *Structures are not universally present.* The same census shows the release is not uniformly
387 complete. Sacrum, hips and femora are present in all 802 records. One record carries no S1,
388 one lacks T12, and nine lack an L5 identifier, in every case because the structure lies outside
389 the field of view or, for L5, because the lowest lumbar segment is labeled L6 or incorporated
390 into the sacrum. Users filtering on the presence of the caudal anchor should filter explicitly
391 rather than assume it.

392 VI. DISCUSSION

393 *Comparison with related material.* Against the collections in Table I, the contribution
394 is the crosswalk placing spine and pelvis on the same series, together with classes for the
395 anatomy an enumeration anomaly actually produces. VerSe⁸ comes closest in intent: it
396 also enriched its cohort for transitional anatomy and graded each case by Castellvi. The
397 difference is what a graded case leaves behind. VerSe names the anomaly but does not
398 segment the sacrum, so the fused or articulating structure the grade describes has no mask
399 to go with it. This release segments that structure directly, as L6 or as a separately carved
400 S1, so a user has the anatomy itself to measure rather than only a grade that describes
401 it. TotalSegmentator²³ segments far more of the body than any source used here but has
402 no class for a sixth lumbar vertebra or a rib on a lumbar body; what this release adds is
403 not more anatomy per scan but the classes an enumeration anomaly needs in order to be
404 recorded accurately.

405 *Limitations.* The cohort comes from one acquisition protocol of a colorectal-screening
406 population aged 50 and over, so how well it generalizes to a pathologic population is untested.
407 The rib layer’s human review was triaged by an automated rule rather than performed on
408 every bone, so it guarantees only the absence of the failure modes that rule detects.

409 *Future directions.* The version history shows the scheme being extended more than once:
410 femora and the S1 anchor, then per-level ribs, then lumbar ribs and surgical hardware, were
411 each folded into the existing release rather than requiring a new dataset. Imaging beyond
412 this cohort’s protocol, and the C2 anchor it cannot supply, are the natural next additions.

413 VII. CONCLUSION

414 CTSpinoPelvic1K places spine, pelvis, per-level ribs and femora on one coordinate frame
415 in 802 records, and gives the anatomy that makes lumbar numbering ambiguous a class of
416 its own rather than recording it as something it is not. Because both counting anchors are
417 explicit, a level can be identified from a field of view that cannot support the conventional
418 count downward from C2, which is every record here.

419 DATA AVAILABILITY

420 The release described here is v8, permanently archived at [https://doi.org/10.5281/](https://doi.org/10.5281/zenodo.22293878)
421 [zenodo.22293878](https://doi.org/10.5281/zenodo.22293878); the concept DOI <https://doi.org/10.5281/zenodo.22139642> resolves
422 to the current version. The archive carries the loader, the label scheme, the quality-control
423 tables and the analysis code under an open-source licence. The repository URL identifies
424 the authors; it is on the title page and will be restored here for the camera-ready version.

425 ACKNOWLEDGEMENTS

426 Acknowledgements are given on the title page, per the journal's double-anonymized re-
427 view policy. This work received no funding.

428 CONFLICT OF INTEREST

429 The authors have no relevant conflicts of interest to disclose.

430 ETHICS

431 This work uses publicly available, de-identified imaging and annotations derived from it.
432 The authors' institutional review board determined that it is not human participant research
433 under 45 CFR 46 and requires no oversight.

434 REFERENCES

- 435 ¹J. Lian, N. Levine, and W. Cho, A review of lumbosacral transitional vertebrae and
436 associated vertebral numeration, *Eur. Spine J.* **27**, 995–1004 (2018), doi:10.1007/s00586-
437 018-5554-8.
- 438 ²Expert Panel on Neurological Imaging, T. A. Hutchins, M. Peckham, et al., ACR Ap-
439 propriateness Criteria low back pain: 2021 update, *J. Am. Coll. Radiol.* **18**, S361–S379
440 (2021), doi:10.1016/j.jacr.2021.08.002.

³G. P. Konin and D. M. Walz, Lumbosacral transitional vertebrae: classification, imaging findings, and clinical relevance, *AJNR Am. J. Neuroradiol.* **31**, 1778–1786 (2010), doi:10.3174/ajnr.A2036.

⁴Y. Deng et al., CTSpine1K: a large-scale dataset for spinal vertebrae segmentation in computed tomography, *Mach. Learn. Biomed. Imaging* **3**, 824–832 (2025), doi:10.59275/j.melba.2025-gf84.

⁵P. Liu et al., Deep learning to segment pelvic bones: large-scale CT datasets and baseline models, *Int. J. Comput. Assist. Radiol. Surg.* **16**, 749–756 (2021), doi:10.1007/s11548-021-02363-8.

⁶H. Möller, R. Graf, J. Schmitt, et al., SPINEPS: automatic whole spine segmentation of T2-weighted MR images using a two-phase approach to multi-class semantic and instance segmentation, *Eur. Radiol.* **35**, 1178–1189 (2025), doi:10.1007/s00330-024-11155-y.

⁷H. Möller, H. Schoen, R. Graf, et al., VERIDAH: solving enumeration anomaly aware vertebra labeling across imaging sequences, 2026, arXiv:2601.14066. doi:10.48550/arXiv.2601.14066.

⁸A. Sekuboyina et al., VerSe: a vertebrae labelling and segmentation benchmark for multi-detector CT images, *Med. Image Anal.* **73**, 102166 (2021), doi:10.1016/j.media.2021.102166.

⁹H. Liebl et al., A computed tomography vertebral segmentation dataset with anatomical variations and multi-vendor scanner data, *Sci. Data* **8**, 284 (2021), doi:10.1038/s41597-021-01060-0.

¹⁰Y. Otake, S. Schafer, J. W. Stayman, et al., Automatic localization of vertebral levels in x-ray fluoroscopy using 3D-2D registration: a tool to reduce wrong-site surgery, *Phys. Med. Biol.* **57**, 5485–5508 (2012), doi:10.1088/0031-9155/57/17/5485.

¹¹S.-F. L. Lo, Y. Otake, V. Puvanesarajah, et al., Automatic localization of target vertebrae in spine surgery: clinical evaluation of the LevelCheck registration algorithm, *Spine (Phila. Pa. 1976)* **40**, E476–E483 (2015), doi:10.1097/BRS.0000000000000814.

¹²T. De Silva, S.-F. L. Lo, N. Aygun, et al., Utility of the LevelCheck algorithm for decision support in vertebral localization, *Spine (Phila. Pa. 1976)* **41**, E1249–E1256 (2016), doi:10.1097/BRS.0000000000001589.

¹³D. S. Kreiner, W. O. Shaffer, J. L. Baisden, et al., An evidence-based clinical guideline for the diagnosis and treatment of degenerative lumbar spinal stenosis (update), *Spine J.*

473 **13**, 734–743 (2013), doi:10.1016/j.spinee.2012.11.059.

474 ¹⁴American College of Surgeons Committee on Trauma, ACS TQIP best practices guidelines
475 in imaging, American College of Surgeons, Chicago, 2018.

476 ¹⁵K. Smith et al., Data from CT COLONOGRAPHY, The Cancer Imaging Archive, 2015,
477 doi:10.7937/K9/TCIA.2015.NWTESAY1.

478 ¹⁶K. Clark et al., The Cancer Imaging Archive (TCIA): maintaining and operating a public
479 information repository, *J. Digit. Imaging* **26**, 1045–1057 (2013), doi:10.1007/s10278-013-
480 9622-7.

481 ¹⁷M. R. Zindrick et al., Analysis of the morphometric characteristics of the thoracic and
482 lumbar pedicles, *Spine (Phila. Pa. 1976)* **12**, 160–166 (1987), doi:10.1097/00007632-
483 198703000-00012.

484 ¹⁸M. M. Panjabi et al., Thoracic human vertebrae: quantitative three-dimensional anatomy,
485 *Spine (Phila. Pa. 1976)* **16**, 888–901 (1991), doi:10.1097/00007632-199108000-00006.

486 ¹⁹M. M. Panjabi et al., Human lumbar vertebrae: quantitative three-dimensional anatomy,
487 *Spine (Phila. Pa. 1976)* **17**, 299–306 (1992), doi:10.1097/00007632-199203000-00010.

488 ²⁰E. C. Benzell, *Biomechanics of Spine Stabilization*, Thieme, New York, 3rd edition, 2015.

489 ²¹R. J. Hughes and A. Saifuddin, Numbering of lumbosacral transitional vertebrae on
490 MRI: role of the iliolumbar ligaments, *AJR Am. J. Roentgenol.* **187**, W59–W65 (2006),
491 doi:10.2214/AJR.05.0415.

492 ²²L. Jin et al., RibSeg v2: a large-scale benchmark for rib labeling and
493 anatomical centerline extraction, *IEEE Trans. Med. Imaging* **43**, 570–581 (2024),
494 doi:10.1109/TMI.2023.3313627.

495 ²³J. Wasserthal et al., TotalSegmentator: robust segmentation of 104 anatomic structures
496 in CT images, *Radiol. Artif. Intell.* **5**, e230024 (2023), doi:10.1148/ryai.230024.

497 ²⁴H. Möller et al., Automated thoracolumbar stump rib detection and analysis in a large
498 CT cohort, 2025, arXiv:2505.05004. doi:10.48550/arXiv.2505.05004.

499 ²⁵[Reference to the authors’ own work, removed for double-anonymized review.]

500 ²⁶F. Isensee, P. F. Jaeger, S. A. A. Kohl, J. Petersen, and K. H. Maier-Hein, nnU-Net:
501 a self-configuring method for deep learning-based biomedical image segmentation, *Nat.*
502 *Methods* **18**, 203–211 (2021), doi:10.1038/s41592-020-01008-z.

503 ²⁷R. Vialle, N. Levassor, L. Rillardon, A. Templier, W. Skalli, and P. Guigui, Radiographic
504 analysis of the sagittal alignment and balance of the spine in asymptomatic subjects, *J.*

505 Bone Joint Surg. Am. **87**, 260–267 (2005), doi:10.2106/JBJS.D.02043.

506 ²⁸A. E. Castellvi, L. A. Goldstein, and D. P. K. Chan, Lumbosacral transitional vertebrae
507 and their relationship with lumbar extradural defects, Spine (Phila. Pa. 1976) **9**, 493–495
508 (1984), doi:10.1097/00007632-198407000-00014.

509 ²⁹P. J. Pickhardt, B. D. Pooler, T. Lauder, A. M. del Rio, R. J. Bruce, and N. Binkley,
510 Opportunistic screening for osteoporosis using abdominal computed tomography scans
511 obtained for other indications, Ann. Intern. Med. **158**, 588–595 (2013), doi:10.7326/0003-
512 4819-158-8-201304160-00003.

513 ³⁰N. E. Epstein, A perspective on wrong level, wrong side, and wrong site spine surgery,
514 Surg. Neurol. Int. **12**, 286 (2021), doi:10.25259/SNI.402_2021.

515 ³¹E. Beaulieu et al., A review of finite element analysis in spine surgery decision-making, J.
516 Clin. Med. **15**, 2584 (2026), doi:10.3390/jcm15072584.

517 FIGURE CAPTIONS

518 **Figure 1.** How the dataset was built. n counts scans, or corrections, at each step.

519 **Figure 2.** The two anchors, from the released labels. Red, lowest rib-bearing vertebra and
520 its rib; blue, S1; yellow, the rib-free vertebrae between. (b) and (c) both count four.

521 **Figure 3.** Instrumentation in the release, one exemplar each, drawn through bone. (a)
522 Hip arthroplasty, 80, $n = 8$. (b) Femoral-neck screws, 82, $n = 1$. (c) Sacroiliac screws, 81,
523 $n = 1$. (d) Interbody cages, 77, $n = 1$. Bar, 5 cm.

524 **Figure 4.** Count-free measures. (a) Vertebrae between the lowest rib and the sacrum. (b)
525 Lowest rib length over the rib above. (c) Lowest lumbar transverse span against its gap to
526 the ala; transitional labels marked.

527 **Figure 5.** Spinopelvic measures against standing reference bands (mean $\pm$ SD)²⁷.

528 **Figure 6.** Morphometry by level: median, interquartile range, 5th–95th percentile, n at
529 right. (a) Body height. (b) Endplate width. (c) Canal. (d) Pedicle width. (e) Disc height.
530 (f) Trabecular attenuation.

531 **Figure 7.** Records carrying each vertebral level, cranial to caudal. Counts per identifier in
532 Table S1.